CS442
Assignment 2

University of Waterloo

Winter 2023

This assignment tests your understanding of the content of Modules 2 and 3. You will implement a λ-calculus reducer in GNU Smalltalk. The assignment is divided into four parts, but you should submit only one file: `a2.st`.

In this and all assignments, any behavior which we do not explicitly define will not be tested, so you may define it however you wish, or allow your program to fail. Make sure though that it actually *is* undefined; ask on Piazza if you're unsure.

Submit all code via UWaterloo submit, e.g.:

```
submit cs442 a2 .
```

This assignment is due on Friday, March 3rd, by 12PM *NOON, NOT MIDNIGHT*, Eastern time.

The λ-Calculus in Smalltalk

You are provided with `lambda.st`, an implementation of the types for the syntax tree of the λ-calculus in GNU Smalltalk. It implements the following classes:

```
Object subclass: LambdaExpr [
    isVar.
    isAbs.
    isApp.
    ifVar: varBlock ifAbs: absBlock ifApp: appBlock.
    reduceWith: block steps: steps.
    freeVars.
    printString.
]

LambdaExpr subclass: LambdaVar [
    | name |
    LambdaVar class >> withName: name.
    dup.
    isVar.
    name.
    freeVars: map.
    displayString.
]
```

```
LambdaExpr subclass: LambdaAbs [
    | var body |
    LambdaAbs class >> withVar: var body: body.
    dup.
    isAbs.
    var.
    body.
    freeVars: map.
    displayString.
]

LambdaExpr subclass: LambdaApp [
    | rator rand |
    LambdaApp class >> withRator: rator rand: rand.
    dup.
    isApp.
    rator.
    rand.
    freeVars: map.
    displayString.
]

Object subclass: LambdaParser [
    | text index tok |
    LambdaParser class >> new: text.
    LambdaParser class >> parse: text.
    parse.
]
```

The `LambdaExpr` class serves as a superclass for all λ-calculus expressions. Each of its methods are only meant to be used on subclasses (that is, in C++ or Java terms, its methods are all abstract, though Smalltalk has no abstract methods). Its `isVar`, `isAbs`, and `isApp` methods each return `true` if the expression is of the named type, or `false` otherwise. The `ifVar:ifAbs:ifApp:` method lets you easily "switch" based on the type; each block will be evaluated only if the value is of the given type. The `freeVars` method returns the free variables of the expression, as a `Dictionary` mapping the name of the free variable to the `LambdaVar` that uses it in the expression; generally, all that you will need from this dictionary is the presence of a free variable, i.e., `includesKey:`. `printString` is overridden to fall through to `displayString`, and `displayString` is overridden in all child types to return a λ-calculus expression as a string, with λ replaced by the caret or circumflex symbol, `^`.

The `reduceWith:steps:` method is a simple method to reduce the expression using the block given as the first argument, maximally as many steps as is given as the second argument. Its purpose will become clearer in the assignment parts on reduction.

We will not repeat the behavior of the shared methods, only the unique methods per each subclass of `LambdaExpr`.

Every subclass of `LambdaExpr` includes a `dup` method, which duplicates the expression, including all children.

The `LambdaVar` class represents a variable. A `LambdaVar` is constructed by `LambdaVar class >> withName:`, which expects the variable name as an argument. The variable name should be a string starting with an alphabetical character, unless this expression is using deBruijn indices and the variable is bound, in which case it should be a number (*not* as a string). The `name` method returns this name.

The `LambdaAbs` class represents an abstraction. A `LambdaAbs` is constructed by `LambdaAbs class >> withVar:body:`, which expects a string variable name (*not* a `LambdaVar`) and a `LambdaExpr` (the body) as arguments. `var` returns the variable, and `body` returns the body. If using deBruijn indices, the variable should be `nil`.

The `LambdaApp` class represents an application. A `LambdaApp` is constructed by `LambdaApp class >> withRator:rand:`, which expects the rator and rand as `LambdaExpr`s. `rator` and `rand` return the rator and rand, respectively.

Finally, the `LambdaParser` class is a parser for λ-calculus strings, which converts such strings into `LambdaExpr`s. The λ-calculus strings accepted by this parser should have λ replaced by `^`. It can be used in one of two ways: either by constructing an instance of `LambdaParser` and calling `parse` on it, or by using `LambdaParser class >> parse:` directly.

For instance, to parse the string `'^f.^x.f (f x)'`, you can use either

```
(LambdaParser new: '^f.^x.f (f x)') parse
```

or

```
LambdaParser parse: '^f.^x.f (f x)'
```

The latter is, of course, implemented in terms of the former. Note that since this implementation of the λ-calculus allows multi-character variable names, `f x` is not the same as `fx`. The first is an application of a variable to a variable, and the second is a variable. Be careful to separate variables with spaces for this reason.

Several examples are included in the "demonstration" section of this document.

1   de Bruijn

Write a file, **a2.st**, which defines at least the following class:

```
Object subclass: Lambda [
    Lambda class >> new: exp.
    toDeBruijn.
]
```

You may (and should!) include more methods and/or classes to help your implementation.

A `Lambda` will eventually be a λ-calculus reducer. For the time being, it's just a de-Bruijn-translator. A `Lambda` is created with a `LambdaExpr` (`exp`), which we will call "the internal expression", which should be stored and updated by all methods.

The `toDeBruijn` method converts a non-de-Bruijn-indexed expression to a de-Bruijn-indexed one, and both returns the de-Bruijn-indexed expression and sets the internal expression to it. All abstractions should be replaced with abstractions that have `nil` as the variable (*not* the string `'nil'`), and all variables should be replaced with numbers at the appropriate depth. Remember than de Bruijn indices are one-indexed, so $\lambda x.\,x$ translates as $\lambda.\,1$, not $\lambda.\,0$.

Any free variables in the expression should be left unchanged. Note that as a consequence of this fact, `deBruijn` can be safely repeated (i.e., it's idempotent): if de Bruijn indices are interpreted as non-de-Bruijn variable names, they will always be free, and so won't be replaced.

`toDeBruijn` may mutate the expression and any of its child objects, or it may create a new expression, or any combination thereof.

Because of how substitution is performed in the λ-calculus, `toDeBruijn` will be used in testing *all* of the methods defined in the rest of this assignment, so make sure you get it right!

Hints

It would be difficult to implement this with `toDeBruijn` alone. You will need to implement something that carries the list of currently-defined variables, so that you can look up a variable in that list and replace it.

You may want to implement that as a method on `LambdaExpr`'s children. If so, don't modify **lambda.st**, as you will not be submitting that file. You can extend an existing class using GNU Smalltalk's `extend` syntax, e.g.:

```
LambdaVar extend [
    toDeBruijn: map [
        ...
    ]
]
```

You would be wise not to *replace* any existing methods in `LambdaVar`, since our tests could use any of them, but you may safely *add* any methods you please. In fact, you can add methods to any class you want, even internal classes!

2 Applicative Order Evaluation

Extend `a2.st`, adding at least the following methods to `Lambda`:

```
Object subclass: Lambda [
    ...
    aoe.
    aoe: steps.
]
```

The `aoe` method performs a *single* reduction step using applicative order evaluation on the expression, returns the reduced expression, and updates the internal expression to match the returned expression. Since the internal expression is updated, if `aoe` is called repeatedly, multiple steps of reduction are taken. If there is no reduction step for AOE (i.e., reduction is complete), then it should return `nil`, and set the `Lambda`'s internal expression to `nil` (which should cause any further calls to `aoe` to fail).

The `aoe:` method performs a specified number of steps of AOE, and returns the result. If fewer than the specified number of steps ends the reduction, then the final reduced expression should be returned, *not* `nil`. The internal expression should be updated to the same result. Note that this means the internal expression should never be updated to `nil` by `aoe:`, even if `aoe` would have updated it to `nil`!

The `aoe:` method is best implemented by `reduceWith:steps:`, perhaps indirectly. The `reduceWith:steps:` method calls the one-argument block passed to `reduceWith:` the number of times specified by `steps:`. The argument to the block on its first evaluation is the `LambdaExpr` that `reduceWith:steps:` was called on, and each subsequent evaluation of the block takes the previous return from the block as its argument, so that it can be used to iteratively reduce an expression. If the block returns `nil`, then `reduceWith:steps:` stops, returning the *previous* value. Read `lambda.st` for more details. When using `reduceWith:steps:`, the block passed to `reduceWith:` should perform a single step of reduction on its argument and return the reduced version. You will probably want to create an internal helper method, used by both `aoe` and `aoe:`, so that their different returns can be accounted for.

You may perform reduction directly on a λ-calculus expression, or you may perform de Bruijn rewriting first. The result of `aoe` may or may not use de Bruijn indices, by your preference. However, the input to `new:` will never use de Bruijn indices, so if you wish to use de Bruijn indices, you should probably modify `new:` (or an `init:` method it uses) to perform this step. There is no method to directly extract the expression, so it's safe to keep your internal expression in either form. You are recommended *not* to use de Bruijn indices, because the complication they avoid during substitution (renaming) is replaced by more subtle complications (renumbering).

`aoe` and `aoe:` may mutate the expression and any of its children, or may create a new expression, or any combination thereof. If you use mutation, make sure to use `dup` during substitution; having two references to the same subexpression will cause some extremely confusing behavior!

Hints

Be *very* careful about reduction steps: `aoe` should represent a *single* application of AOE's →, not two or three! Order is also important, and you will be tested on whether you've reduced the correct part of the expression. Note that in future assignments, we will usually not ask for individual steps in this way; we are asking for this only for the λ-calculus.

Reduction requires substitution, so you will presumably want to implement a substitution method. Substitution requires the ability to create a "new" variable name. Our version of the λ-calculus in this assignment is more forgiving in variable names than in Module 2; in particular, they may be of any length. One simple technique to generate fresh variable names is to carry a counter, and increment it every time you generate a new name. An even simpler technique uses a Smalltalk-specific trick: every object has a *hash*, and the hash is unique to that object. So, if you need a new name, you can use the `hash` method to get a unique number, so long as you're calling it on an adequately unique object. For instance, the canonical implementation of substitution on `LambdaAbs` includes these statements:

```
nvar := var , (self hash asString).
body := body substitute: var for: (LambdaVar withName: nvar).
var := nvar.
```

As with `toDeBruijn`, you may find reduction easier to do by extending the children of `LambdaExpr`.

Remember when testing that most demonstrations of the λ-calculus use shorthand, e.g. ⟦true⟧ , which isn't truly part of the λ-calculus, and so isn't supported by our reducer. If you have an expression E that uses the shorthand $x = M$, you can rewrite E as $(\lambda x.\,E)M$. For example, we can rewrite $\lambda l.\,l$ ⟦true⟧ as $(\lambda true.\,\lambda l.\,l\ true)(\lambda x.\,\lambda y.\,x)$. See the "demonstration" section for some larger examples.

3 Normal Order Reduction

Extend `a2.st`, adding at least the following methods to `Lambda`:

```
Object subclass: Lambda [
    ...
    nor.
    nor: steps.
]
```

The `nor` and `nor:` methods behave like `aoe` and `aoe:`, but using normal order reduction instead of applicative order evaluation. Like `aoe` and `aoe:`, they may mutate the expression, or create a new one.

Note that although it may be strange to do so, mixing and matching `aoe` and `nor` steps is perfectly valid. You must assure that nothing prevents this.

Hints

You must reduce the outermost *reducible* expression. To know whether to reduce the current expression or recurse deeper, you need only to check some types. Remember, an expression is a redex if it's an application and its rator is an abstraction. The `isAbs` method is there for exactly this check!

4 η-Reduction

Extend `a2.st`, adding at least the following methods to `Lambda`:

```
Object subclass: Lambda [
    ...
    eta.
    eta: steps.
]
```

The `eta` and `eta:` methods behave like `aoe` and `aoe:` or `nor` and `nor:`, but using η-reduction instead of β-reduction. Reduce the leftmost, innermost η-reducible expression.

Demonstration

The following demonstrates a possible interaction with `Lambda`. Note that exact variable names within λ-expressions may differ between your implementation and this demonstration, but the result of `toDeBruijn` should always be the same.

```
st> | x s l |
st> x := LambdaParser parse: '(^mul.^two.mul two two) (^m.^n.^f.m(n f)) (^f.^x.f (f x))'.
((((^mul.(^two.((mul two) two))) (^m.(^n.(^f.(m (n f)))))) (^f.(^x.(f (f x)))))
st> l := Lambda new: x.
a Lambda
st> l aoe.
((^two.(((^m.(^n.(^f.(m (n f))))) two) two)) (^f.(^x.(f (f x)))))
st> l aoe.
((((^m.(^n.(^f.(m (n f))))) (^f.(^x.(f (f x))))) (^f.(^x.(f (f x)))))
st> l aoe.
((^n.(^f.(((^f.(^x.(f (f x)))) (n f)))) (^f.(^x.(f (f x)))))
```

```
st> x := l aoe dup.
(^f.((^f.(^x.(f (f x)))) ((^f.(^x.(f (f x)))) f)))
st> l toDeBruijn.
(^.((^.(^.(2 (2 1)))) ((^.(^.(2 (2 1)))) 1)))
st> l := Lambda new: x.
a Lambda
st> l aoe.
nil
st> x := LambdaParser parse: '(^mul.^two.mul two two) (^m.^n.^f.m(n f)) (^f.^x.f (f x))'.
((((^mul.(^two.((mul two) two))) (^m.(^n.(^f.(m (n f)))))) (^f.(^x.(f (f x))))))
st> l := Lambda new: x.
a Lambda
st> x := l aoe: 1000.
(^f.((^f.(^x.(f (f x)))) ((^f.(^x.(f (f x)))) f)))
st> s := x displayString.
'(^f.((^f.(^x.(f (f x)))) ((^f.(^x.(f (f x)))) f)))'
st> l toDeBruijn.
(^.((^.(^.(2 (2 1)))) ((^.(^.(2 (2 1)))) 1)))
st> x := LambdaParser parse: s , 'f x'.
((((^f.((^f.(^x.(f (f x)))) ((^f.(^x.(f (f x)))) f))) f) x)
st> l := Lambda new: x.
a Lambda
st> l aoe.
((((^f.(^x.(f (f x)))) ((^f.(^x.(f (f x)))) f)) x)
st> l aoe.
((((^f.(^x.(f (f x)))) (^x.(f (f x)))) x)
st> l aoe.
((^x.((^x.(f (f x))) ((^x.(f (f x)) x))) x)
st> l aoe.
((^x.(f (f x))) ((^x.(f (f x)) x))
st> l aoe.
((^x.(f (f x))) (f (f x)))
st> l aoe.
(f (f (f (f x))))
st> x := LambdaParser parse: '(^head.^tail.^cons.^zero.^two.(^succ.(^pred.(pred two)) (^n.(tail(n(^p.cons(succ(head p))
    (head p))(cons zero zero))))) (^n.^f.^x.n f(f x))) (^l.l^x.^y.x) (^l.l^x.^y.y) (^h.^t.^s.s h t) (^f.^x.x) (^f.^x.f
    (f x)) f x'.
(((((((((^head.(^tail.(^cons.(^zero.(^two.((^succ.((^pred.(pred two)) (^n.(tail ((n (^p.((cons (succ (head p))) (head
    p)))) ((cons zero) zero)))))) (^n.(^f.(^x.((n f) (f x)))))))))) (^l.(l (^x.(^y.x))))) (^l.(l (^x.(^y.y))))) (^
    h.(^t.(^s.((s h) t))))) (^f.(^x.x))) (^f.(^x.(f (f x))))) f) x)
st> l := Lambda new: x.
a Lambda
st> x := l aoe: 1000.
(f x)
st> l toDeBruijn.
(f x)
st> x := LambdaParser parse: '(^mul.^two.mul two two) (^m.^n.^f.m(n f)) (^f.^x.f (f x))'.
((((^mul.(^two.((mul two) two))) (^m.(^n.(^f.(m (n f)))))) (^f.(^x.(f (f x))))))
st> l := Lambda new: x.
a Lambda
st> l nor.
((^two.(((^m.(^n.(^f.(m (n f))))) two) two)) (^f.(^x.(f (f x)))))
st> l nor.
(((^m.(^n.(^f.(m (n f))))) (^f.(^x.(f (f x))))) (^f.(^x.(f (f x)))))
st> l nor.
((^n.(^f.((^f.(^x.(f (f x)))) (n f)))) (^f.(^x.(f (f x)))))
st> l nor.
(^f.((^f.(^x.(f (f x)))) ((^f.(^x.(f (f x)))) f)))
st> l nor.
(^f.(^x.(((^f.(^x.(f (f x)))) f) (((^f.(^x.(f (f x)))) f) x))))
st> l nor.
(^f.(^x.((^x.(f (f x))) (((^f.(^x.(f (f x)))) f) x))))
st> l nor.
(^f.(^x.(f (f (((^f.(^x.(f (f x)))) f) x)))))
st> l nor.
(^f.(^x.(f (f ((^x.(f (f x))) x)))))
st> x := l nor dup.
(^f.(^x.(f (f (f (f x))))))
st> l toDeBruijn.
(^.(^.(2 (2 (2 (2 1)))))))
st> l := Lambda new: x.
a Lambda
st> l nor.
nil
st> x := LambdaParser parse: '(^mul.^two.mul two two) (^m.^n.^f.m(n f)) (^f.^x.f (f x))'.
((((^mul.(^two.((mul two) two))) (^m.(^n.(^f.(m (n f)))))) (^f.(^x.(f (f x))))))
st> l := Lambda new: x.
a Lambda
st> x := l nor: 1000.
(^f.(^x.(f (f (f (f x))))))
st> s := x displayString.
'(^f.(^x.(f (f (f (f x))))))'
st> l toDeBruijn.
(^.(^.(2 (2 (2 (2 1)))))))
```

```
st> x := LambdaParser parse: s , 'f x'.
((((^f.(^x.(f (f (f (f x)))))) f) x)
st> l := Lambda new: x.
a Lambda
st> l nor: 1000.
(f (f (f (f x))))
st> x := LambdaParser parse: '(^head.^tail.^cons.^zero.^two.(^succ.(^pred.(pred two)) (^n.(tail(n(^p.cons(succ(head p))
    (head p))(cons zero zero))))) (^n.^f.^x.n f(f x))) (^l.l^x.^y.x) (^l.l^x.^y.y) (^h.^t.^s.s h t) (^f.^x.x) (^f.^x.f
    (f x)) f x'.
((((((((^head.(^tail.(^cons.(^zero.(^two.((^succ.((^pred.(pred two)) (^n.(tail ((n (^p.((cons (succ (head p))) (head
    p)))) ((cons zero) zero)))))) (^n.(^f.(^x.((n f) (f x)))))))))))) (^l.(l (^x.(^y.x)))) (^l.(l (^x.(^y.y)))) (^
    h.(^t.(^s.((s h) t))))) (^f.(^x.x))) (^f.(^x.(f (f x)))) f) x)
st> l := Lambda new: x.
a Lambda
st> x := l nor: 1000.
(f x)
st> l toDeBruijn.
(f x)
st> x := LambdaParser parse: '(^head.^tail.^cons.^isNull.^nil.^zero.^succ.(^Y.^F.(^len.(len (cons zero (cons zero nil))
    )) (Y F)) (^f.(^x.f(x x))(^x.f(x x))) (^f.^l.(isNull l) zero (succ (f(tail l))))) (^l.l(^x.^y.x)) (^l.l(^x.^y.y))
    (^h.^t.^s.s h t) (^l.l^h.^t.^x.^y.y) (^s.^x.^y.x) (^f.^x.x) (^n.^f.^x.n f(f x))'.
((((((((^head.(^tail.(^cons.(^isNull.(^nil.(^zero.(^succ.(((^Y.(^F.((^len.(len ((cons zero) ((cons zero) nil)))) (Y F)
    ))) (^f.((^x.(f (x x))) (^x.(f (x x)))))) (^f.(^l.(((isNull l) zero) (succ (f (tail l))))))))))))))) (^l.(l (^x.(^
    y.x))))) (^l.(l (^x.(^y.y))))) (^h.(^t.(^s.((s h) t))))) (^l.(l (^h.(^t.(^x.(^y.y))))))) (^s.(^x.(^y.x)))) (^f.(^
    x.x))) (^n.(^f.(^x.((n f) (f x))))))
st> l := Lambda new: x.
a Lambda
st> l nor: 1000.
(^f.(^x.(f (f x))))
st> x := LambdaParser parse: '^m.^n.^f.^x.m(n f)x'.
(^m.(^n.(^f.(^x.((m (n f)) x)))))
st> l := Lambda new: x.
a Lambda
st> l eta.
(^m.(^n.(^f.(m (n f)))))
st> l eta.
nil
st> x := LambdaParser parse: '^m.^n.^f.^x.n m f x'.
(^m.(^n.(^f.(^x.(((n m) f) x)))))
st> l := Lambda new: x.
a Lambda
st> l eta.
(^m.(^n.(^f.((n m) f))))
st> l eta.
(^m.(^n.(n m)))
st> l eta.
nil
st>
```

A Note on Testing

Because substitution creates fresh variables names, there are multiple (in fact, infinite) correct values for the result of any step of reduction which includes bound variables. You are not required to keep *any* bound variable names, even if they are unambiguous. You may choose to use de Bruijn indices for all reduction, but you may also rename any bound variable if you wish to. As such, testing will require `toDeBruijn`, as that's the only way of assuring that all names are consistent (in that there aren't any names). If your `toDeBruijn` method doesn't work, your grade on `aoe`, `nor`, and `eta` may suffer, simply because we cannot reach a baseline for testing; assignments are hand graded in this course, but we still need the ability to run tests. Make sure `toDeBruijn` works!

Rights