

CS442

Module 3: Formal Semantics

University of Waterloo

Winter 2022

In the last module, we described the syntax for a simplistic programming language, the λ -calculus. We also described how to compute using the λ -calculus, and the problems that arise depending on how you choose to compute (in particular, the reduction order). While we defined the *form* of λ -calculus quite formally, the *behavior* we described informally.

One of the major endeavors of modern programming languages research is in formalizing our understanding of how language constructs behave, both on their own and in interaction with each other. It is not sufficient to formalize the syntax (shape, form) of programs; we are interested in formalizing the *meanings* of the various elements of a programming language, and ultimately the language itself. This discipline is called *formal semantics*. In studying formal semantics, our goal is to formulate a model capable of precisely describing the behavior of every program in a given language. Such a model provides us tools to prove program correctness, program termination, or other critical properties. Furthermore, we can also use such a model to prove certain properties of the language itself, such as that it's powerful enough to express universal programs (is equivalent to a Turing machine), or to show equivalence of programs in different programming languages. The knowledge gained could even help build compilers and interpreters to produce more efficient implementations of the language, by proving that certain transformations do not affect the behavior.

In Module 2, we described the syntax of the λ -calculus formally, and described substitution formally, but described the behavior informally. In Module 2, we described, for instance, the AOE strategy in plain English, as “always choosing the leftmost, innermost redex that is not in an abstraction”. A prosaic definition like this will usually not suffice. We would like to have a small set of (usually) syntax-directed rules that describe the behavior of the elements of a language's syntax in a formal, mathematical setting.

A semantic model usually comes with a set of *observables*, which describes the valid outputs of the model. Such outputs could be the produced value by following a number of rules, the set of all types of the language, or simply whether a program returns an error or not. In each case, we would choose an appropriate set of observables, and then build a semantic model to match.

In this course's formal semantics, we are primarily going to study the *operational semantics* of various programming languages. Operational semantics is a system for specifying how a program executes and possibly how to extract a result from it. More specifically, as our main goal for the course is to understand how programs interact with data and code in various programming paradigms, we are more concerned at the *small-step* operational semantics of such paradigms. A small-step operational semantics builds an imaginary “machine” and succinctly describes how this machine might take individual steps, rather than describing the entire computation in one step.

As mentioned, we have already introduced the operational semantics of λ -calculus in Module 2 in an informal way. The goal of this module is to revisit λ -calculus with formal, small-step, operational semantics, demonstrate that it can be used to prove some properties, and show ways of extending it with added primitives.

1 Semantics and Category Theory

We will be describing the reduction steps in our programming languages by formally describing an arrow (\rightarrow) operator, which maps a program state to the “next” program state. This is described within the context of *category theory*, in which our language is a *category*, and \rightarrow is a *morphism* over that category. You are not expected to have seen these terms before, so we will briefly introduce category theory here. This course will not look deeply into category theory, but, since programming language semantics are described as categories in category theory, knowing some of the language from category theory will help to contextualize formal semantics.

In CS courses, you have undoubtedly seen *sets* and *set theory*. *Group theory* extends set theory by describing groups, which are sets that correspond to, and are described by, certain *axioms* (defined for given groups), and generalizes the language of functions between and within groups. It is from group theory that words such as *isomorphic* and *homomorphic* arise, to describe certain properties of these functions. Category theory abstracts beyond this by describing categories which may not obviously be describable as groups or sets; in particular, one can describe entire mathematical calculi as categories. For instance, one can describe set theory *itself* as a category, with expressions in set theory as the *objects* described, and the functions being equivalences (or reductions, or expansions, etc) between them; for instance, the resolution of the expression $\{1, 2\} \cup \{3, 4\}$ to $\{1, 2, 3, 4\}$ is a function, probably a function that more generally describes the resolution of all expressions of the form $X \cup Y$. We call these functions between objects *morphisms*.

In set theory, we don’t have any computation, and so there’s no ordering, so any morphism that maps $\{1, 2\} \cup \{3, 4\}$ to $\{1, 2, 3, 4\}$ should be reversible. These two expressions are equivalent, and the morphism shouldn’t “prefer” one to the other. In computing, even in languages that we don’t explicitly specify an order of computation, any given expression is only computed in one direction (e.g., we never replace 4 with $2 + 2$), and so our morphisms will describe steps of computation.

Where category theory becomes particularly relevant is its abstraction over itself. In the language of category theory, we can describe an entire category—which, recall, can be a mathematical calculi, a language—as an object within the category of categories, and describe a morphism mapping that category to another category. For instance, it is possible to reversibly map the language of sets to the language of predicate logic. By doing so, whole bodies of mathematical literature and proofs can be mapped into other contexts, allowing for a sort of generalization of proofs that was not possible before category theory. Mappings between categories like this are called *functors*, but they’re really just morphisms given a funny name because mathematicians aren’t as accustomed to this kind of abstraction as we are.

Aside: We introduced categories as “not obviously be describable as groups or sets”. In fact, since category theory describes how categories can be mapped to other categories, and category theory is itself a category, it is perfectly possible to map any category to *some* kind of set: for instance, we describe the set of all valid program states. Hence “not obviously”, rather than “not”.

We describe our own languages in terms of a morphism, which maps program states to program states. At this point we will describe program states in purely the same syntax as the language itself, but in future modules, we will add extra syntax for additional state; in either case, the syntax is our calculus. Morphisms are usually shown as arrows, often with text to specify exactly which morphism is being described; in fact, we’ve already seen a few morphisms, such as \rightarrow_β , but didn’t call them such at the time.

When describing morphisms, we are free to use other categories to do so. For instance, we could say that $(x + y \rightarrow z)$ if $(x + y = z)$, and we are now describing our language in terms of the language of arithmetic. It’s important to be clear, in such cases, what language is being described and what language is being used; for instance, in this case, it’s important to realize that the first $+$ was part of the syntax of *our* language, and the second $+$ was part of the syntax of *arithmetic*. Usually, this sort of mapping is too narrow to be called a functor—we haven’t actually described a complete rewriting of our language in terms of arithmetic, merely a step within our language—but in some cases, languages are actually described in terms of functors, by describing how to rewrite one language into another language. In fact, that’s a compiler, and proving things about compiler correctness involves proving the functor correct.

We won’t go any more deeply than this into category theory, because we’re not usually proving more broad things about categories. We’re only narrowly interested in proving things about our particular languages. But, you

should now have some idea of the formal underpinnings we're using to describe languages: programming language semantics are not an ad hoc invention, they are described in the language of categories.

2 Semantics and Reality

Our goal is to formally model programming languages, to prove things about them mathematically. But, programming languages aren't mathematical abstractions, they're practical tools for solving real problems. Is there anything to guarantee that the semantics we formally model are the same as the semantics implemented in real programming language implementations? The short answer is "usually not".

Our formal semantics are a mathematical abstraction, not software. There are systems that make formal semantics executable, but the resulting interpreters are usually unusably slow. The purpose of these systems is to have a ground truth for writing test cases, not to use as practical language implementations. Even that is imperfect, however, since it's always possible to write formal semantics which are consistent, but not what you intended; an inconsistency between a real implementation and a formal semantics can be a bug in the implementation, but it can also be a bug in the semantics.

There are also aspects of real implementations which are usually intentionally ignored in formal semantics. For instance, we won't discuss what happens when the program state is too large to hold in memory. And, in later modules, we won't discuss garbage collection, even though it's crucial to a correct implementation of many systems.

In a much later module on systems programming, we will discuss one counterexample, which successfully uses a formally-defined version of C both as a formal semantics and as a real compiler.

In reality, it's impossible to prove, in the mathematical sense of the word, anything about how a program will behave on a real system. Aside from the pitfalls mentioned above, no formal system can model "a disgruntled employee took a pickaxe to my server". Since we're proving things about abstract calculi, rather than a real implementations, we can actually *prove* things, with all the rigor of mathematics. But, since we're not proving things about a real implementation, it is the job of the designer of a formal semantics to argue that the semantics correctly reflects the design of the language, and/or of some implementation of the language.

Put differently, don't get too dependent on formal semantics and formal proof. Although we will describe languages formally, a language without a good implementation is useless, and the connection between the formal system and the implementation will always be tenuous.

3 Review: Post System

If you are already keen on theoretical computer science (or just still remember the Post system introduced in CS 245), congratulations, you may skip this section. Otherwise, please read along.

The *Post system*, named after Emil Post, is an example of a deductive formal system, which can be used to reason about programming languages. There are three components to a Post system: a set of *signs* (which forms the *alphabet* of the system), a set of *variables*, and a set of *productions*. A term is a string of signs and variables, and a production is an expression of the form:

$$\frac{t_1 \quad t_2 \quad t_3 \quad \cdots \quad t_n}{t}$$

Where $t, t_1, \dots, t_n (n \geq 0)$ are all terms. The t_i are called the *premises* of the production, and the t is the *conclusion*. Thus, a production with the form

$$\frac{\text{premises}}{\text{conclusion}}$$

is read as “if premises are true, then the conclusion holds”. A production without premises is permitted, and simply means that the conclusion is true by definition. Productions without premises are called *axioms*.

Productions are the definitions within our system, so it is outside the scope of the Post system to prove that the productions themselves are correct. In our case, the conclusions will be our \rightarrow morphism, and thus the definition of how our programming languages are evaluated; in essence, each conclusion is a step we can take, and the premises are the context in which we can take those steps.

Post systems are used to prove conclusions, where a proof is constructed from proofs of its premises. Proofs based on Post systems are constructed using the following rules:

1. An instance of an axiom is a proof of its conclusion;
2. If P_1, P_2, \dots, P_n are proofs of t_1, t_2, \dots, t_n respectively, and

$$\frac{t_1 \quad t_2 \quad t_3 \quad \dots \quad t_n}{t}$$

is an instance of a production, then

$$\frac{P_1 \quad P_2 \quad P_3 \quad \dots \quad P_n}{t}$$

is a proof of t .

Thus, given a final conclusion, a proof of that conclusion can be formed by proving its premises, until no unproven premises remain. The result of such a proof is an upside down tree with the root (final conclusion) at the bottom and the leaves (axioms) at the top.

Example 1. As an example of a Post system, we can encode the logical operations of ‘and’, \wedge , and ‘or’, \vee , using the following three rules:

$$\frac{A}{A \vee B}$$

$$\frac{B}{A \vee B}$$

$$\frac{A \quad B}{A \wedge B}$$

Using this small system, it is possible to show that the proof of $(A \vee B) \wedge (A \vee C)$ follows from a proof of A alone:

$$\frac{\frac{A}{A \vee B} \quad \frac{A}{A \vee C}}{(A \vee B) \wedge (A \vee C)}$$

Post systems are used extensively for describing formal semantics. You will see that formal semantics of programming languages, including type systems, are often described in Post systems.

4 Operational Semantics for (Plain) λ -calculus

We have already discussed the semantics of λ -terms in Module 2, when we discussed free and bound variables, substitution, α -conversion, and β -reduction. Assuming that we have already established the notion of binding, substitution, and α -conversion, β -reduction seems to be a suitable candidate for operational semantics, for it specifies a procedure for carrying out computation.

Let’s rewrite β -reduction as a formal set of rules. First of all, all expressions that have β -redex in the outermost level can be directly reduced, with no premises:

$$\overline{(\lambda x. M)N \rightarrow_{\beta} M[N/x]}$$

This rule corresponds in the first part of the definition. However, the next part of the definition¹, which describes the reduction of β -redices within an expression, cannot be simply interpreted using a single rule. We have to rely on the structure of the λ -expressions. Recall that λ -expressions are either abstractions, applications, or variables. A variable itself certainly doesn't need any rules for β -reduction, as a variable can't directly be reduced, but we can have reductions happening inside abstractions and applications. The above rule is the special case where the rator is an abstraction.

We still need to take the case where there is reduction happening inside an abstraction, or within the rator or the rand of an application. The following rule captures the first of these cases, reduction within an abstraction:

$$\frac{M \rightarrow_{\beta} P}{\lambda x. M \rightarrow_{\beta} \lambda x. P}$$

According to this rule, in order to show that $\lambda x. M \rightarrow_{\beta} \lambda x. P$, we must either provide a proof of $M \rightarrow_{\beta} P$, or there must exist a rule that states $M \rightarrow_{\beta} P$ is an axiom.

For applications, remember that in the original description of β -reduction, we didn't specify a reduction order. That is, we can choose to start our reduction either in the rator and the rand. For those two cases, we need separate rules:

$$\frac{M \rightarrow_{\beta} P}{MN \rightarrow_{\beta} PN} \qquad \frac{N \rightarrow_{\beta} P}{MN \rightarrow_{\beta} MP}$$

We have just described how computation proceeds in λ -calculus. However, because a λ -calculus expression may match more than one of these conditions, our description is non-deterministic; we haven't describe a *particular* way of computing, but *all valid* ways of computing. In the previous module, we made this deterministic by focusing on the selection of redices, and we will now do the same formally.

5 Defining Evaluation Order

As we mentioned earlier, evaluation order is in fact very important, since most programming languages have a deterministic order of execution, and even those that don't have at least *some* restrictions on the order of execution. If we were to model actual programming languages using our calculus, it is crucial to choose a reduction strategy. In this section, we are going to discuss the operational semantics of λ -calculus under Normal Order Reduction and Applicative Order Evaluation. Let's first consider NOR.

Definition 1. (Small-Step Operational Semantics of the Untyped λ -Calculus, NOR)

Let the metavariable M range over λ -expressions. Then a semantics of λ -terms in NOR is given by the following rules:

$$\begin{array}{c} \frac{}{(\lambda x. M_1)M_2 \rightarrow M_1[M_2/x]} \qquad \frac{M \rightarrow M'}{\lambda x. M \rightarrow \lambda x. M'} \\[10pt] \frac{M_1 \rightarrow M'_1 \quad \forall x. \forall M_3. M_1 \neq \lambda x. M_3}{M_1 M_2 \rightarrow M'_1 M_2} \qquad \frac{M_2 \rightarrow M'_2 \quad \forall M'_1. M_1 \not\rightarrow M'_1 \quad \forall x. \forall M_3. M_1 \neq \lambda x. M_3}{M_1 M_2 \rightarrow M_1 M'_2} \end{array}$$

The first and the second rule stayed the same as in β -reduction. Similar to non-deterministic β -reduction, we can reduce an expression that is either a redex, or an abstraction which contains a redex. However, in order to enforce NOR, we have to add additional restrictions to the third and fourth reduction rules. First of all, if M_1 can be reduced further, we should reduce M_1 instead; this is the reason we introduced the premise $\forall M'_1. M_1 \not\rightarrow M'_1$. We also explicitly state that the third and fourth rule do not apply if the first rule would have applied, with the premise $\forall x. \forall M_3. M_1 \neq \lambda x. M_3$ (that is if M_1 is an abstraction).

¹Recall from Module 2: $C[(\lambda x. M)N] \rightarrow_{\beta} C[M[N/x]]$

Now let's look at AOE:

Definition 2. (Small-Step Operational Semantics of the Untyped λ -Calculus, AOE)

Let the metavariable M range over λ -expressions. Then a semantics of the λ -terms in AOE is given by the following rules:

$$\frac{\forall M'_2. M_2 \not\rightarrow M'_2}{(\lambda x. M_1) M_2 \rightarrow M_1[M_2/x]} \qquad \frac{M_1 \rightarrow M'_1}{M_1 M_2 \rightarrow M'_1 M_2}$$

$$\frac{M_2 \rightarrow M'_2 \quad \forall M'_1. M_1 \not\rightarrow M'_1}{M_1 M_2 \rightarrow M_1 M'_2}$$

For AOE, the first rule has the added condition that the rand (i.e. the argument) can be applied only if it can't be reduced further. Also, the abstraction rule is removed, since we can not reduce within an abstraction; similarly, in most programming languages, you can not evaluate inside a function you didn't yet call. The last rule has the premise $\forall x. M_1 \neq \lambda x. M_2$ removed since again we want the argument to be fully reduced before substituting itself into an abstraction first.

6 Terminal Values

In the previous section, we used the language of predicate logic (specifically, for-all) to conditionalize productions. While this is mathematically valid, it complicates the description of the language, and makes it more difficult to prove that a particular production is the right one to use: to show that we can use the first rule of AOE, for instance, we need to demonstrate that the rand cannot be reduced. Formally proving this involves proving that it does *not* match any existing rule, which can in turn require proving more negatives. But, the point of this for-all was simply “use this reduction if another reduction is not possible”. The rules and conditions become much clearer and easier to prove if we can specify when an expression cannot be reduced any further *syntactically*, i.e., based only on the form of the expression. We call such irreducible expressions *terminal*, or *final*.

There are a few choices that we could make for possible sets of terminal values in λ -calculus: we could choose β -normal form, weak normal form, or even head normal form (only the leftmost expression is required to be in normal form). If we use anything other than β -normal form, we are losing the guarantee given by the Church-Rosser Theorem. Even if we use β -normal form as the set of terminal value, we still need to be able to answer some important questions. For example, what is the semantics of $(\lambda x. xx)(\lambda x. xx)$? The only response we can give is “no semantics”, since it does not have a normal form. And what about $(\lambda x. \lambda y. y)((\lambda x. xx)(\lambda x. xx))$? It has a terminal value, but not all possible legal derivations will lead to it. Should this expression be given a final value of $\lambda y. y$ since there is a possible reduction to it, or we should say that there is no meaning? In fact, the answer depends on what one needs to achieve by designing the semantics.

For now, we focus on the steps themselves rather than the possible terminal values. As a result, we can just let our final values be “the set of values our operational semantics would produce”. Terminal values become more important when we introduce types, so in the next module, we will discuss terminal values in greater detail.

7 Showcase: A Simple Proof

In this section, we will show that $(\lambda x. \lambda y. y)((\lambda x. xx)(\lambda x. xx))x$ indeed terminates and evaluates to x under NOR. This is not a formal proof by any means; however, we will use this example to give you an idea to how programming language theorists work with semantics.

The formal way to specify that the former expression terminates and evaluates to the latter is going to look like this:

$$(\lambda x. \lambda y. y)((\lambda x. xx)(\lambda x. xx))x \rightarrow (\lambda y. y)x \rightarrow x$$

This example is quite short. However, what should we do if we are dealing with larger examples? The answer is that we need a \rightarrow^* operator. It might be useful to formally define the \rightarrow^* operator so we can show every single step at once, instead of splitting them into separate proofs:

Definition 3. (Sequencing) Let the metavariables M range over λ -expressions and \rightarrow be the operator of “one step” in any small-step operational semantics. Then \rightarrow^* is defined as so:

$$\frac{}{M \rightarrow^* M} \qquad \frac{M_1 \rightarrow^* M_2 \quad M_2 \rightarrow M_3}{M_1 \rightarrow^* M_3}$$

Aside: \rightarrow^* is the *reflexive and transitive closure* of \rightarrow .

To keep the proof text short, we will make the following definitions:

$$A = (\lambda x. \lambda y. y), B = (\lambda x. xx)(\lambda x. xx)$$

Now we can actually start our “proof”.

$$\frac{\frac{ABx \rightarrow^* ABx \quad ABx \rightarrow (\lambda y. y)x}{ABx \rightarrow^* (\lambda y. y)x} \quad (\lambda y. y)x \rightarrow x}{(\lambda x. \lambda y. y)((\lambda x. xx)(\lambda x. xx))x \rightarrow^* x}$$

Although this proof is of course trivial, with proper abstraction, we can prove similar properties of entire classes of programs. We will look at some of those properties in the next module.

Aside: There are also many, many other kinds of semantics. In this aside, we showcase two of them since they are also used in the field of programming languages. One of the variations of operational semantics is *big-step operational semantics*, which describes the terminal values every expression will evaluate to directly, rather than as the transitive closure of smaller steps. For example, this is the big-step operational semantics for the λ -calculus under AOE:

$$\frac{}{V \Downarrow V} \qquad \frac{M[V_1/x] \Downarrow V_2}{(\lambda x. M)V_1 \Downarrow V_2}$$

$$\frac{M_1 \Downarrow V_1 \quad V_1 M_2 \Downarrow V_2}{M_1 M_2 \Downarrow V_2} \qquad \frac{M_1 \Downarrow V_1 \quad V_2 V_1 \Downarrow V_3}{V_2 M_1 \Downarrow V_3}$$

In this example, V is the metavariable over values.

Another kind of semantics is *denotational semantics*. Denotational semantics are used to show the correspondence from language constructs to familiar mathematical objects. Our definition of functional language constructs in terms of λ -expressions is an example of a denotational semantics, where the map between expressions and observables (i.e. the λ -expressions) is generally denoted by double squared brackets. A denotational semantics is formally a functor, while an operational semantics is a morphism over the language described.

8 Adding Primitives

At the end of Module 2, we discussed the λ -calculus implementations of commonly seen data types. While those discussions are very useful in showcasing the power of λ -calculus in representing computation, the implementations presented are not particularly practical; furthermore, it is much more efficient to make use of the computer architectures we have and implement those features in their terms. For instance, since all computer architectures support integers (of some limited range) natively, it would be absurd to implement integers as Church numerals in a real language. As a result, in the practice of modeling real programming languages, we tend to model those as *primitives*. To be specific, those data types will be treated as intrinsic (i.e. built-in) values of our language. In the λ -calculus, the only real values were functions; we now extend the λ -calculus by adding other kinds of values, and the other functionality necessary to make use of them. Most formal semantics for programming languages have some kind of primitives, so in the remainder of this module, we will describe the semantic rules required to add common primitives. These rules will be referenced or reused in future modules.

8.1 Booleans and Conditionals

We will first introduce the syntactic elements, in Backus Normal Form (BNF). Note that we are adding new kinds of expressions in the definition of $\langle Expr \rangle$; We will use “...” to denote the part of the definitions of expressions that was defined in Module 2.

$$\begin{aligned}\langle Boolexp \rangle &::= \text{true} \mid \text{false} \\ &\mid \text{not } \langle Expr \rangle \\ &\mid \text{and } \langle Expr \rangle \langle Expr \rangle \\ &\mid \text{or } \langle Expr \rangle \langle Expr \rangle \\ \langle Expr \rangle &::= \dots \\ &\mid \langle Boolexp \rangle \\ &\mid \text{if } \langle Expr \rangle \text{ then } \langle Expr \rangle \text{ else } \langle Expr \rangle\end{aligned}$$

These syntactic elements are very similar to the ones you have seen from Module 2. However, they are now actually part of the syntax of our extended language; that is, rather than describing how to convert them to λ -calculus, we extend λ -calculus with a new type of expression. Programs in the λ -calculus with Boolean primitives are simply λ -calculus expressions with additional syntax for Boolean expressions, like so:

$\lambda x. \lambda y. \text{ if } x \text{ then } y \text{ else } false$

Aside: Our grammar is actually ambiguous, because in “and” and “or”, there’s no clean delineation between the two subexpressions, and nested “if”s can create the same ambiguity with “else”. This problem is actually common in real languages though, with almost all languages in the C family having an ambiguity with “if”. We’ll just sidestep this problem by only writing code that’s unambiguous; this isn’t a parsing course!

We will now describe the operational semantics for this new language. Let the metavariable E range over all expressions. We will start with “not”:

$$\frac{}{\text{not true} \rightarrow \text{false}} \qquad \frac{}{\text{not false} \rightarrow \text{true}} \qquad \frac{E \rightarrow E'}{\text{not } E \rightarrow \text{not } E'}$$

For “and” and “or”, we want the computation of the first parameter to happen first. In addition, like in most programming languages, we would like short-circuiting behavior for them; i.e., the evaluation of the second operand should not proceed if the synthesis is known from the first.

$$\begin{array}{c}
\frac{}{\text{and false } E \rightarrow \text{false}} \qquad \frac{}{\text{and true false} \rightarrow \text{false}} \qquad \frac{}{\text{and true true} \rightarrow \text{true}} \\
\frac{E_1 \rightarrow E'_1}{\text{and } E_1 \ E_2 \rightarrow \text{and } E'_1 \ E_2} \qquad \frac{E \rightarrow E'}{\text{and true } E \rightarrow \text{and true } E'}
\end{array}$$

The last two rules are how we describe “the first argument must be fully evaluated before the second one”. The first rule describes the short-circuiting behavior: whether the second argument is evaluated or not, as long as the first argument evaluates to “false”, the whole “and” evaluates to false.

Exercise 1. Write the semantic rules for “or”.

Now we add the rules for if statements:

$$\begin{array}{c}
\frac{}{\text{if true then } E_1 \text{ else } E_2 \rightarrow E_1} \qquad \frac{}{\text{if false then } E_1 \text{ else } E_2 \rightarrow E_2} \\
\frac{E_1 \rightarrow E'_1}{\text{if } E_1 \text{ then } E_2 \text{ else } E_3 \rightarrow \text{if } E'_1 \text{ then } E_2 \text{ else } E_3}
\end{array}$$

Assuming you’ve followed the exercise to write the semantics for “or”, our semantics are now complete, but you may have noticed that they have a problem: there are programs that are syntactically correct but don’t match any of our rules! For instance, if $\lambda x. x \text{ true false}$. Actually, we had a similar case before—you can’t reduce a variable—but now the problem arises for another reason: in order to reduce our new expressions, certain subexpressions need to reduce to either true or false. For the time being, we can’t solve this problem; we’ll come back to it in the next module.

8.2 Numbers

Note that we will restrict our definition to natural numbers. Also, we are working in an imaginary machine, so we don’t care about overflows (i.e., we assume that we can represent numbers of an infinite range). We will make the following definitions in our syntax, again in BNF:

$$\begin{aligned}
\langle \text{Num} \rangle &::= 0 \mid 1 \mid \dots \\
&\quad \mid \langle \text{NumBinOps} \rangle \langle \text{Expr} \rangle \langle \text{Expr} \rangle \\
\langle \text{NumBinOps} \rangle &::= + \mid - \mid * \mid / \\
\langle \text{Expr} \rangle &::= \dots \\
&\quad \mid \langle \text{Num} \rangle
\end{aligned}$$

We will now consider the semantics of binary operations. Let M, N range over expressions, a, b range over natural numbers. We will start with addition:

$$\begin{array}{c}
\frac{a + b = c}{(+ \ a \ b) \rightarrow c} \qquad \frac{M \rightarrow M'}{(+ \ M \ N) \rightarrow (+ \ M' \ N)} \qquad \frac{M \rightarrow M'}{(+ \ a \ M) \rightarrow (+ \ a \ M')}
\end{array}$$

Note that this set of rules forces the first argument (i.e., the left-hand side) to be evaluated before the second argument is evaluated. Also note that we’re describing our language, the λ -calculus with numbers, in terms of the language of arithmetic, with the predicate $a + b = c$.

Aside: Why are we using M and N here and E elsewhere? Partially it's because we expect these particular expressions to evaluate to numbers, and M and N are common for that, but mostly it's just to remind you that none of this is set in stone. Make sure you know how any work you're reading is using its variables!

Let's now look at subtraction.

$$\frac{a - b = c \quad c \in \mathbb{N}}{(- a b) \rightarrow c} \qquad \frac{M \rightarrow M'}{(- M N) \rightarrow (- M' N)} \qquad \frac{M \rightarrow M'}{(- a M) \rightarrow (- a M')}$$

The semantics for subtraction is almost the same as with addition, but there is one difference: to actually compute $a - b$, we need to make sure that $a - b$ is a natural number. With this semantics, there is no rule to match expressions like $(- 2 3)$, and so such expressions cannot be reduced. We describe this phenomenon as “getting stuck”, and in the next module, we will dive into this issue and discuss the significance of an expression getting stuck. Another way of handling this is to actually allow such subtraction, but define the result as something arbitrary and perhaps counter-intuitive, such as 0:

$$\frac{a - b = c \quad c \notin \mathbb{N}}{(- a b) \rightarrow 0}$$

Although this definition is unintuitive, it is not incorrect. It's not correct, either; it's a definition. We are defining how our language behaves, and if that doesn't match perfectly with the $-$ of arithmetic, then so be it; that is part of the definition. Indeed, in real programming languages with integers of a limited size, no mathematical operations match their arithmetic definitions perfectly, because of overflow, but these languages are still valid and well defined.

This is one of the many places where real implementations tend to diverge from formal semantics. It is rare, albeit not unheard of, for formal semantics to model integer bounds and similar constraints that arise from real hardware.

Exercise 2. Write the semantic rules for $*$ and $/$ (use integer division; think about how to handle zero division.)

Exercise 3. Propose changes to the syntax rules and add new semantic rules, so we have pred and succ , which are unary functions for getting predecessor and successor of a number, in our language. Note: $\text{pred } 0 = 0$.

8.3 Lists

In this section we will discuss lists. We will use the representation you should be familiar with: a list containing 1, 2, 3 will be

$$(\text{cons } 1 (\text{cons } 2 (\text{cons } 3 \text{ empty}))) = [1, 2, 3]$$

We will use a shorthand in mathematics to make our semantic rules compact: $L_1 + L_2$ will be operator to append L_1 to the start of L_2 . For example: $[1] + [2] = [1, 2]$. We will also assume that $+$ works for “empty”, the empty list.

Again, we will list the syntactic elements of lists here:

$$\begin{aligned} \langle \text{ListExpr} \rangle &::= \text{empty} \mid (\text{cons } \langle \text{Expr} \rangle \langle \text{ListExpr} \rangle) \\ &\quad \mid [\langle \text{Expr} \rangle \langle \text{ListRest} \rangle] \\ \langle \text{ListRest} \rangle &::= \varepsilon \mid , \langle \text{Expr} \rangle \langle \text{ListRest} \rangle \\ \langle \text{Expr} \rangle &::= \dots \\ &\quad \mid \langle \text{ListExpr} \rangle \\ &\quad \mid \text{first } \langle \text{ListExpr} \rangle \mid \text{rest } \langle \text{ListExpr} \rangle \end{aligned}$$

The recursive definition of lists is essentially identical to the recursive data definition of the Racket list you saw in first-year courses.

Here are the semantic rules. Let the metavariables L, E range over list expressions and λ -expression respectively:

$$\frac{L_2 = [E] + L_1 \quad \forall E_1. E \not\rightarrow E_1}{(\text{cons } E \ L_1) \rightarrow L_2} \qquad \frac{L_1 = [E] + L_2}{(\text{first } L_1) \rightarrow E} \qquad \frac{L_1 = [E] + L_2}{(\text{rest } L_1) \rightarrow L_2}$$

Note that the premises in the form $L_1 = [E] + L_2$ implies that L_1 is not empty.

Last, don't forget that we want expressions to reduce inside those built-in functions:

$$\frac{E_1 \rightarrow E'_1}{(\text{cons } E_1 \ E_2) \rightarrow (\text{cons } E'_1 \ E_2)} \qquad \frac{E_1 \rightarrow E'_1}{(\text{first } E_1) \rightarrow (\text{first } E'_1)}$$

$$\frac{\forall E_3. E_1 \not\rightarrow E_3 \quad E_2 \rightarrow E'_2}{(\text{cons } E_1 \ E_2) \rightarrow (\text{cons } E_1 \ E'_2)} \qquad \frac{E_1 \rightarrow E'_1}{(\text{rest } E_1) \rightarrow (\text{rest } E'_1)}$$

Note that our definition of lists has been slightly less formal than our previous definitions, as we relied on an informally described mathematical language of lists for our predicates. It is not uncommon for formal semantics to have some quasi-formal “holes” like this, though obviously it is preferable to define everything as precisely as possible.

8.4 Sets

A set is a mathematical collection of distinct objects. In real programming languages, it is usually implemented by a hash-map. However, when formulating a semantics for sets, we usually do not need to worry about their actual implementation; we can just treat it as a mathematical object. So long as the implementation provides the same *observable* behavior, it is correct.

The syntax of sets will be as follows:

$$\begin{aligned} \langle \text{SetExpr} \rangle &::= \text{empty} \mid \{ \langle \text{Expr} \rangle \langle \text{SetRest} \rangle \} \\ &\quad \mid \text{insert } \langle \text{Expr} \rangle \ \langle \text{SetExpr} \rangle \\ &\quad \mid \text{remove } \langle \text{Expr} \rangle \ \langle \text{SetExpr} \rangle \\ \langle \text{SetRest} \rangle &::= \varepsilon \mid , \langle \text{Expr} \rangle \langle \text{SetRest} \rangle \\ \langle \text{Expr} \rangle &::= \dots \\ &\quad \mid \langle \text{SetExpr} \rangle \\ \langle \text{Boolexp} \rangle &::= \dots \\ &\quad \mid \text{contains? } \langle \text{Expr} \rangle \ \langle \text{SetExpr} \rangle \end{aligned}$$

Let metavariables S, E range over be set expressions and all λ -expressions respectively:

$$\frac{\forall E_1. E \not\rightarrow E_1}{(\text{insert } E \ \text{empty}) \rightarrow \{E\}} \qquad \frac{\forall E_1. E \not\rightarrow E_1}{(\text{remove } E \ \text{empty}) \rightarrow \text{empty}} \qquad \frac{\forall E_1. E \not\rightarrow E_1 \quad E \in S}{(\text{contains? } E \ S) \rightarrow \text{true}}$$

$$\frac{\forall E_1. E \not\rightarrow E_1 \quad S' = S \cup \{E\}}{(\text{insert } E \ S) \rightarrow S'} \qquad \frac{\forall E_1. E \not\rightarrow E_1 \quad S' = S \setminus \{E\}}{(\text{remove } E \ S) \rightarrow S'} \qquad \frac{\forall E_1. E \not\rightarrow E_1 \quad E \notin S}{(\text{contains? } E \ S) \rightarrow \text{false}}$$

Exercise 4. Write the semantic rules for set where at least one argument is not fully reduced.

Note that again, we have described our own sets in terms of the language of set theory.

9 Fin

In the next module, we will introduce *types*, which allow us to prove certain properties of languages, including that the semantics do not “get stuck”, by categorizing the kinds of values that may undergo certain operations.

Rights

Copyright © 2020–2023 Yangtian Zi, Gregor Richards, Brad Lushman, and Anthony Cox.
This module is intended for CS442 at University of Waterloo.
Any other use requires permission from the above named copyright holder(s).