

---

# SE 464

## Week 8

— Replication, Intro to Security —

---

# Availability via Replication

The following content is sourced from Computer Systems Design from MIT OCW

<https://ocw.mit.edu/courses/6-033-computer-system-engineering-spring-2018/pages/week-11/>

**goal:** build reliable systems from unreliable components  
the abstraction that makes that easier is

**transactions**, which provide **atomicity** and **isolation**, while not hindering **performance**

**atomicity** → **shadow copies** (simple, poor performance) or **logs** (better performance, a bit more complex)

**isolation** → **two-phase locking**

we also want transaction-based systems to be **distributed** — to run across multiple machines — and to remain **available** even through failures

**C<sub>1</sub>** **write<sub>1</sub>(X)**

**S<sub>1</sub>**

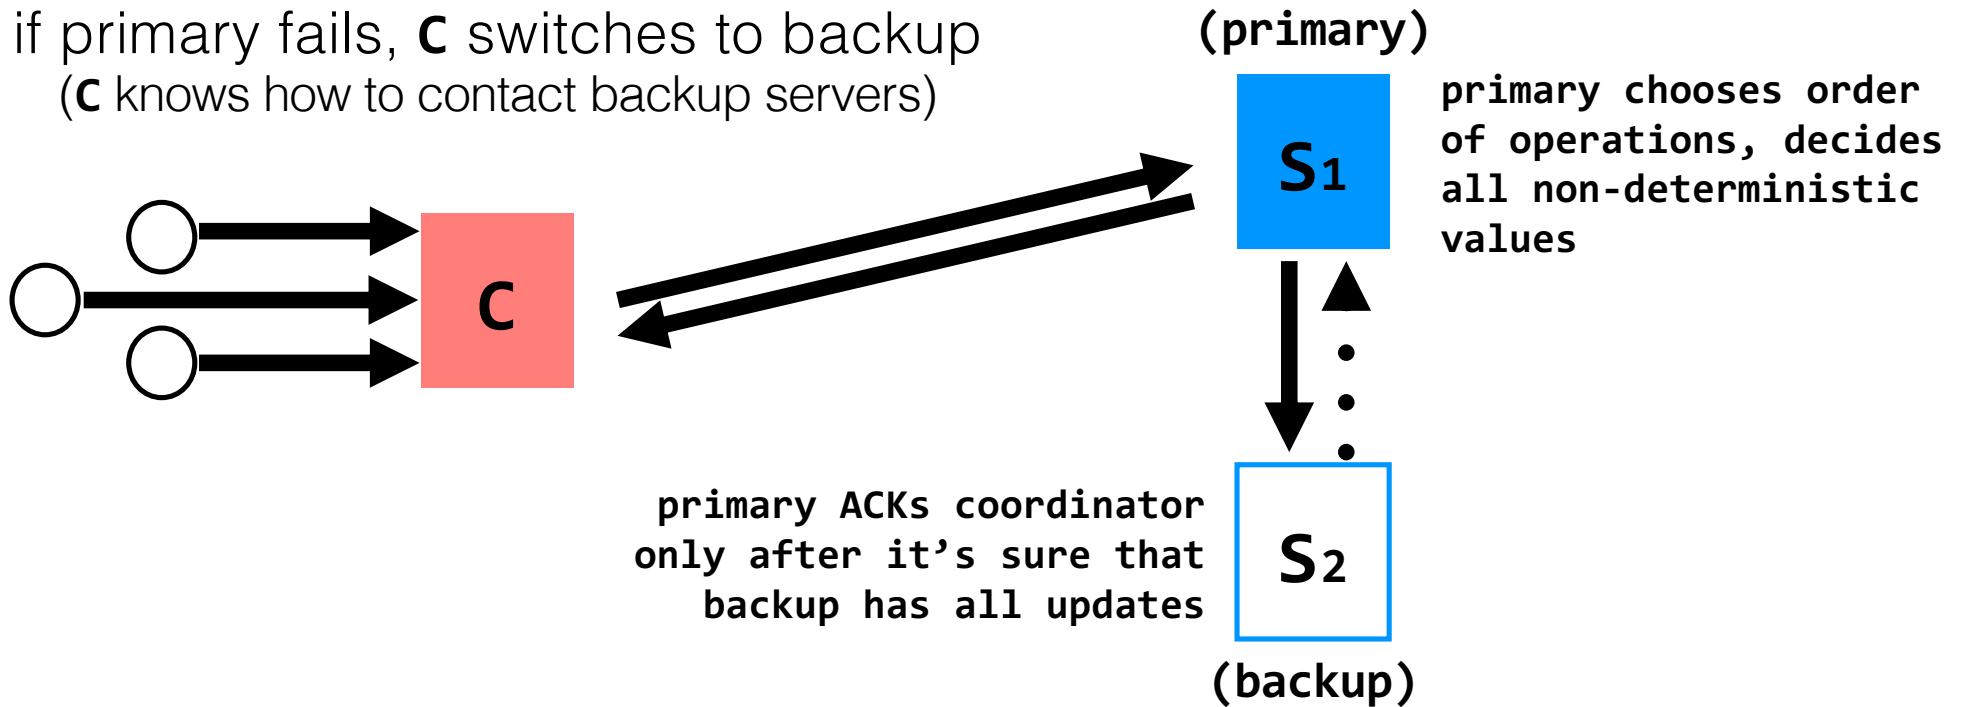
**C<sub>2</sub>** **write<sub>2</sub>(X)**

**S<sub>2</sub>**

(replica of S<sub>1</sub>)


**problem:** replica servers can become inconsistent

if primary fails, **C** switches to backup  
(**C** knows how to contact backup servers)




**attempt:** coordinators communicate with primary servers, who communicate with backup servers

if primary fails, **C** switches to backup  
(**C** knows how to contact backup servers)




**attempt:** coordinators communicate with primary servers, who communicate with backup servers

if primary fails, **C** switches to backup  
(**C** knows how to contact backup servers)



**attempt:** coordinators communicate with primary servers, who communicate with backup servers


# multiple coordinators + the network = problems



**attempt:** coordinators communicate with primary servers, who communicate with backup servers




# multiple coordinators + the network = problems




**attempt:** coordinators communicate with primary servers, who communicate with backup servers

# multiple coordinators + the network = problems




... network partition ...




**attempt:** coordinators communicate with primary servers, who communicate with backup servers

# multiple coordinators + the network = problems




... network partition ...



**C<sub>1</sub>** and **C<sub>2</sub>** are using different primaries;  
**S<sub>1</sub>** and **S<sub>2</sub>** are no longer consistent

**attempt:** coordinators communicate with primary servers, who communicate with backup servers

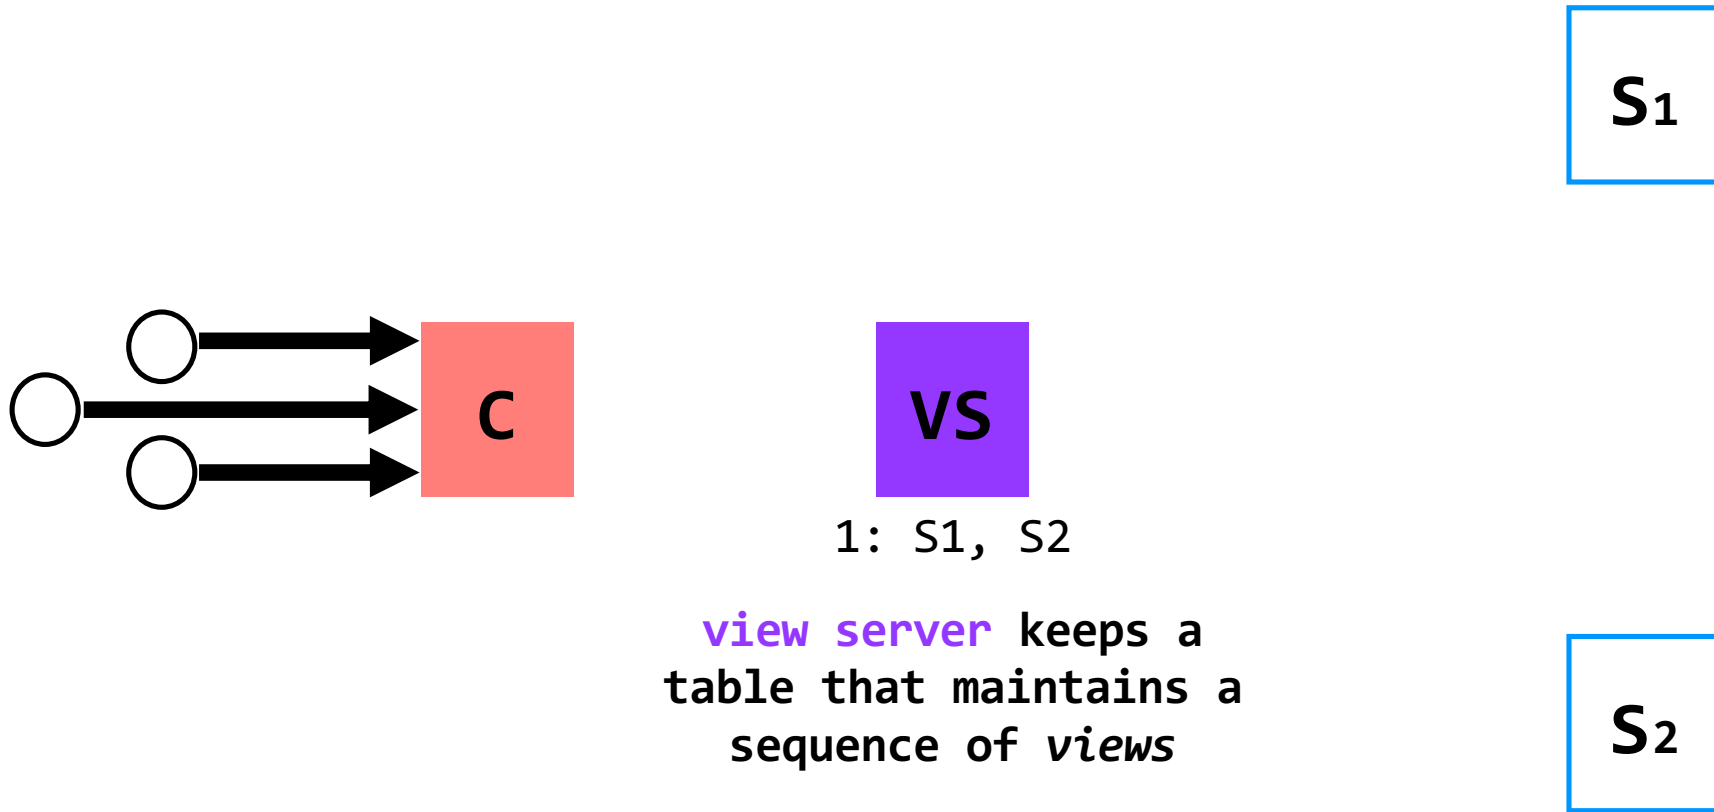


VS

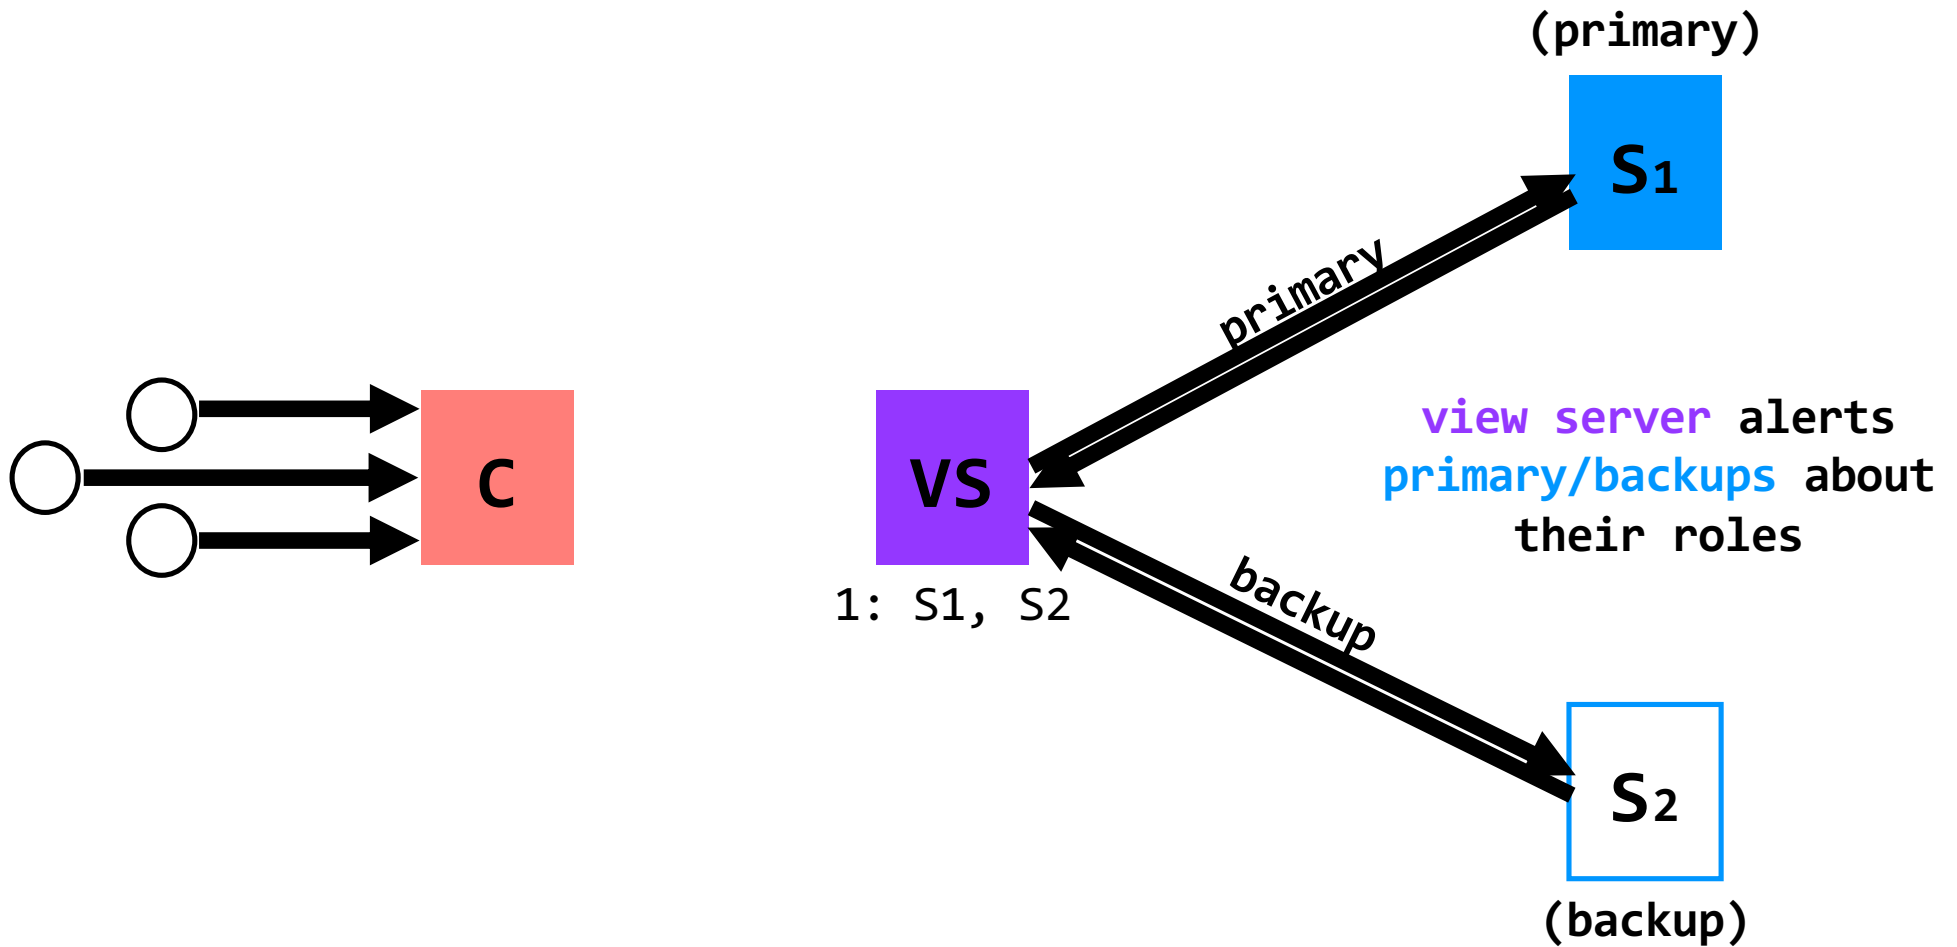
S<sub>1</sub>

S<sub>2</sub>

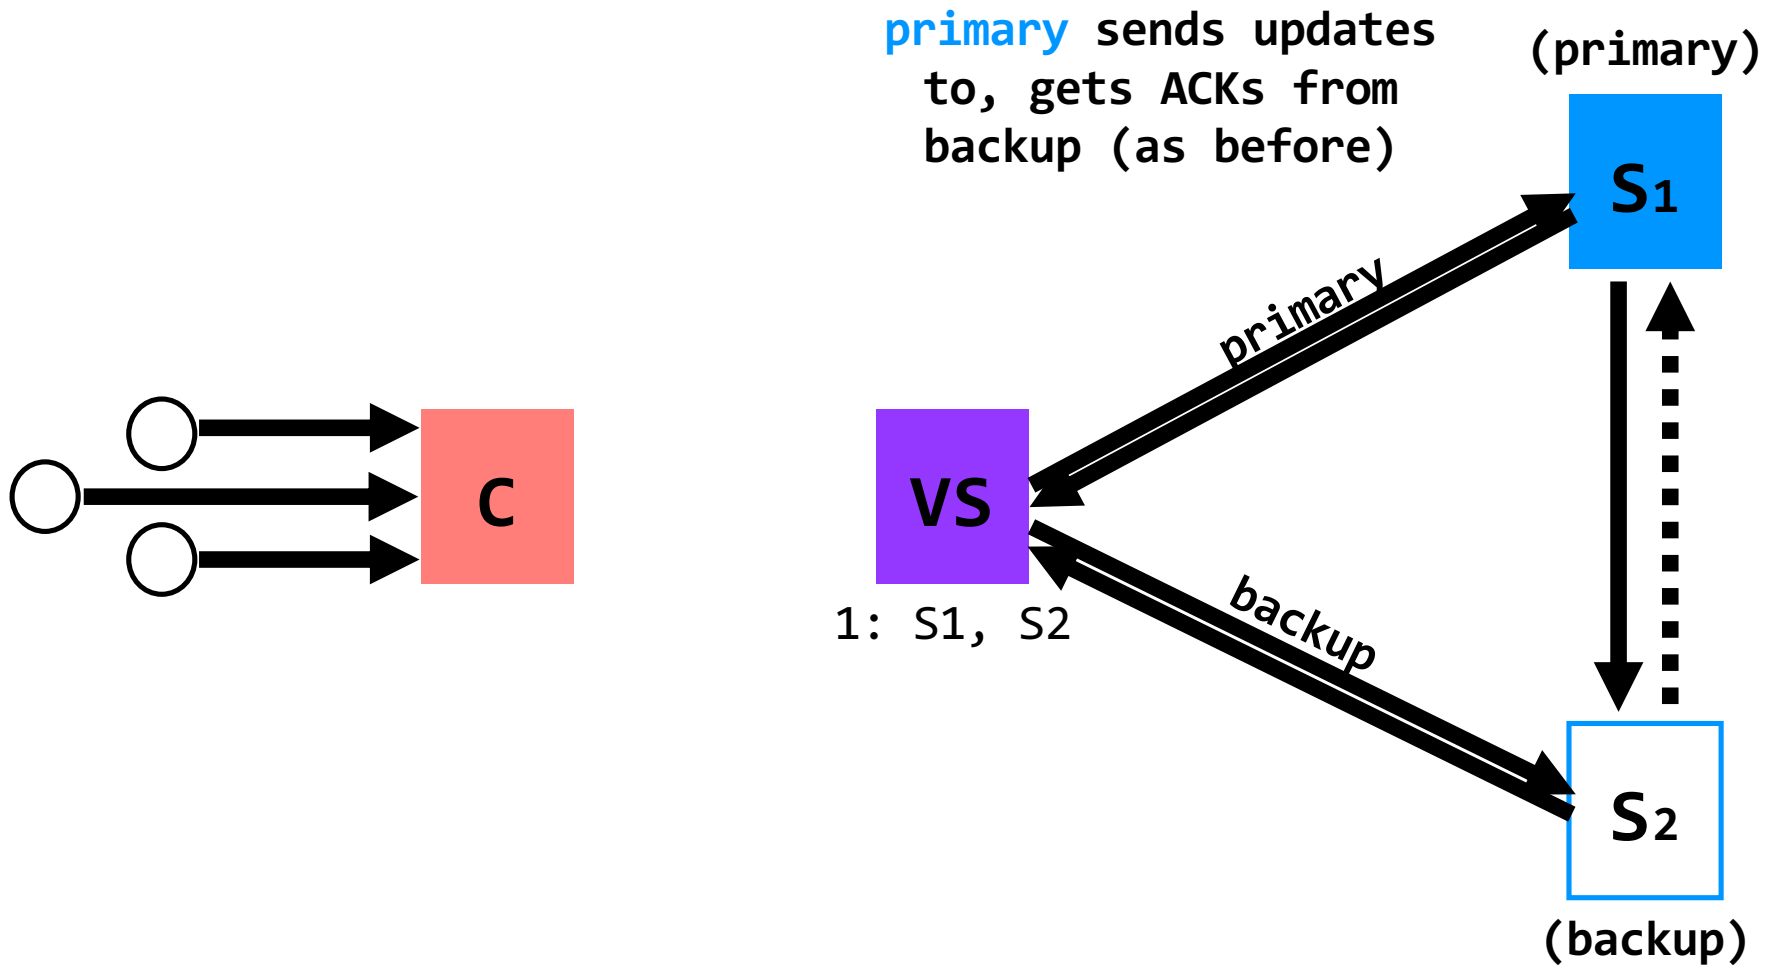
use a **view server**, which determines which replica is the primary



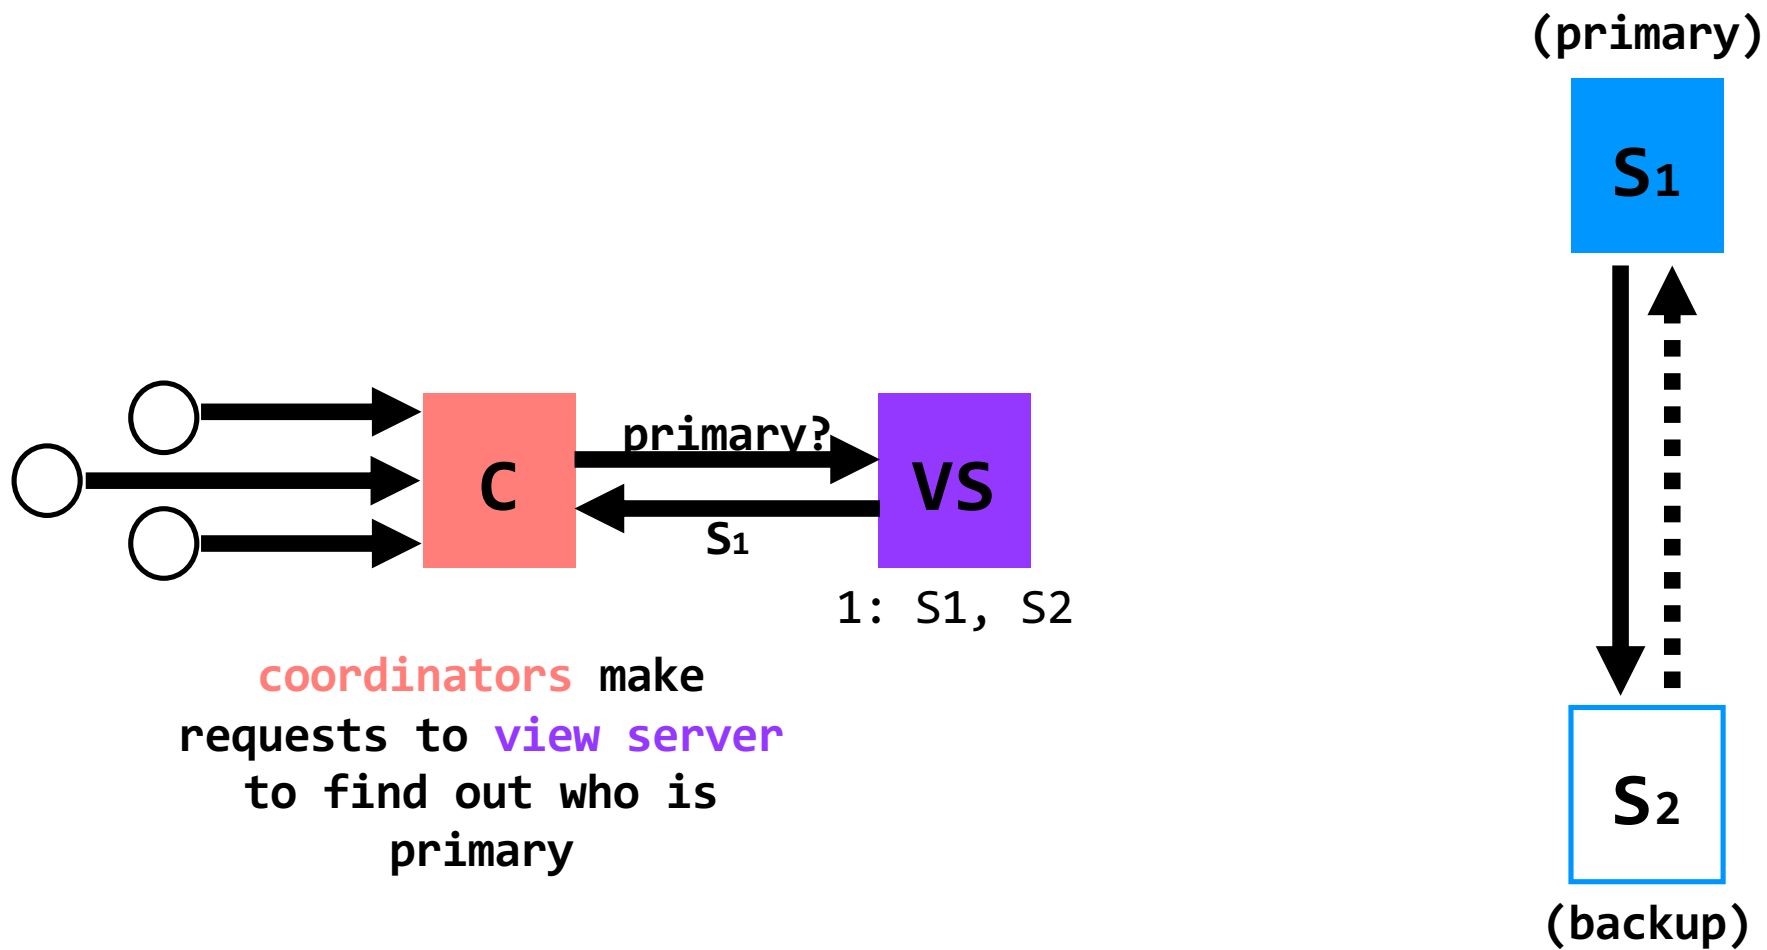
use a **view server**, which determines which replica is the primary



use a **view server**, which determines which replica is the primary




use a **view server**, which determines which replica is the primary



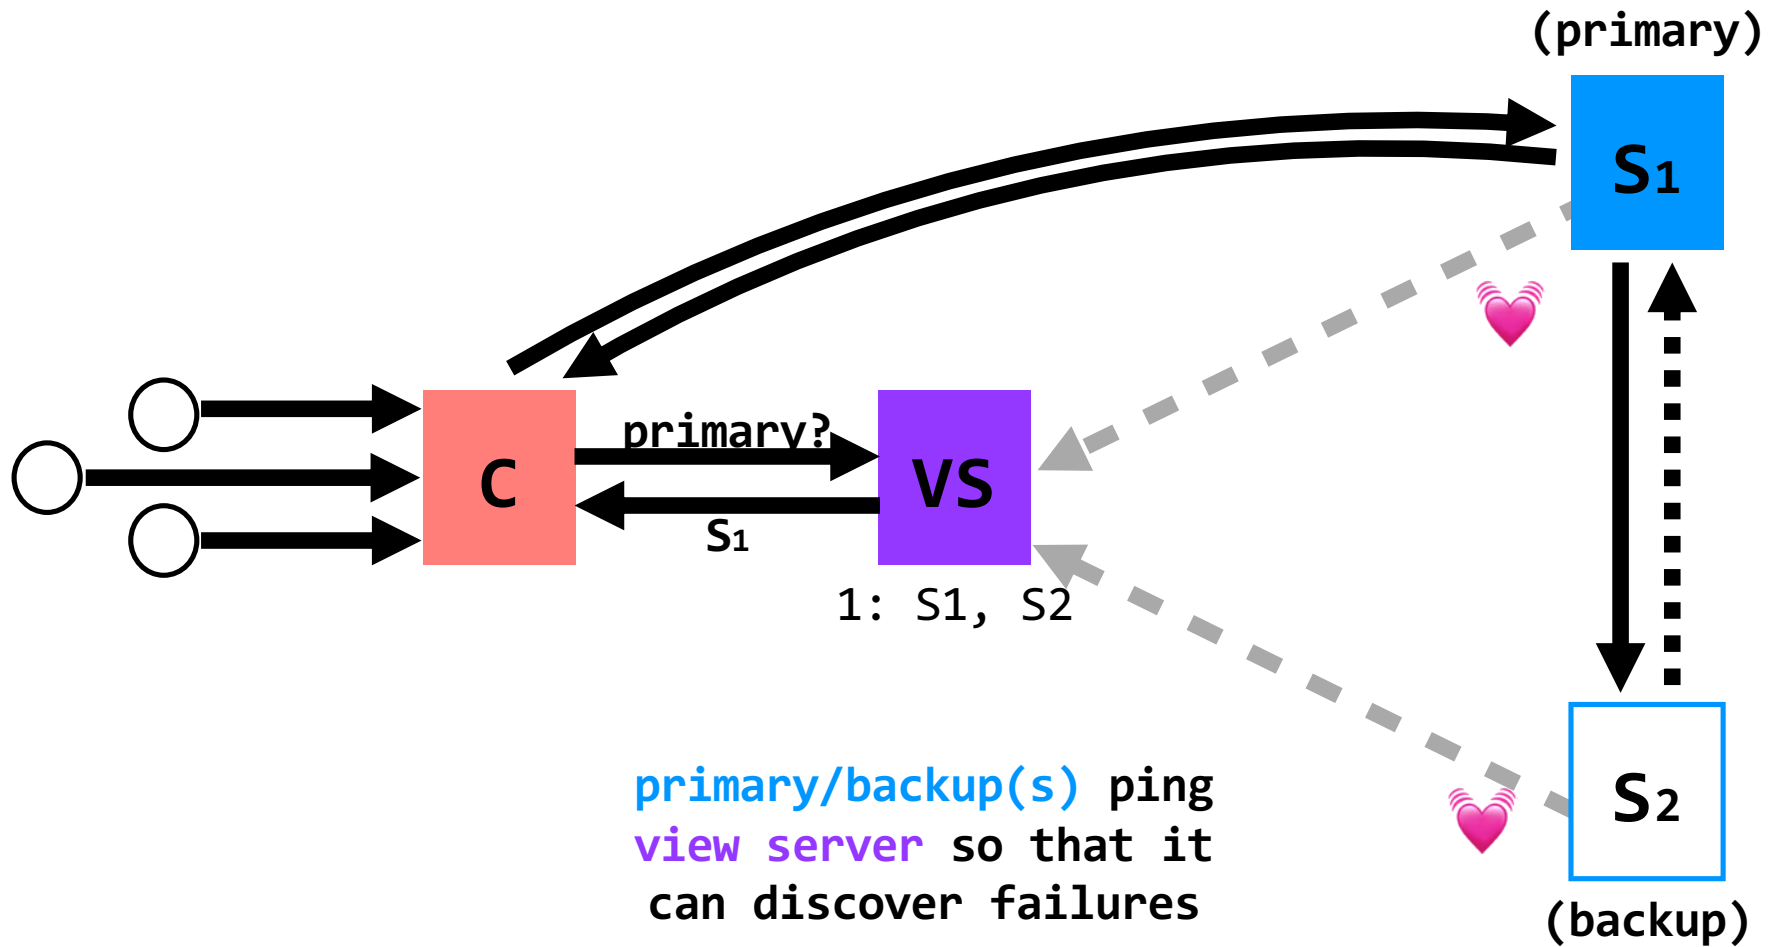
use a **view server**, which determines which replica is the primary



coordinators contact  
primary (as before)





use a **view server**, which determines which replica is the primary




use a **view server**, which determines which replica is the primary

# handling primary failure

(dead)





lack of pings indicates  
to **VS** that **S1** is down



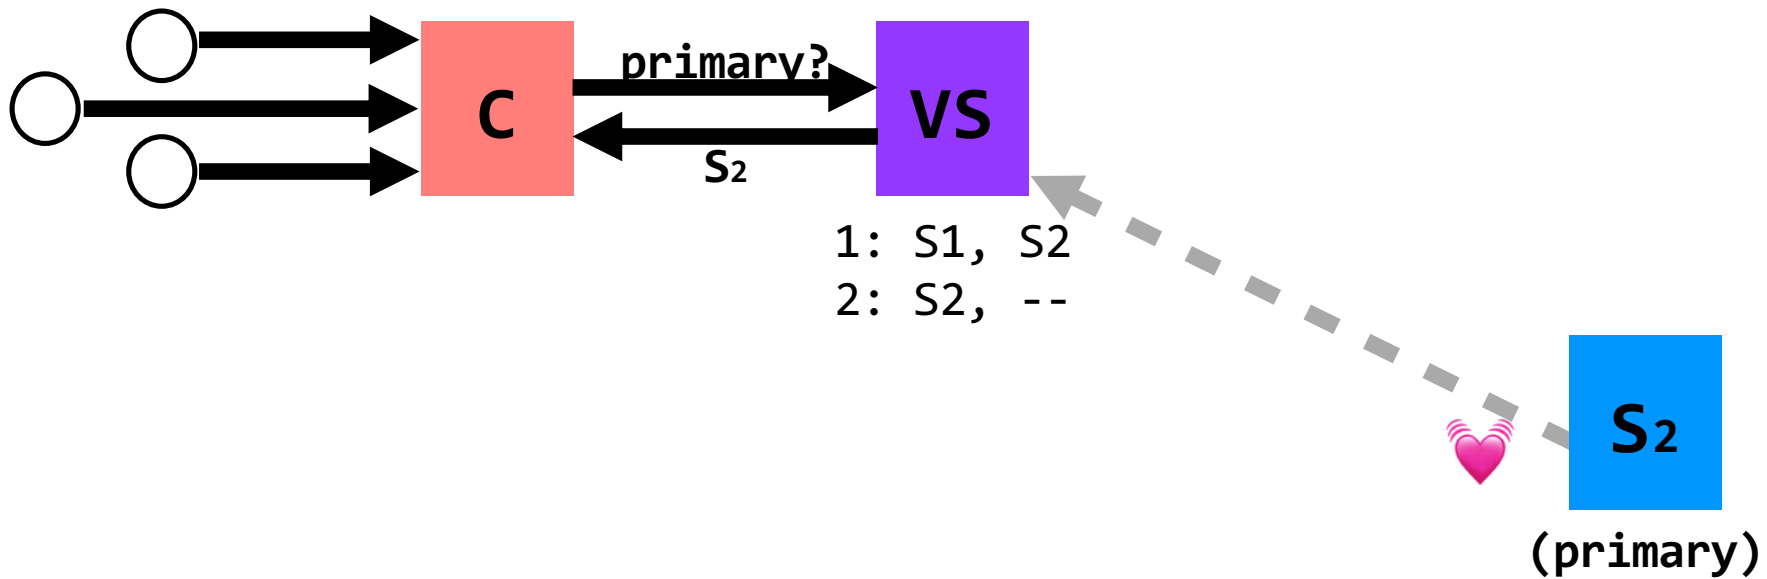
# handling primary failure

(dead)




# handling primary failure

(dead)

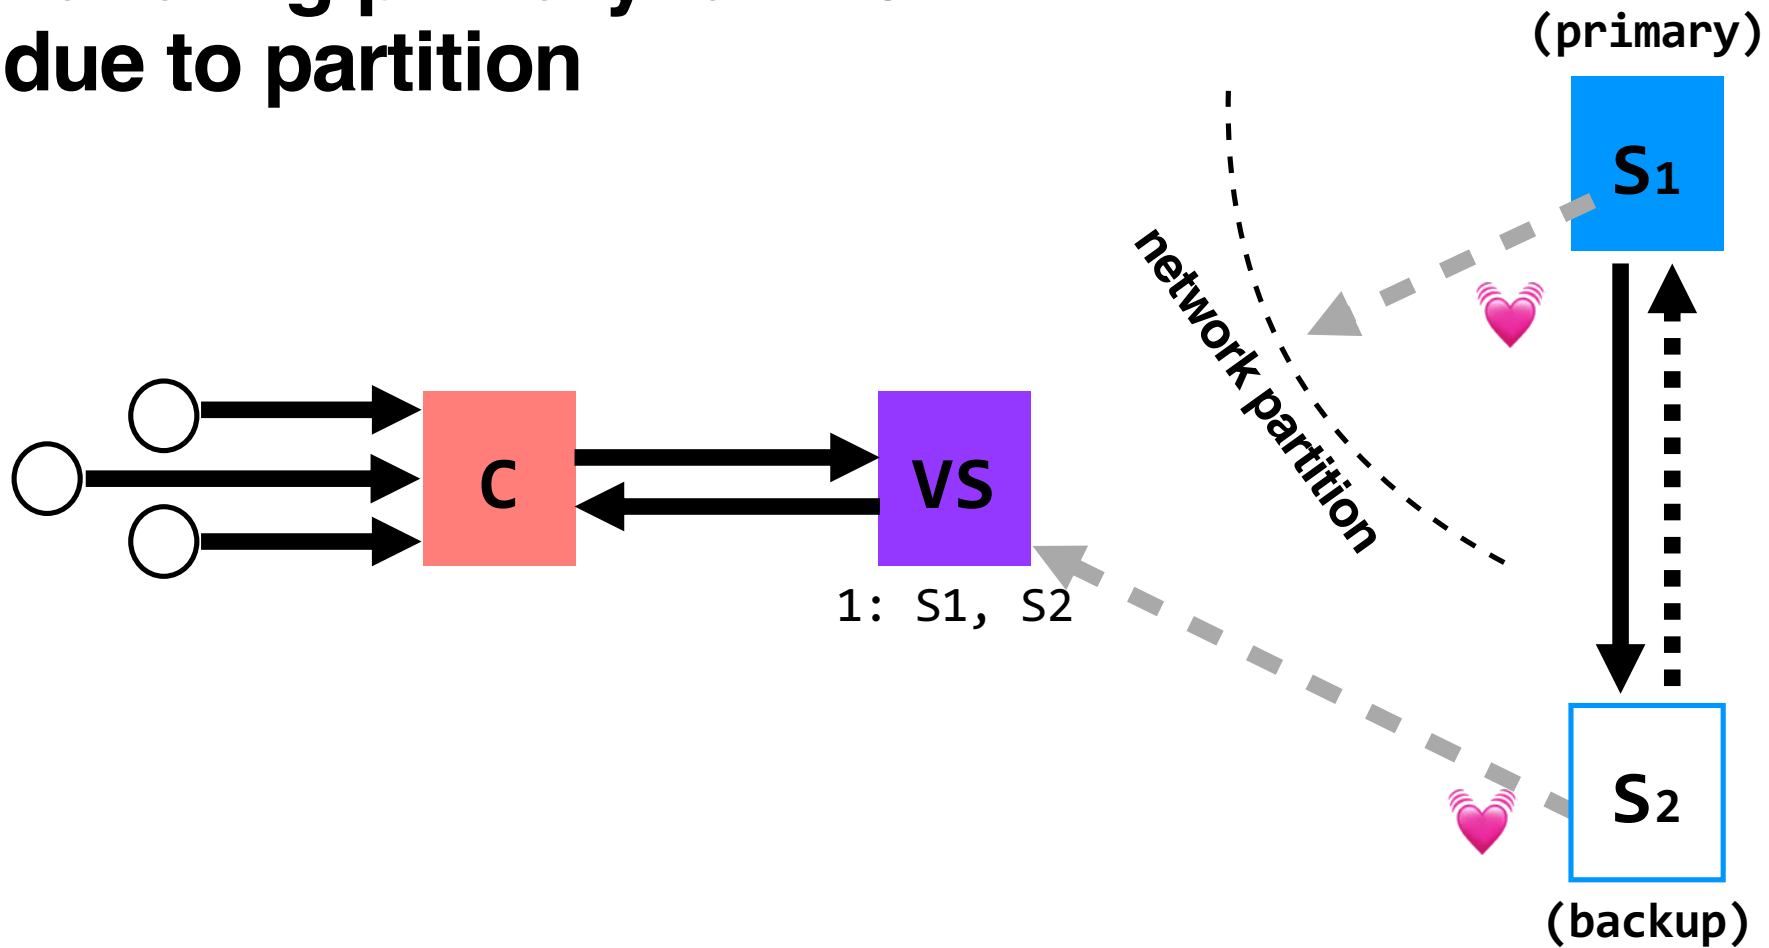


# handling primary failure

(dead)

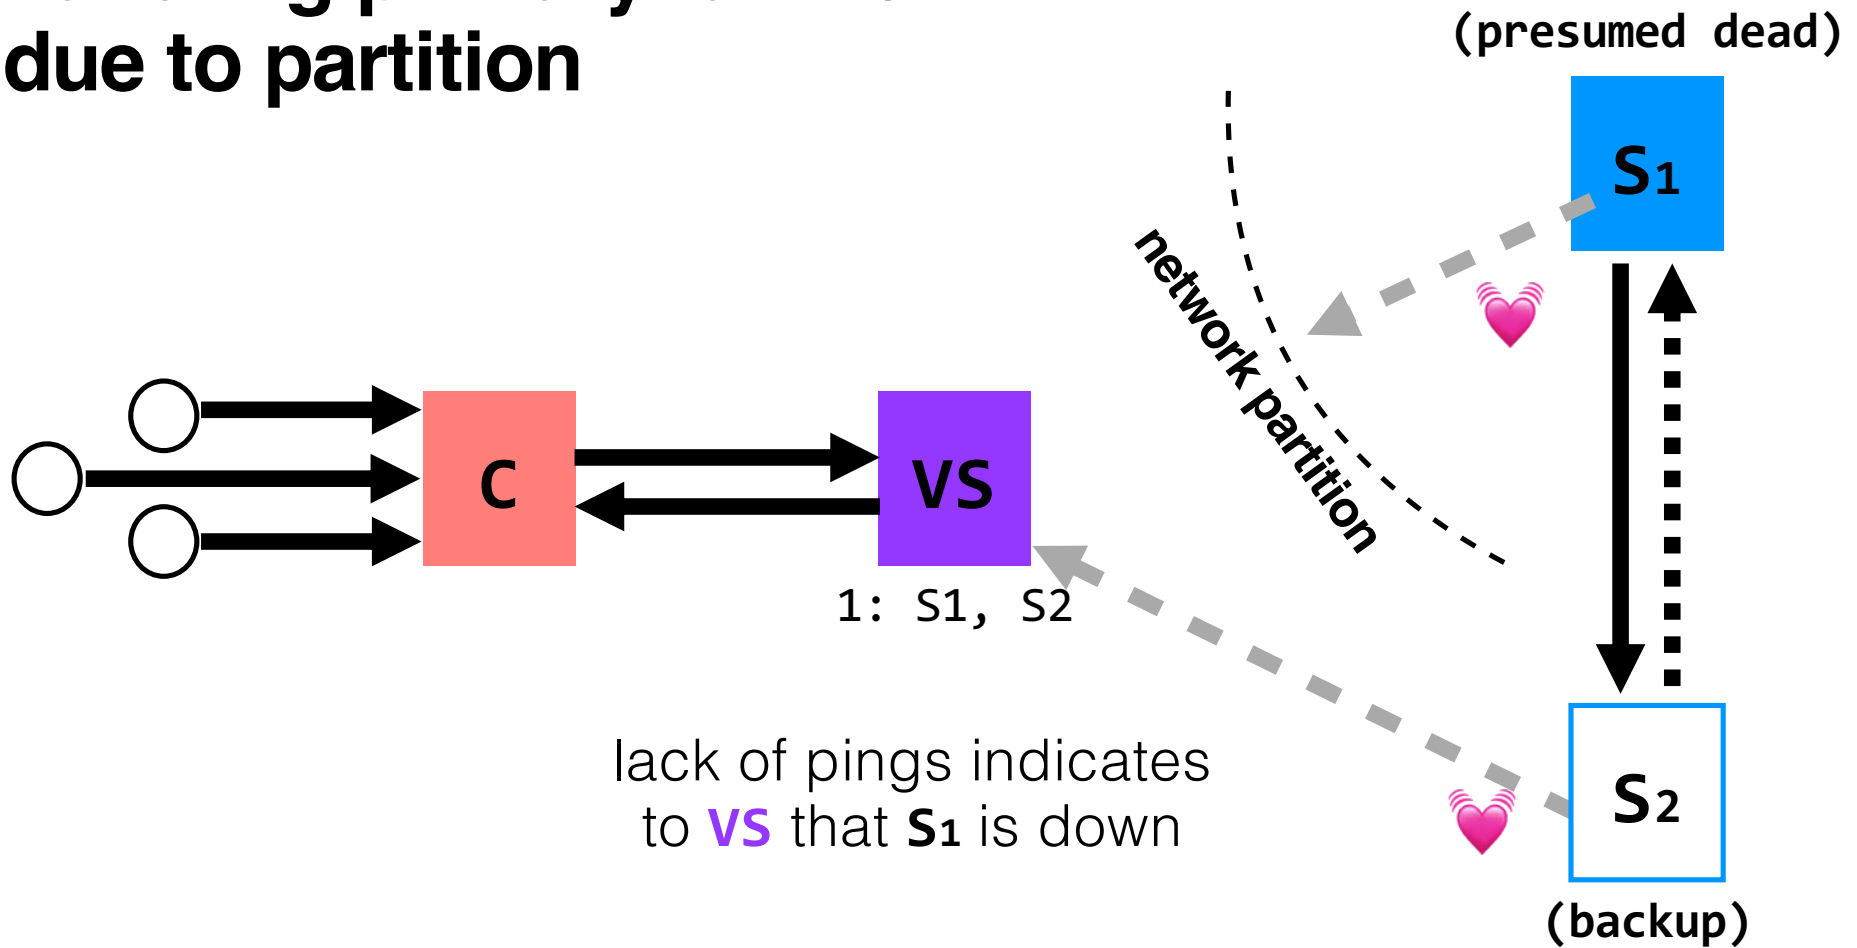


# handling primary failure due to partition



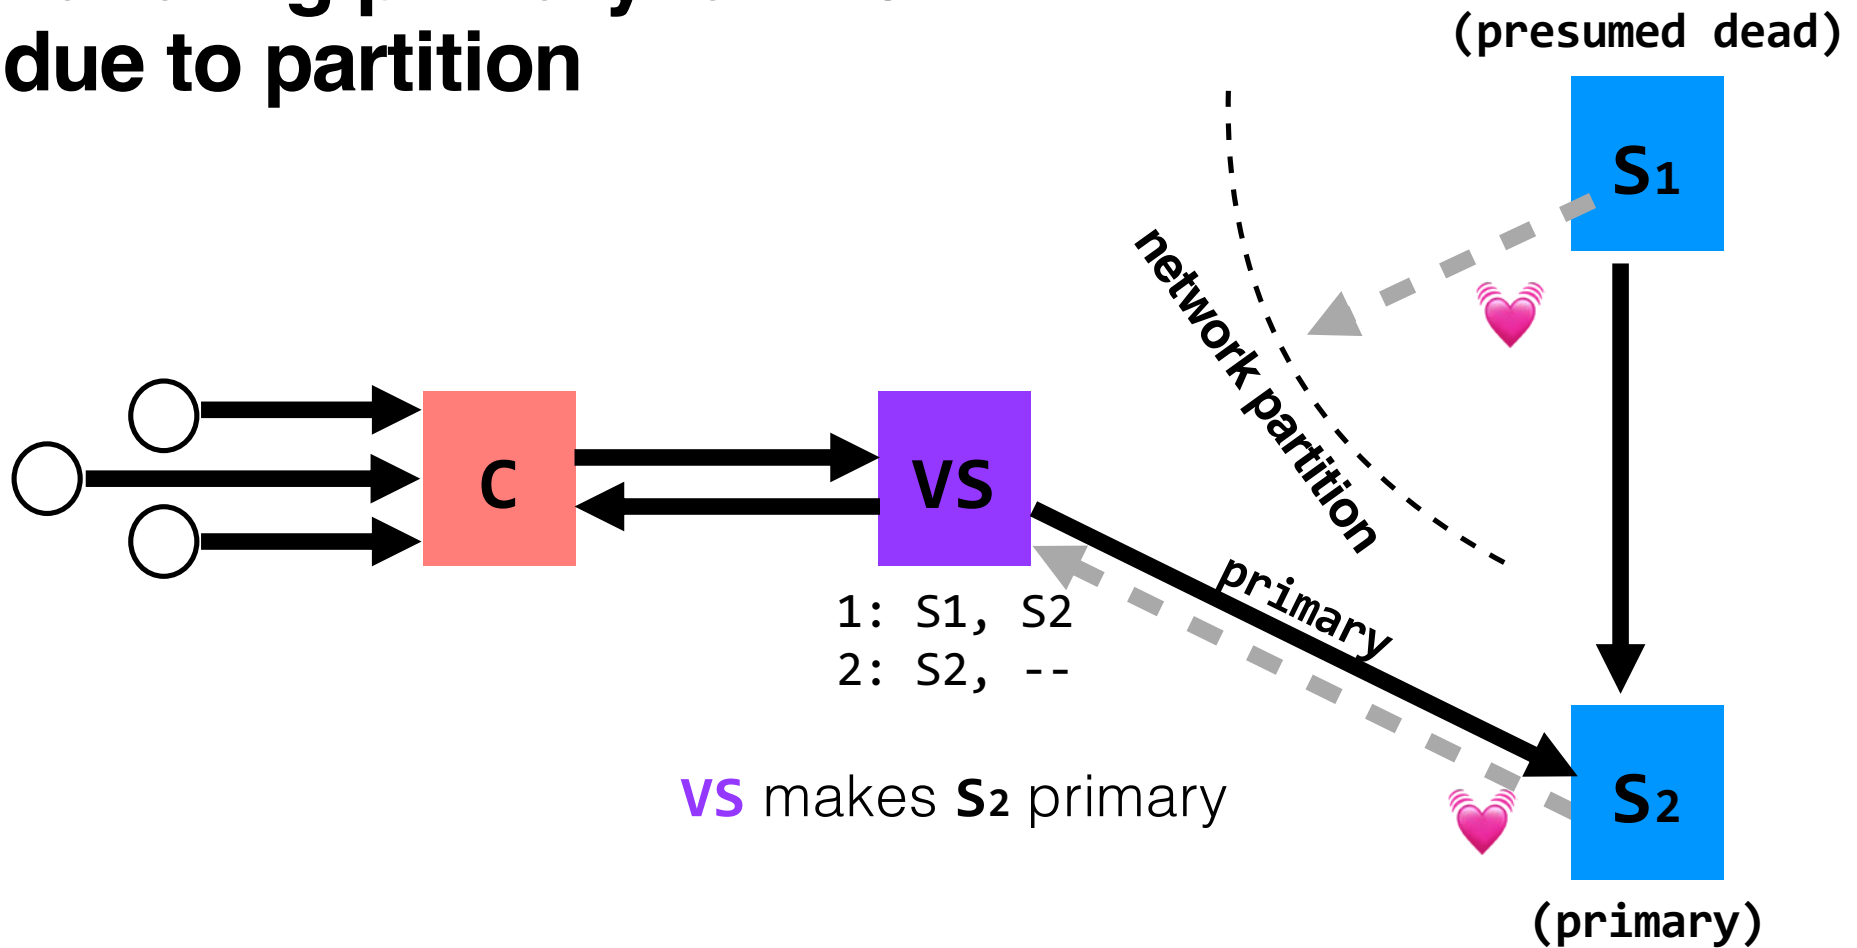
pose a partition keeps **S<sub>1</sub>** from communicating with the **view server**

# handling primary failure due to partition

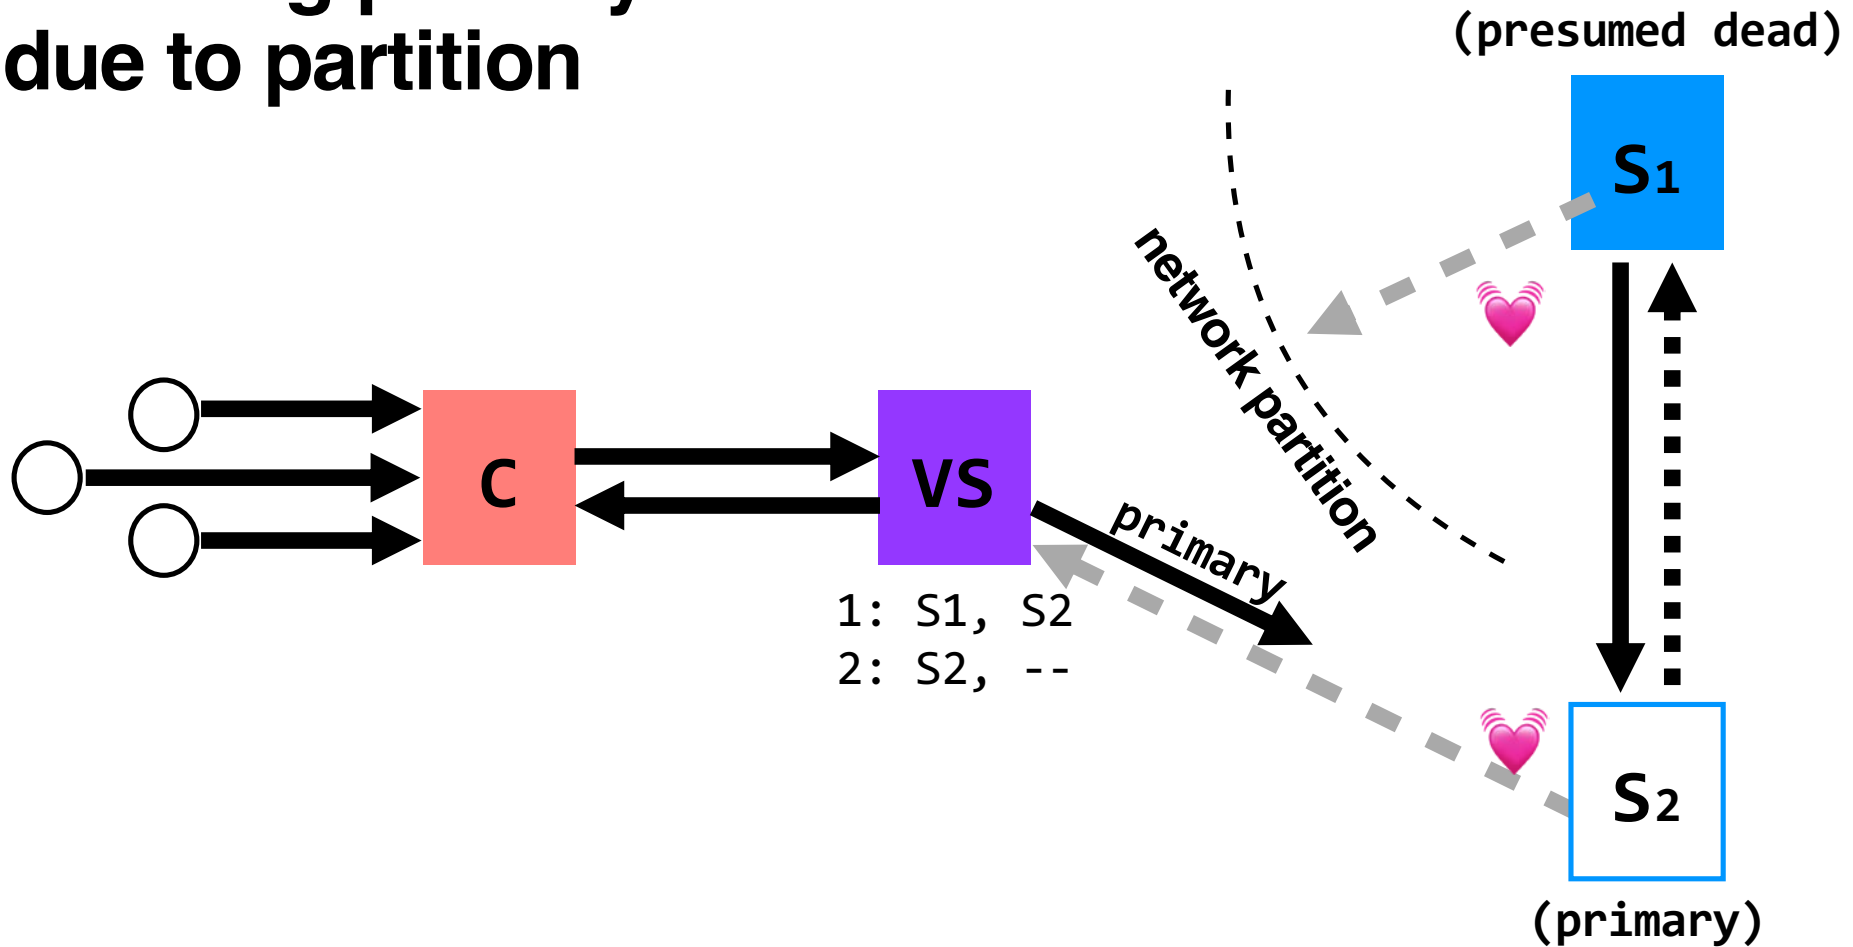




# handling primary failure due to partition

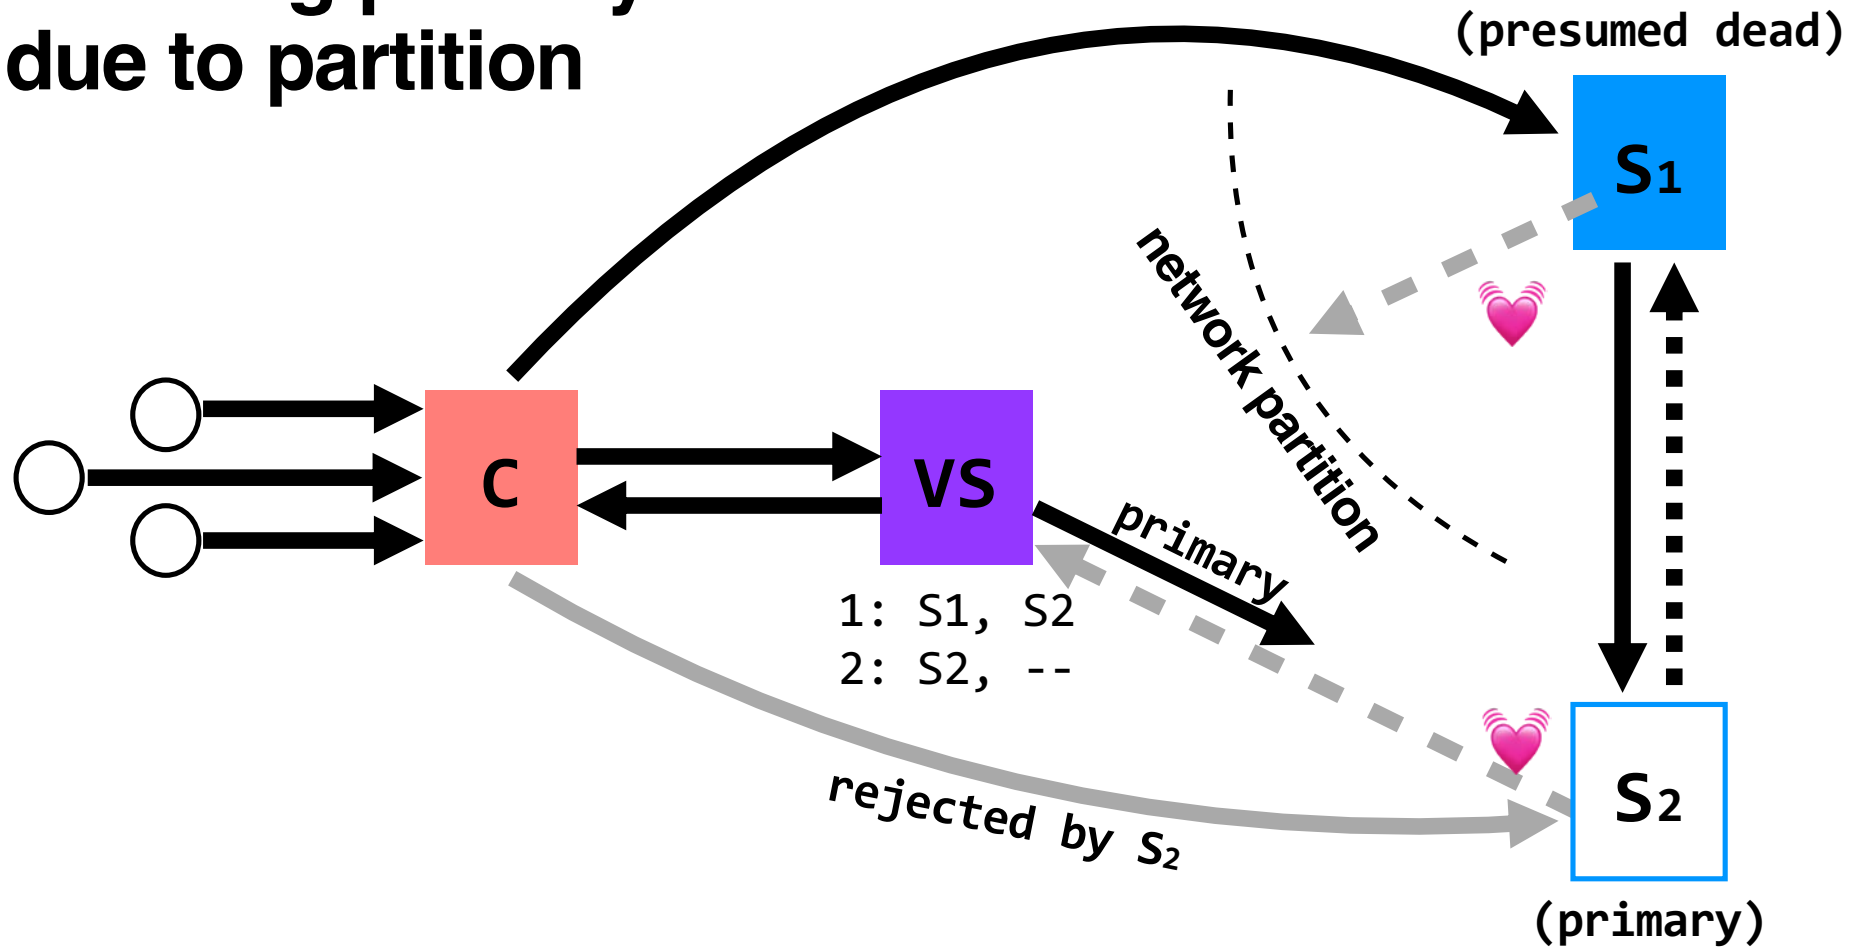


# handling primary failure due to partition



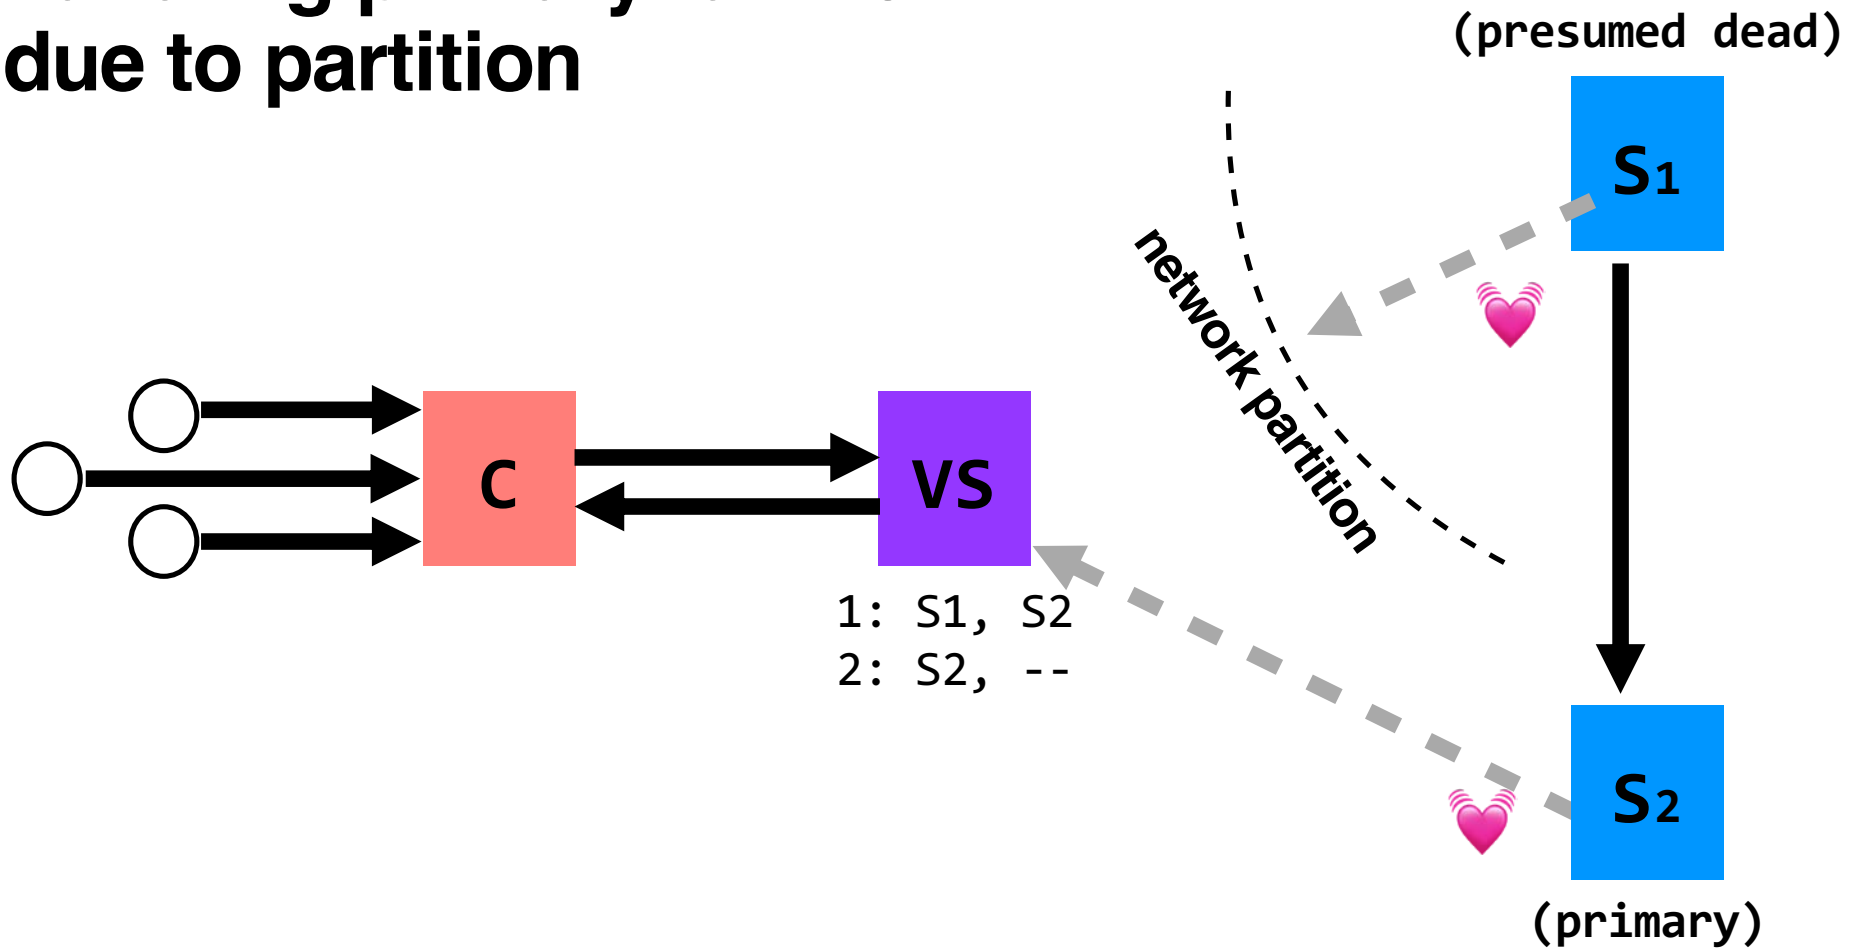
**question:** what happens before  $S_2$  knows  
it's the primary?

# handling primary failure due to partition



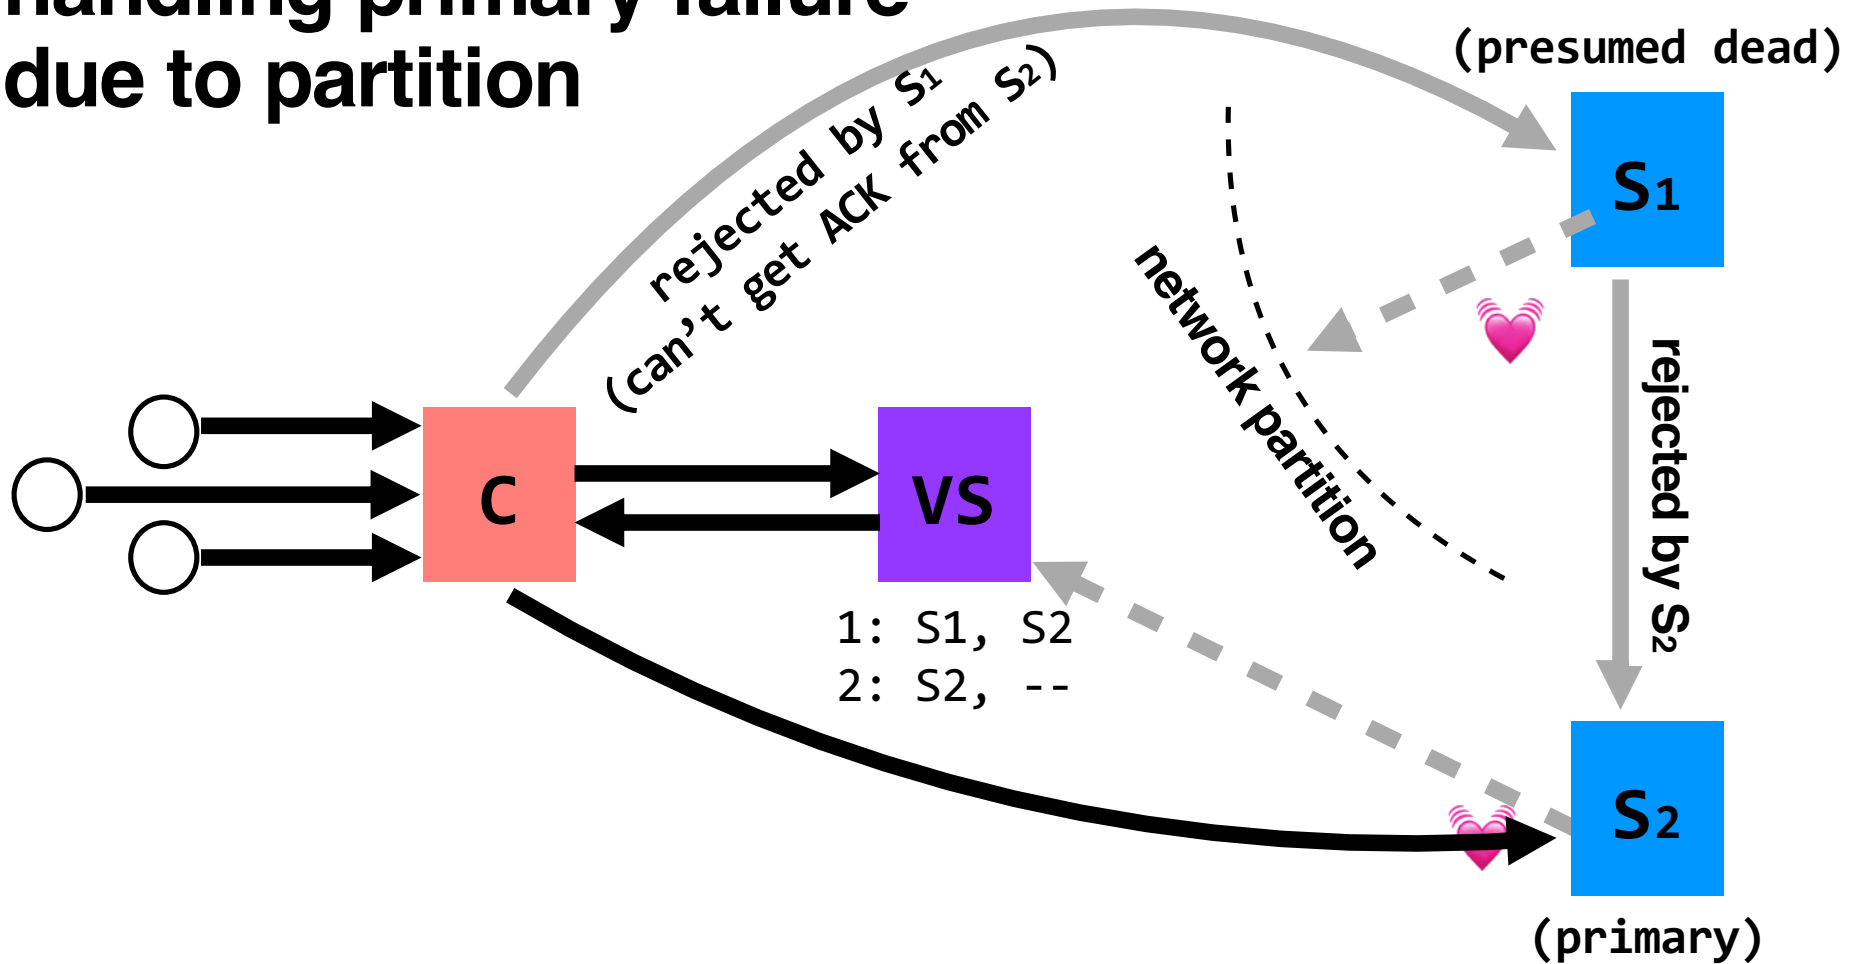
**S<sub>2</sub> will act as backup**  
(accept updates from S<sub>1</sub>, reject coordinator requests)

# handling primary failure due to partition




**question:** what happens after  $S_2$  knows it's the primary, but  $S_1$  also thinks it is?


# handling primary failure due to partition



**S<sub>1</sub> won't be able to act as primary**  
(can't accept client requests because it won't get ACKs from S<sub>2</sub>)



**problem:** what if view server fails?



**problem:** what if view server fails?

**go to recitation tomorrow and find out!**

- **Replicated state machines (RSMs)** provide **single-copy consistency**: operations complete as if there is a single copy of the data, though internally there are replicas.
- RSMs use a **primary-backup** mechanism for replication. The **view server** ensures that only one replica acts as the primary. It can also recruit new backups after servers fail.
- To extend this model to handle view-server failures, we need a mechanism to provide **distributed consensus**; see tomorrow's recitation (on Raft).



# Intro to Security

The following content is sourced from Computer Systems Design from MIT OCW

<https://ocw.mit.edu/courses/6-033-computer-system-engineering-spring-2018/pages/week-11/>

Suspicious event hijacks Amazon traffic for 2 hours, steals cryptocurrency

Secure | <https://arstechnica.com/information-technology/2018/04/suspicious-event-hijacks-amazon-traffic-for-2-hours-steals-cryptocurren...>

ars TECHNICA BIZ & IT TECH SCIENCE POLICY CARS GAMING & CULTURE FORUMS SUBSCRIPTIONS SIGN IN

BORDER GATEWAY PROTOCOL ATTACK —

# Suspicious event hijacks Amazon traffic for 2 hours, steals cryptocurrency

Almost 1,300 addresses for Amazon Route 53 rerouted for two hours.

DAN GOODIN - 4/24/2018, 3:00 PM



Amazon

108

Amazon lost control of a small number of its cloud services IP addresses for two hours on Tuesday morning when hackers exploited a known Internet-protocol weakness that let them to redirect traffic to rogue destinations. By subverting Amazon's domain-resolution service, the attackers masqueraded as cryptocurrency website MyEtherWallet.com and stole about

RISK ASSESSMENT —

# Yahoo says half a billion accounts breached by nation-sponsored hackers

One of the biggest compromises ever exposes names, e-mail addresses, and much more.

DAN GOODIN - 9/22/2016, 4:21 PM





LILY HAY NEWMAN SECURITY 04.18.17 7:00 AM

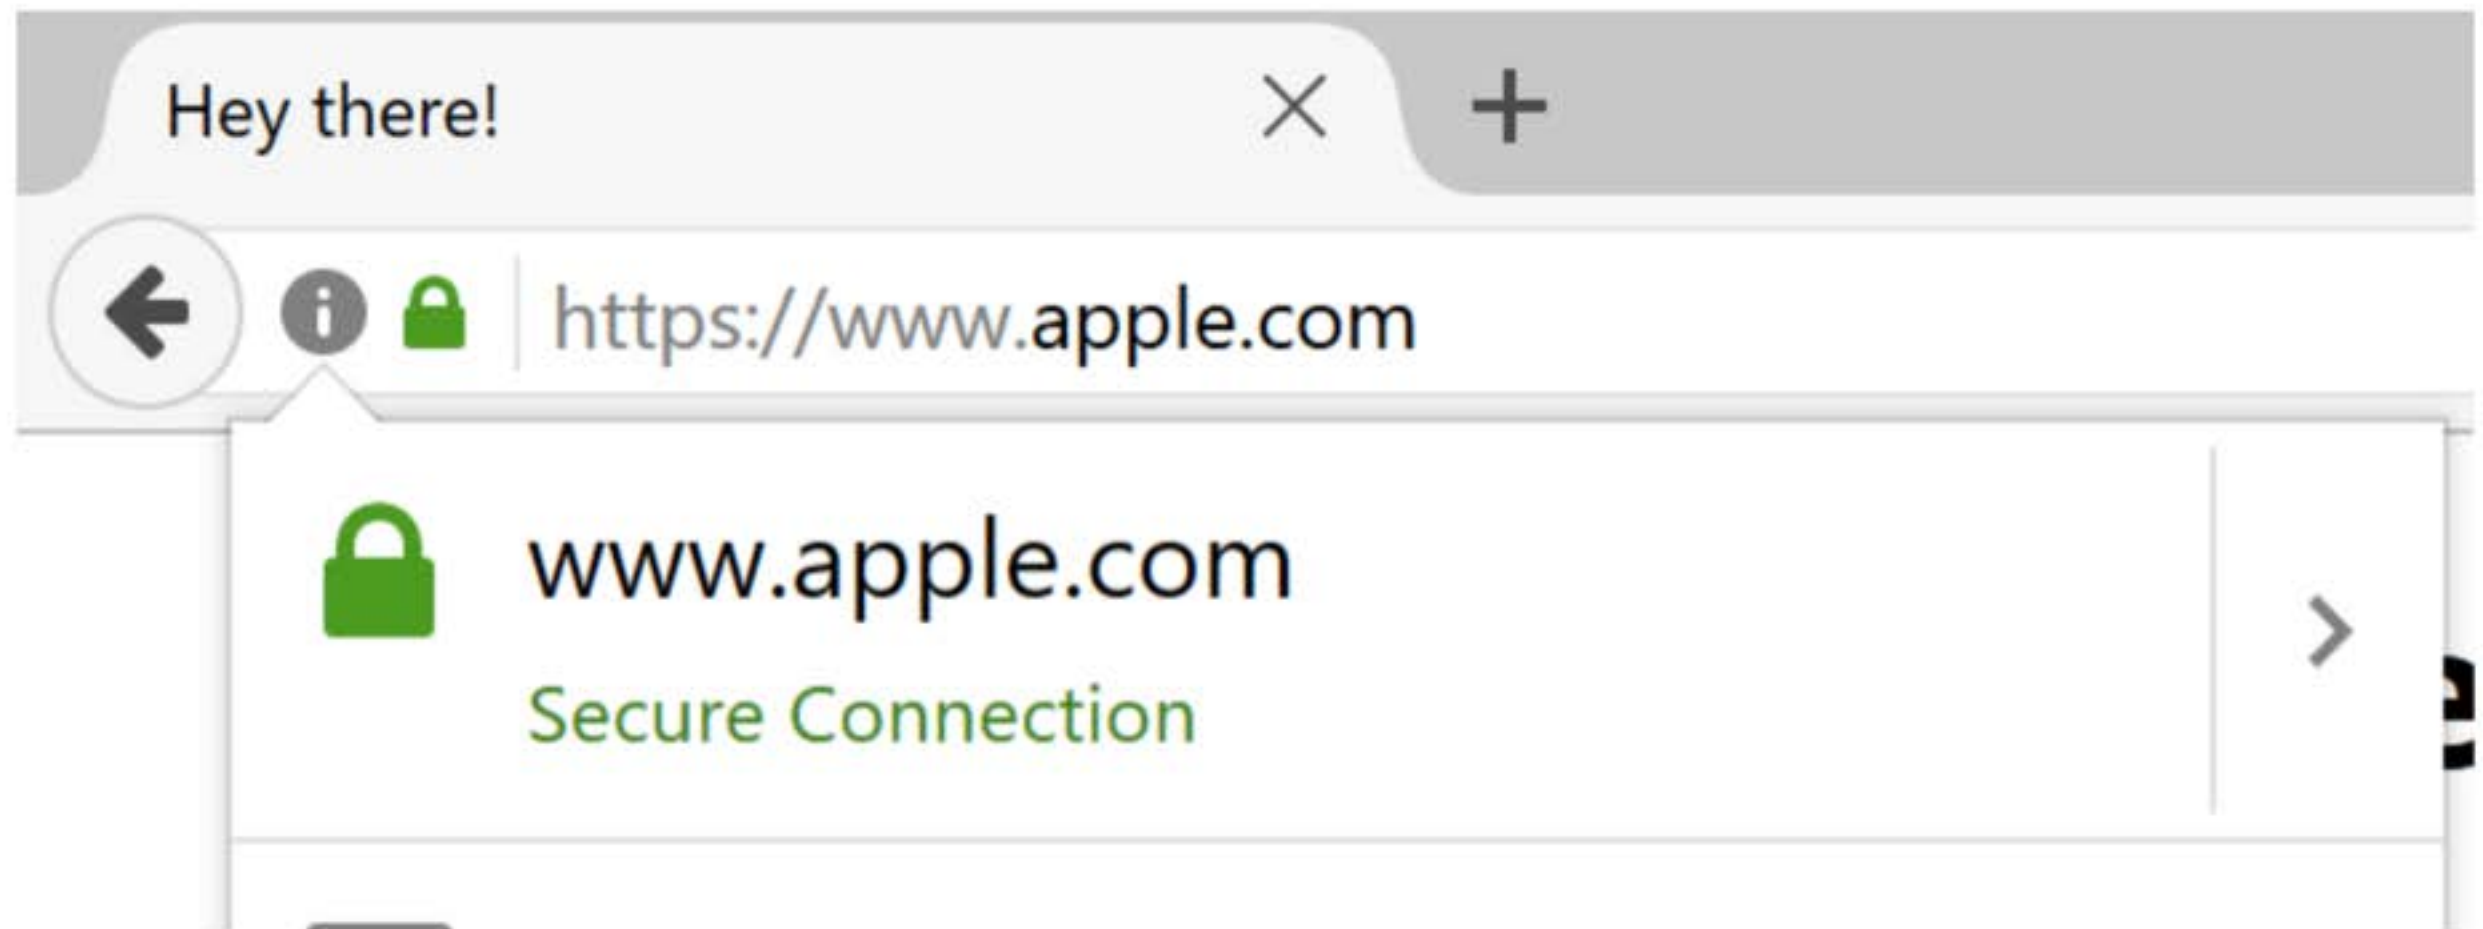
# SNEAKY EXPLOIT ALLOWS PHISHING ATTACKS FROM SITES THAT LOOK SECURE





# Phishing with Unicode Domains

Posted by [Xudong Zheng](#) on April 14, 2017



Before I explain the details of the vulnerability, you should take a look at the [proof-of-concept](#).

[Punycode](#) makes it possible to register domains with foreign characters. It works by converting individual domain label to an alternative format using only ASCII characters. For example, the domain "xn--s7y.co" is equivalent to "短.co".

From a security perspective, Unicode domains can be problematic because many Unicode characters are difficult to distinguish from common ASCII characters. It is possible to register domains such as "xn--pple-

RISK ASSESSMENT —

# BrickerBot, the permanent denial-of-service botnet, is back with a vengeance

New botnet squadrons wage fiercer, more intense attacks on unsecured IoT devices.

DAN GOODIN - 4/24/2017, 4:43 PM





Stuxnet Was Far More Dan... x

www.businessinsider.com/stuxnet-was-far-more-dangerous-than-previous-thought-2013-11

BUSINESS INSIDER

BI INTELLIGENCE EVENTS f t g+ in LOGIN REGISTER US EDITION

Tech Finance Politics Strategy Life Sports Video All Search

MILITARY & DEFENSE More: Stuxnet Iran Israel Cyberwarfare

# The Stuxnet Attack On Iran's Nuclear Plant Was 'Far More Dangerous' Than Previously Thought

MICHAEL B KELLEY NOV. 20, 2013, 12:58 PM 60,330 11

FACEBOOK LINKEDIN TWITTER

The Stuxnet virus that ravaged Iran's Natanz nuclear facility "was far more dangerous than the cyberweapon that is now



In-flight Wi-Fi is "direct li... x

arstechnica.com/security/2015/04/15/in-flight-wi-fi-is-direct-link-to-hackers/

Register Log in

ars technica

MAIN MENU MY STORIES: 25 FORUMS SUBSCRIBE JOBS ARS CONSORTIUM

## RISK ASSESSMENT / SECURITY & HACKTIVISM

### In-flight Wi-Fi is "direct link" to hackers

Report: Planes could be targeted by a malicious hacker on the ground.

by Michael Rundle Apr 15, 2015 11:03am EDT

Share Tweet 88



LATEST FEATURE STORY



FEATURE STORY (2 PAGES)

#### The promise—and massive challenge—of making games for the Apple Watch

How to make 15-second microgames with targets "the size of salad bar ham cubes"

WATCH ARS VIDEO



Meet the e-voting machin... x +

arstechnica.com/tech-policy/2015/04/15/meet-the-e-voting-machine-so-easy-to-hack-it-will-take-your-breath-away/

Register Log in

ars technica

MAIN MENU MY STORIES: 25 FORUMS SUBSCRIBE JOBS ARS CONSORTIUM

# LAW & DISORDER / CIVILIZATION & DISCONTENTS

## Meet the e-voting machine so easy to hack, it will take your breath away

Virginia decertifies device that used weak passwords and wasn't updated in 10 years.

by Dan Goodin - Apr 15, 2015 2:55pm EDT

Share Tweet 156



WINVote

WATCH ARS VIDEO

LATEST FEATURE STORY



FEATURE STORY (2 PAGES)

### The promise—and massive challenge—of making games for the Apple Watch

How to make 15-second microgames with targets "the size of salad bar ham cubes"

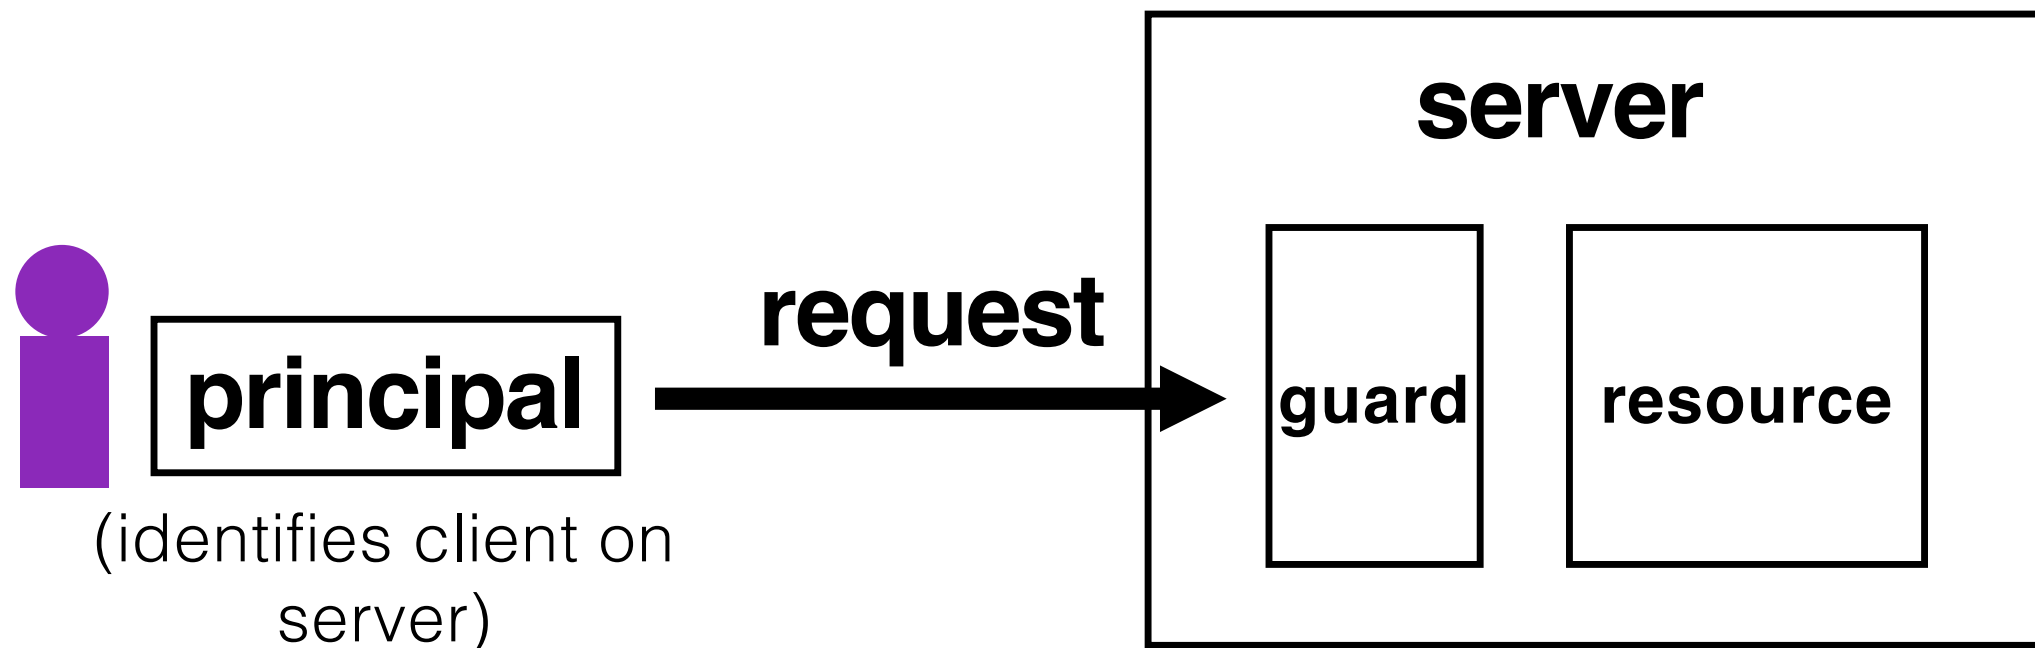
**what makes computer security  
special?**

# **why is security difficult?**

# **steps towards building a more secure system:**

- 1. be clear about goals (**policy**)**
- 2. be clear about assumptions  
(**threat model**)**

**complete mediation:** every request for resource goes through the guard



**authentication:** is the principal who they claim to be?

**authorization:** does principal have access to perform request on resource?

**what can go wrong with the guard  
model?**

# sql injection demo

username	email	public?
karen	karen@fake.com	yes
peter	peter@fake.com	yes
katrina	no	

**SELECT** username, email **FROM** users **WHERE**  
username='<username>' **AND** public='yes'

Let <username> = **katrina' OR username=**

# sql injection demo

username	email	public?
karen	karen@fake.com	yes
peter	peter@fake.com	yes
katrina	no	

```
SELECT username, email FROM users WHERE  
username='katrina' OR username=' ' AND  
public='yes'
```



```
> cd /mit/bob/project  
> cat ideas.txt  
Hello world.  
  
...  
> mail alice@mit.edu < ideas.txt
```

**what can go wrong with the guard  
model?**

- **Adversarial attacks** are different from “normal” failures. They’re targeted, rarely random, and rarely independent. Just one successful attack can bring down a system.
- Securing a system starts by specifying our goals (**policy**) and assumptions (**threat model**).
- The **guard model** provides **complete mediation**. Even though things can still go wrong, systems that use this model avoid common pitfalls.

# Security (Cont.)

The follow slides are used with permission from Professor Werner Dietl from Fall 2022

# Software Architecture & Design

SE 464

Week 5: 04.10.2022  
Security



Werner Dietl

<https://ece.uwaterloo.ca/~wdietl/>

# NFP: Security

“The protection afforded to an automated information system in order to attain the applicable objectives of preserving the **integrity, availability** and **confidentiality** of information system resources (includes hardware, software, firmware, information/data, and telecommunications).”

National Institute of Standards and  
Technology

# NFP: Security

- **Confidentiality:** Preserving the confidentiality of information means preventing unauthorized parties from accessing the information or perhaps even being aware of the existence of the information.
- **Integrity:** Maintaining the integrity of information means that only authorized parties can manipulate the information and do so only in authorized ways.
- **Availability:** Resources are available if they are accessible by authorized parties on all appropriate occasions.

# Design Principles for Security

- **Least Privilege:** give each component only the privileges it requires
- **Fail-safe Defaults:** deny access if explicit permission is absent
- **Economy of Mechanism:** adopt simple security mechanisms
- **Complete Mediation:** ensure every access is permitted
- **Open Design:** do not rely on secrecy for security





# Design Principles for Security

- **Separation of Privilege:** introduce multiple parties to avoid exploitation of privileges
- **Least Common Mechanism:** limit critical resource sharing to only a few mechanisms
- **Psychological Acceptability:** make security mechanisms usable
- **Defense in Depth:** have multiple layers of countermeasures

# Security terminology

<https://www.us-cert.gov/bsi/sdlc/design>

<http://web.mit.edu/Saltzer/www/publications/protection/Basic.html>

# Architectural Access Control Models

Decide whether access to a protected resource should be granted or denied

- Discretionary access control  
Based on the identity of the requestor, the resource, and whether the requestor has permission to access
- Mandatory access control  
Policy based

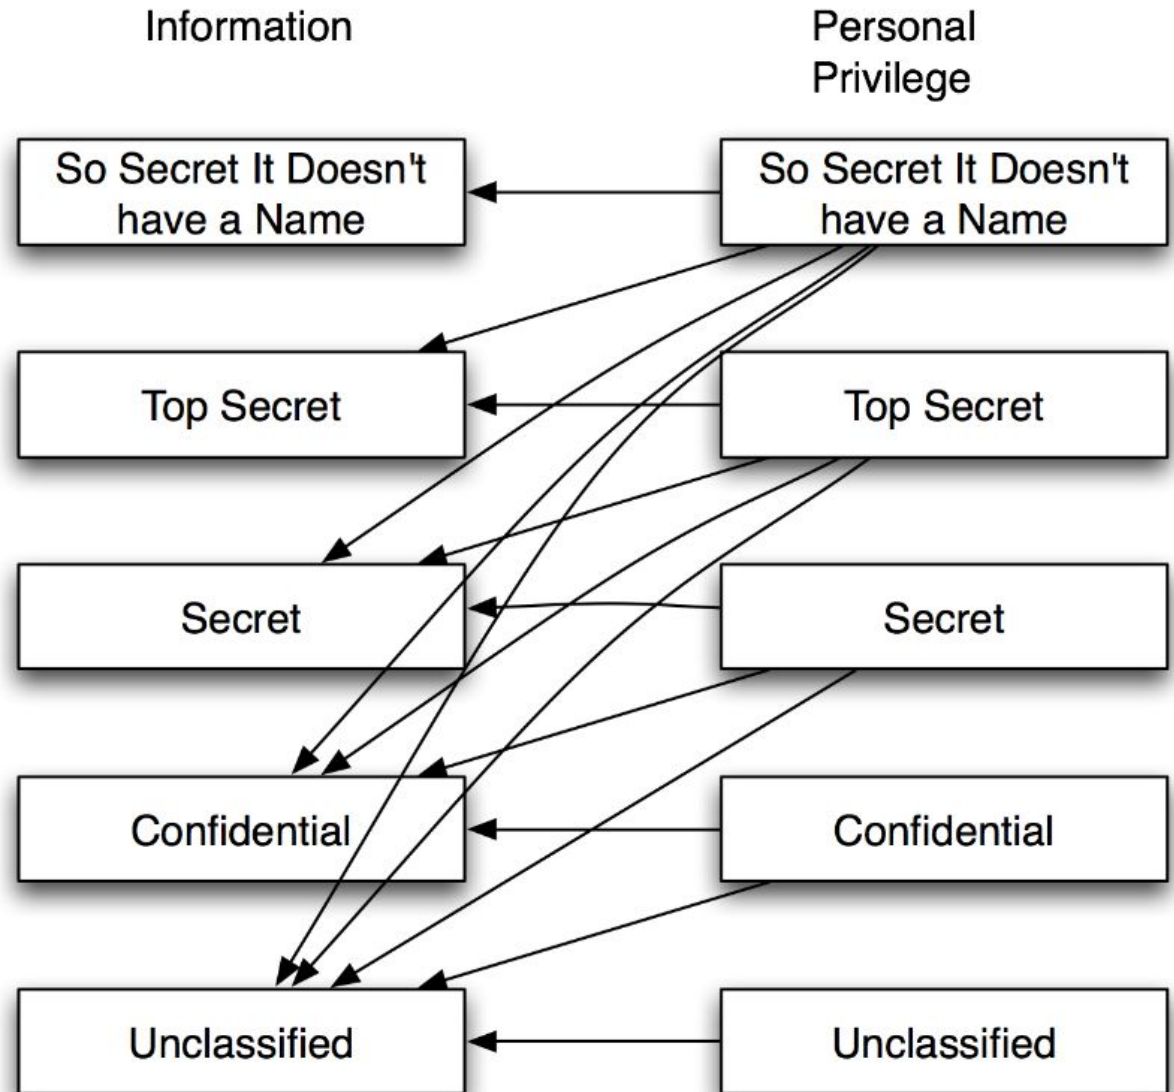
# Discretionary Access Control

	Database A	Component Q	Interface F
<i>Alice</i>	Read-Write; Always	Bend	Yes
<i>Bob</i>	Read-Write; Between 9 and 5	Fold	No
<i>Charles</i>	No access	Spindle	No
<i>Dave</i>	No access	Mutilate	Yes
<i>Eve</i>	Read-only; Always	None	No

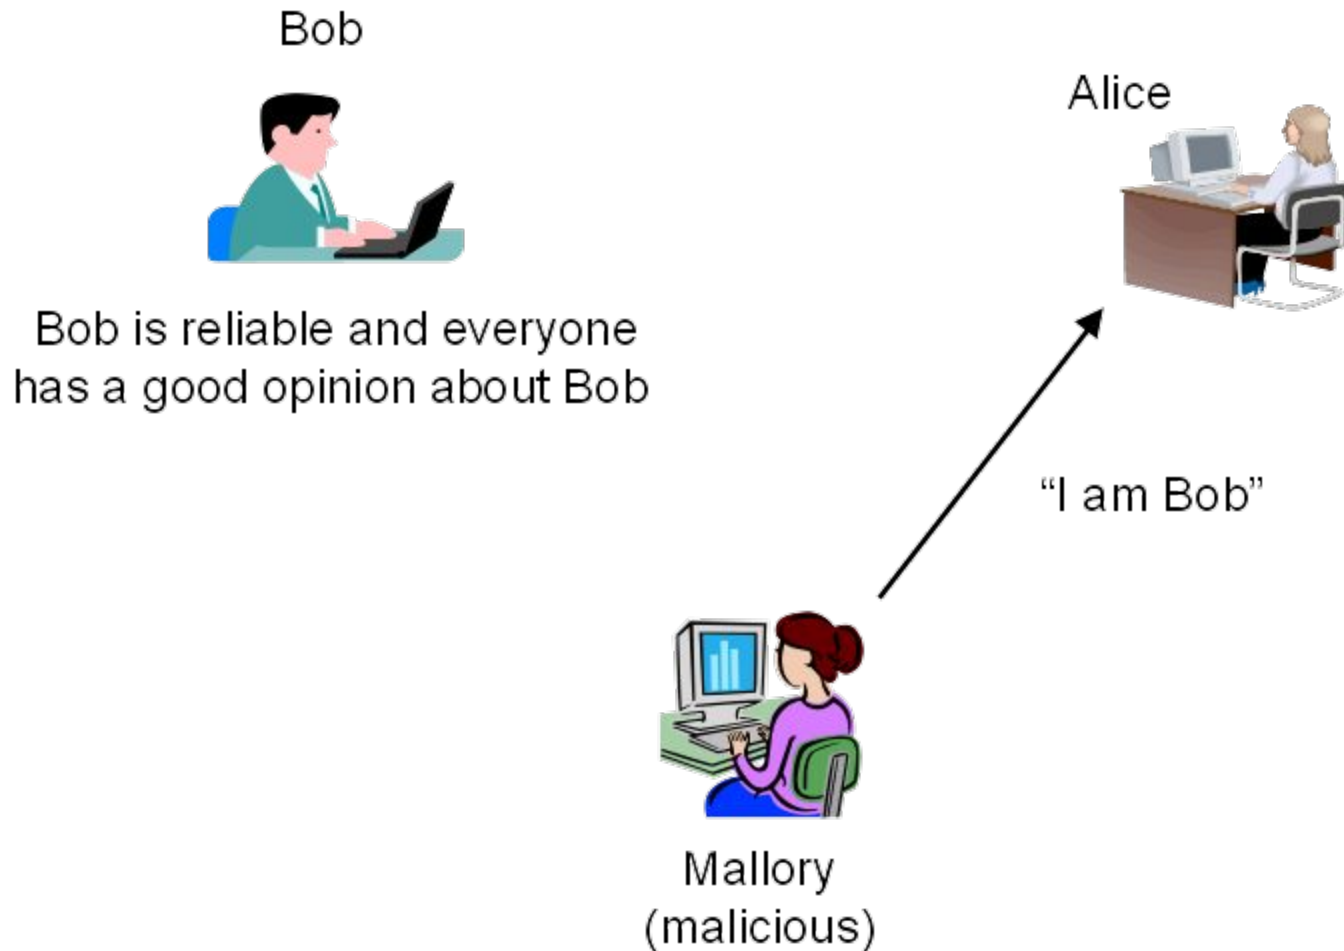


# Mandatory Access Control

Bob: Secret  
Alice: Confidential  
Tom: Top Secret



# Impersonation



# Misrepresentation

Bob



Bob is reliable and everyone  
has a good opinion about Bob

Alice



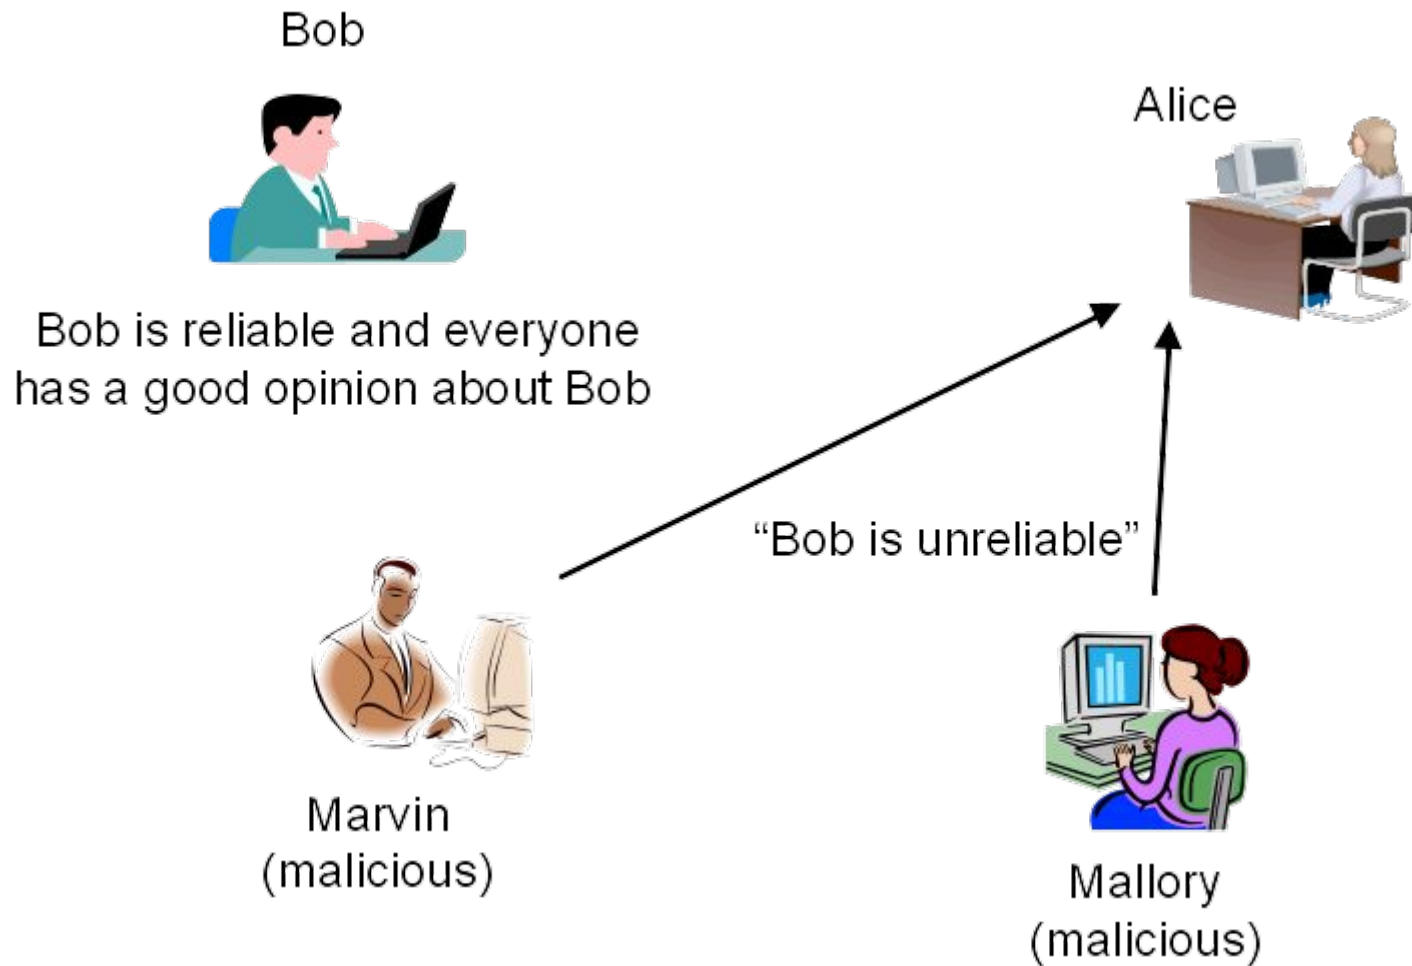
"Bob is unreliable"



Mallory  
(malicious)




# Collusion

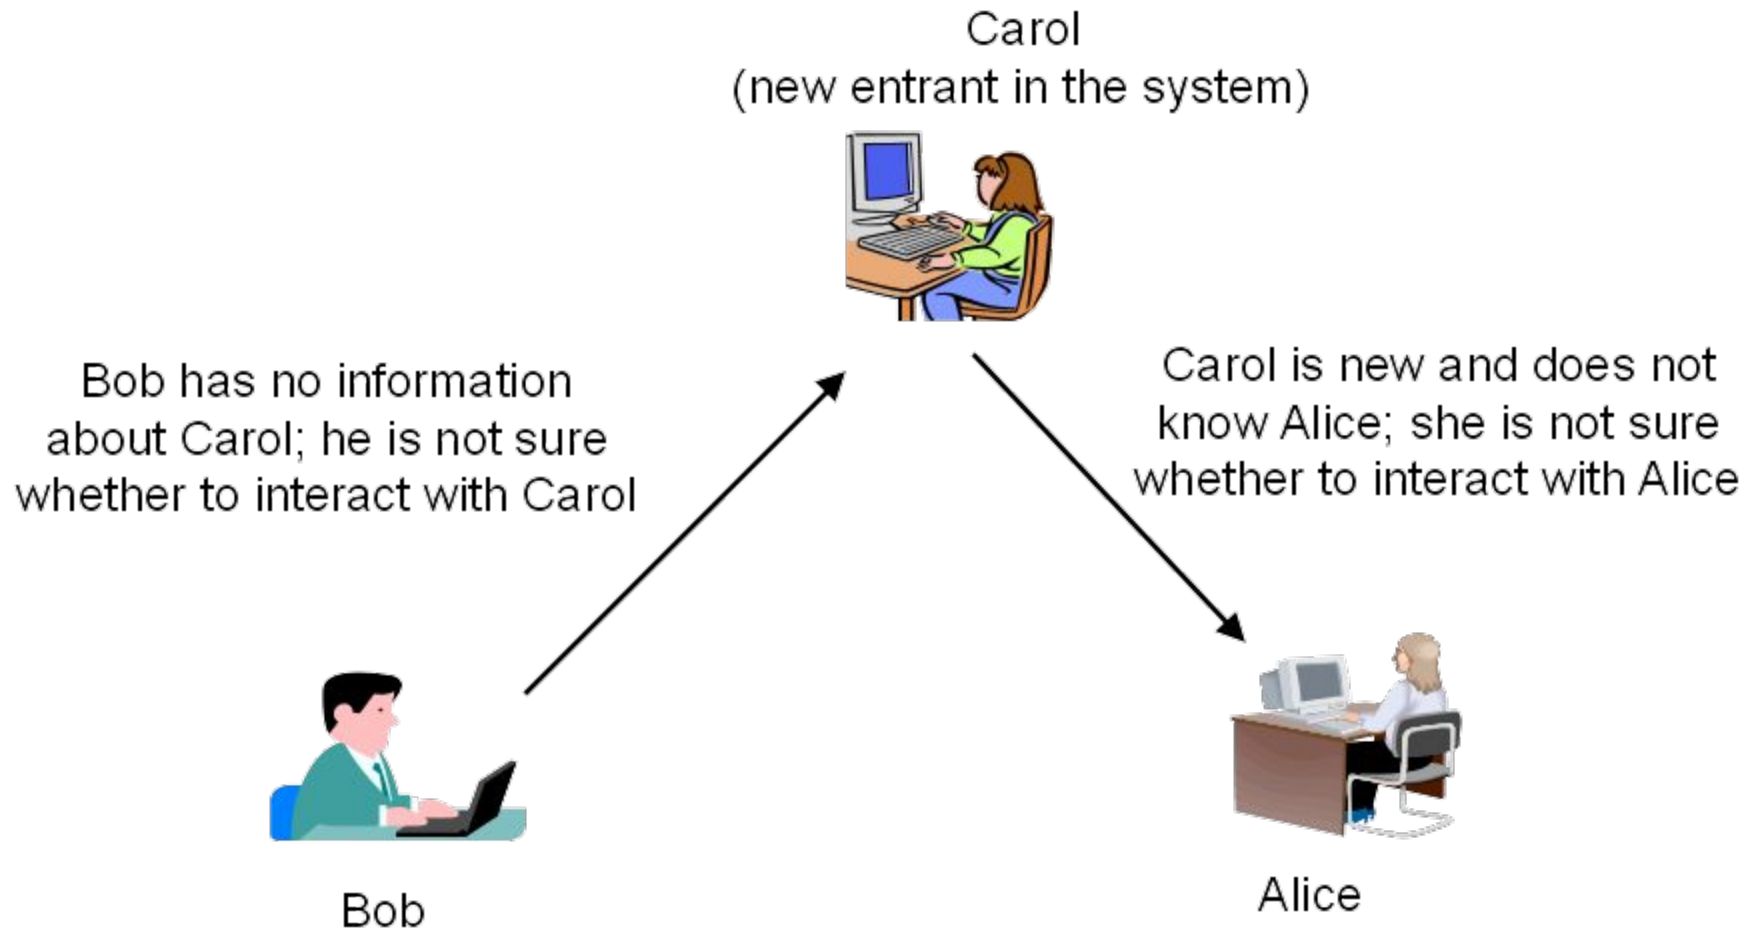




# Fraudulent Actions



# Addition of Unknowns



# Trust Management

- **Trust** is a particular level of the subjective probability with which an agent assesses that another agent will perform a particular action in a context that affects his actions.
- **Reputation** is the expectation about an entity's behavior based on past behavior.  
May be used to determine trust

Two types of trust management systems

- Credential and Policy-based
- Reputation-based

# Architecture in Practice: Chrome



Read:

- Browser Security: Lessons from Google Chrome

<http://queue.acm.org/detail.cfm?id=1556050>

- The Security Architecture of the Chromium Browser

<http://seclab.stanford.edu/websec/chromium/chromium-security-architecture.pdf>



# Security risks

Online content is insecure and can compromise:

- Confidentiality: Leak user data
- Integrity: Read/write arbitrary data on disk
- Availability: Crash host application and/or OS

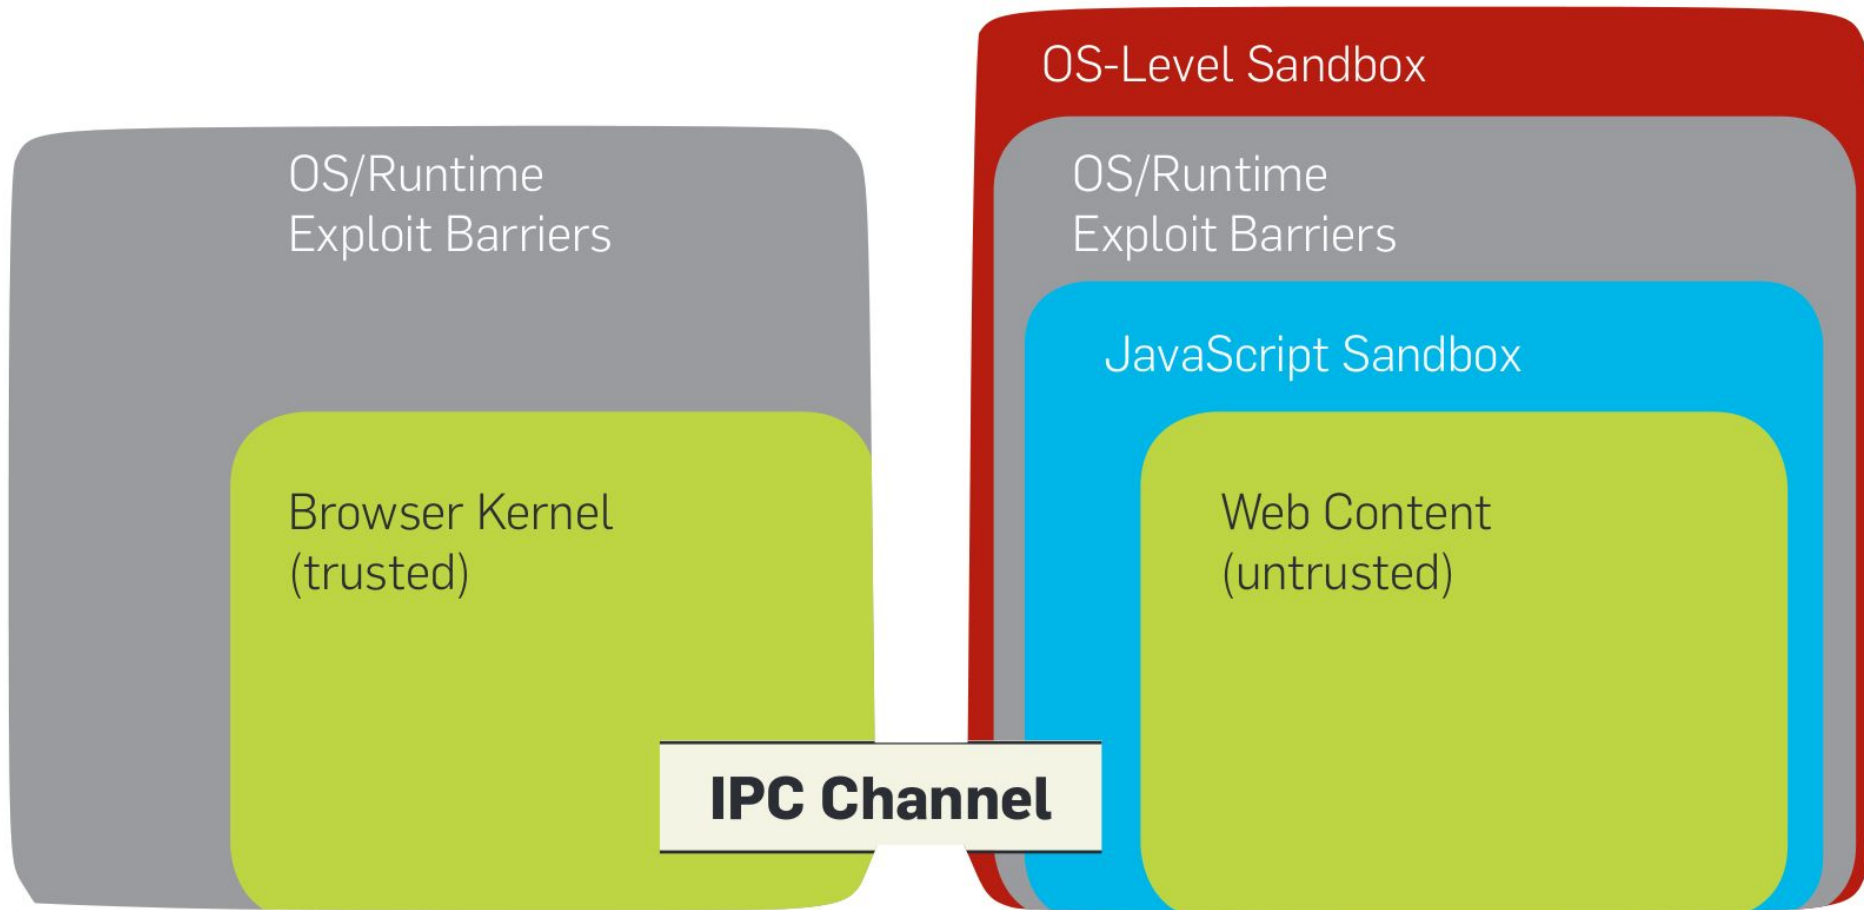
Chrome relies on **least privilege**, **separation of privilege**, and **defence in depth** to securely parse and render insecure content.

# Main factors for Chrome security

## Reduce

- the severity of vulnerabilities
  - Sandboxing
- the window of vulnerability
  - Auto-update (needs automated testing)
- the frequency of exposure
  - Warn about malicious sites

# Chrome architecture




# Authentication and Passwords

The following content is sourced from Computer Systems Design from MIT OCW

<https://ocw.mit.edu/courses/6-033-computer-system-engineering-spring-2018/pages/week-12/>



**complete mediation:** every request for resource goes through the guard



**guard typically provides:**

**authentication:** is the principal who they claim to be?

**authorization:** does principal have access to perform request on resource?

Rank	2011	2012	2013	2014	2015	2016	2017
1	password	password	123456	123456	123456	123456	123456
2	123456	123456	password	password	password	password	password
3	12345678	1234567	12345678	12345	12345678	12345	12345678
4	qwerty	abc123	qwerty	12345678	qwerty	12345678	qwerty
5	abc123	qwerty	abc123	qwerty	12345	football	12345
6	monkey	monkey	123456789	123456789	123456789	qwerty	123456789
7	1234567	letmein	111111	1234	football	1234567890	letmein
8	letmein	dragon	1234567	baseball	1234	1234567	1234567
9	trustno1	111111	iloveyou	dragon	1234567	princess	football
10	dragon	baseball	adobe123	football	baseball	1234	iloveyou
11	baseball	iloveyou	123123	1234567	welcome	login	admin
12	111111	trustno1	admin	monkey	123456789	welcome	welcome
13	iloveyou	1234567	1234567890	letmein	abc123	solo	monkey
14	master	sunshine	letmein	abc123	111111	abc123	login
15	sunshine	master	photoshop	111111	1qaz2wsx	admin	abc123
16	ashley	123123	1234	mustang	dragon	121212	starwars
17	bailey	welcome	monkey	access	master	flower	123123
18	passw0rd	shadow	shadow	shadow	monkey	passw0rd	dragon
19	shadow	ashley	sunshine	master	letmein	dragon	passw0rd
20	123123	football	12345	michael	login	sunshine	master
21	654321	jesus	password1	superman	princess	master	hello
22	superman	michael	princess	696969	qwertyuiop	hottie	freedom
23	qazwsx	ninja	azerty	123123	solo	loveme	whatever
24	michael	mustang	trustno1	batman	passw0rd	zaq1zaq1	qazwsx
25	Football	password	000000	trustno1	starwars	password1	trustno1

© Wikipedia. All rights reserved. This content is excluded from our Creative Commons license.  
For more information, see <https://ocw.mit.edu/help/faq-fair-use>.

**problem:** users pick terrible passwords

<u>username</u>	<u>password</u>
dom	fam1ly
han	dr1ftnNt0ky0
roman	Lamb0s4ever
tej	31173h4ck3r

```
check_password(username, inputted_password):  
    stored_password = accounts_table[ username ]  
    return stored_password == inputted_password
```

**problem:** adversary with access to server can get passwords

<u>username</u>	<u>hash(password)</u>
dom	e5f3c4e1694c53218978fae2c302faf4a817ce7b
han	365dab99ab03110565e982a76b22c4ff57137648
roman	ed0fa63cd3e0b9167fb48fa3c1a86d476c1e8b27
tej	0e0201a89000fe0d9f30adec170dabce8c272f7c

```

check_password (username, inputted_password):
    stored_hash = accounts_table[ username]
    inputted_hash = hash(inputted_password)
    return stored_hash == inputted_hash

```

**problem:** hashes are fast to compute, so adversary could quickly create a “rainbow table”

<u>username</u>	<u>slow_hash(password)</u>
dom	gamynjSAIeYZ4i0BT4ua03r5ub80
han	JXYWVPkpoQ6W1tbA21t6c66G4QUo
roman	Xn5U1QvQz5MG0zdfJWgF80iDFv1q
tej	1o5WIidPPZePoSyMB20.fUz3fLeZ

```

check_password (username, inputted_password):
    stored_hash = accounts_table[ username]
    inputted_hash = slow_hash(inputted_password)
    return stored_hash == inputted_hash

```

**problem:** adversary can still create rainbow tables for the most common passwords


stored in plaintext  
↓

username	salt	slow_hash(password   salt)
dom	LwVx6k04SNY3jPVf0pfYe.	M4ayLRWuzU.sSQtjoteIrIjNXI4UX
han	UbDsyTUST6d0cFpmuhWu.e	Y8ie/A18u9ymrS0FgVh9IOVx2Qe48
roman	CnfkXqUJz5C50fucP/UKIu	3GDJu07gk2iL7mFVqu0zPt3L3IITe
tej	cBGohtI6BwsaVs0SAo0u7.	8/v1Kl6rImUMYVw/.oGmA/BaRA1gC

```
check_password (username, inputted_password)
    stored_hash = accounts_table[ username ].hash
    salt = accounts_table[ username ].salt
    inputted_hash = slow_hash(inputted_password | salt)
    return stored_hash == inputted_hash
```

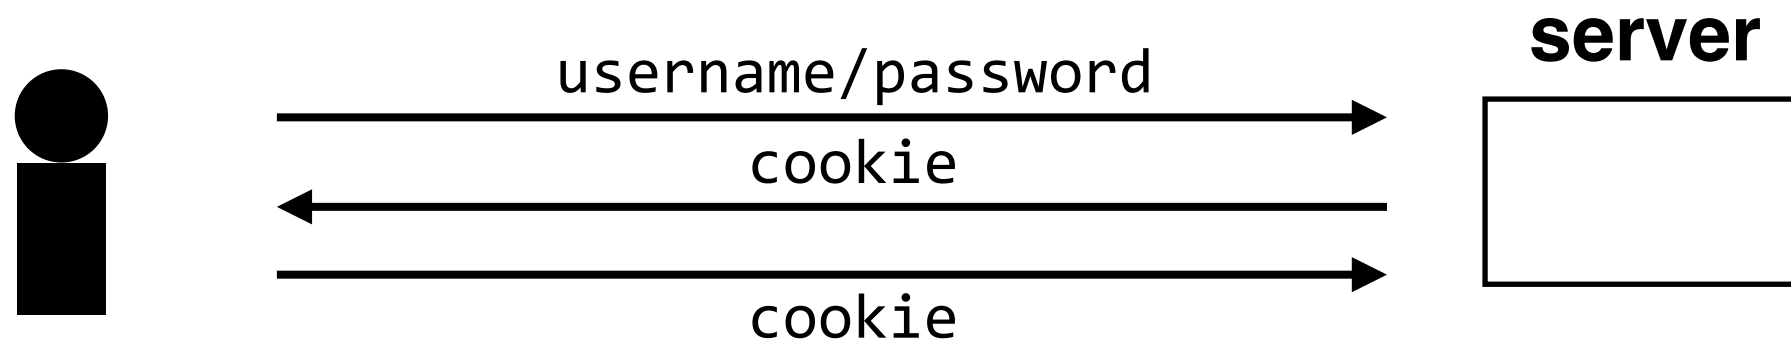
**adversary would need a separate rainbow table for every possible salt**

# how can we avoid transmitting the password over and over?



once the client has been authenticated, the server will send it a “cookie”, which it can use to keep authenticating itself for some period of time

# how can we avoid transmitting the password over and over?

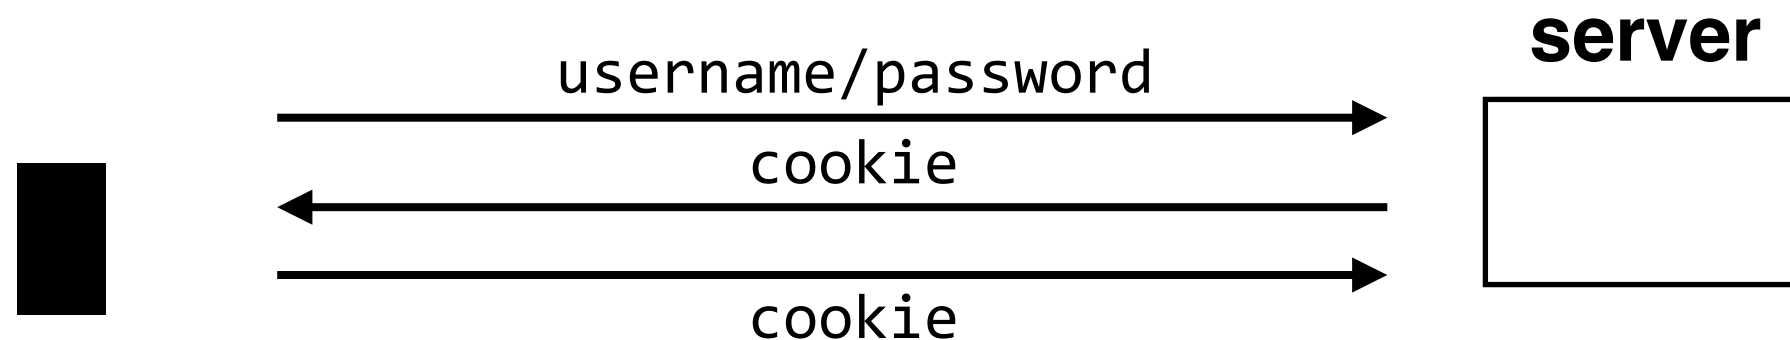


`cookie = {username, expiration} ?`

**problem:** adversaries could easily create their own cookies




# how can we avoid transmitting the password over and over?



`cookie = {username, expiration, H(username | expiration)} ?`

**problem:** adversaries could still easily create their own cookies

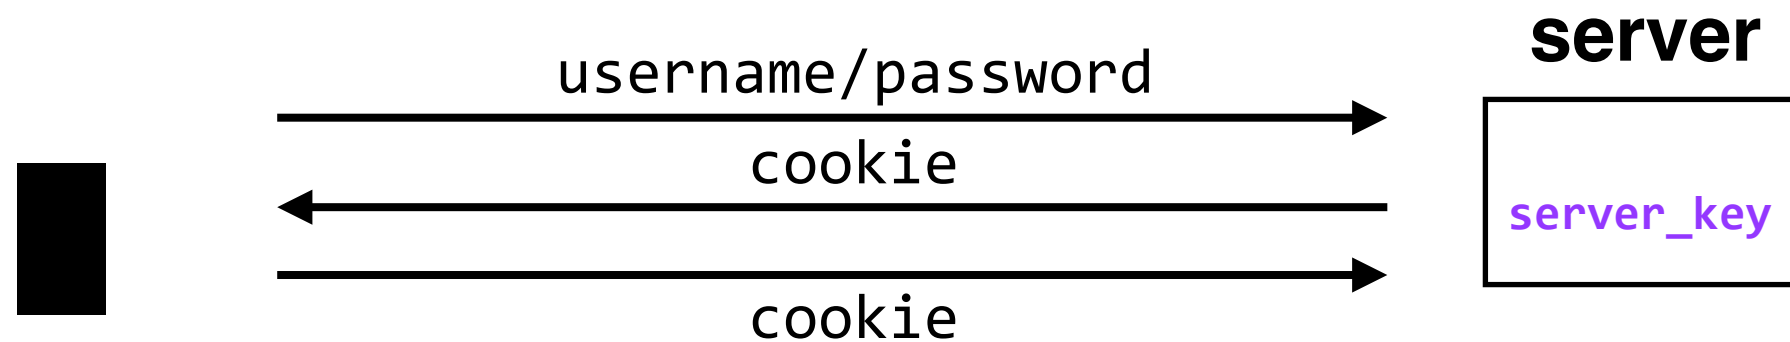
# how can we avoid transmitting the password over and over?



cookie = {username, expiration, server\_key, H(username | expiration)} ?

**problem:** adversaries could *still* easily create their own cookies


# how can we avoid transmitting the password over and over?




{username, expiration, H(server\_key | username | expiration)}

# how can we protect against phishing attacks, where an adversary tricks a user into revealing their password?

must avoid sending the password to the server entirely,  
but still allow valid servers to authenticate users



# challenge-response protocol



(random number)

458653



ccfc38b071124374ea039ff8b40e83fbf4e80d92

= H(fam1ly | 458643)


**password is never sent directly**

**valid server**

<u>username</u>	<u>password</u>
dom	fam1ly
han	dr1ftnNt0ky0
roman	Lamb0s4ever
tej	31173h4ck3r

server computes  
H(fam1ly | 458643) and  
checks

# challenge-response protocol



(random number)


458653



ccfc38b071124374ea039ff8b40e83fbf4e80d92



=  $H(\text{fam1ly} \mid 458643)$

adversary-owned server



adversary only learns  
 $H(\text{fam1ly} \mid 458643)$ ; can't  
recover the password from that

# challenge-response protocol



**password is never sent directly**

**valid server**

<u>username</u>	<u>password</u>
dom	<b>fam1ly</b>
han	dr1ftnNt0ky0
roman	Lamb0s4ever
tej	31173h4ck3r

server computes  
H(**fam1ly** | **458643**) and  
checks

**adversary-owned servers (that don't know passwords) won't learn the password; client never sends password directly**

problems arise when the server stores (salted) hashes — as it should be doing — but there are challenge-response protocols that handle that case

**how do we initially set (bootstrap) or  
reset a password?**



**are there better alternatives to  
passwords?**

- Using passwords securely takes some effort. Storing **salted hashes**, incorporating **session cookies**, dealing with **phishing**, and **bootstrapping** are all concerns.
- Thinking about how to use passwords provides more **general lessons**: consider human factors when designing secure systems, in particular.
- There are always **trade-offs**. Many “improvements” on passwords add security, but also complexity, and typically decrease usability.