

---

# SE 464

## Week 10

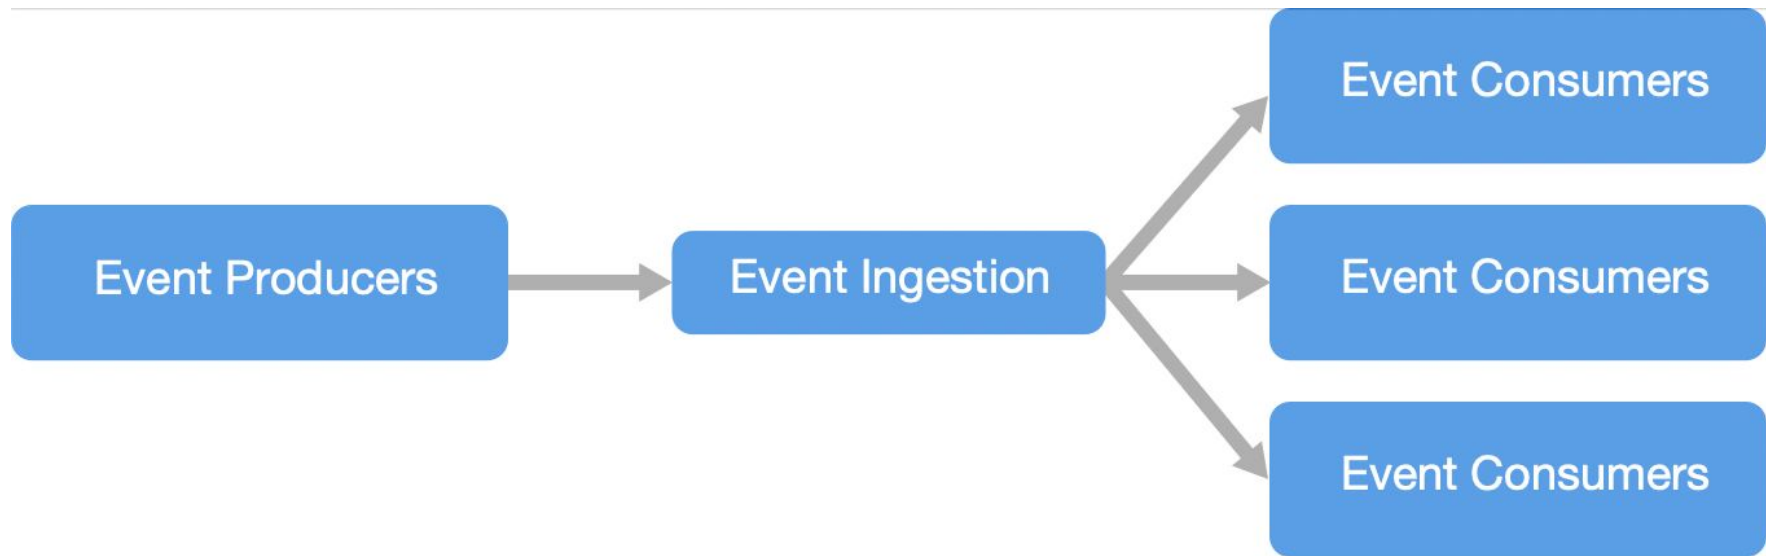
Architectures: Event-Driven, Streaming  
vs Batch, Leader-Follower

---

# Event-Driven Architecture

# Event Driven Architecture

- Definition: Architecture style that orchestrates behavior around events
- Core Concept
- Event producers and consumers
- Events as significant state changes

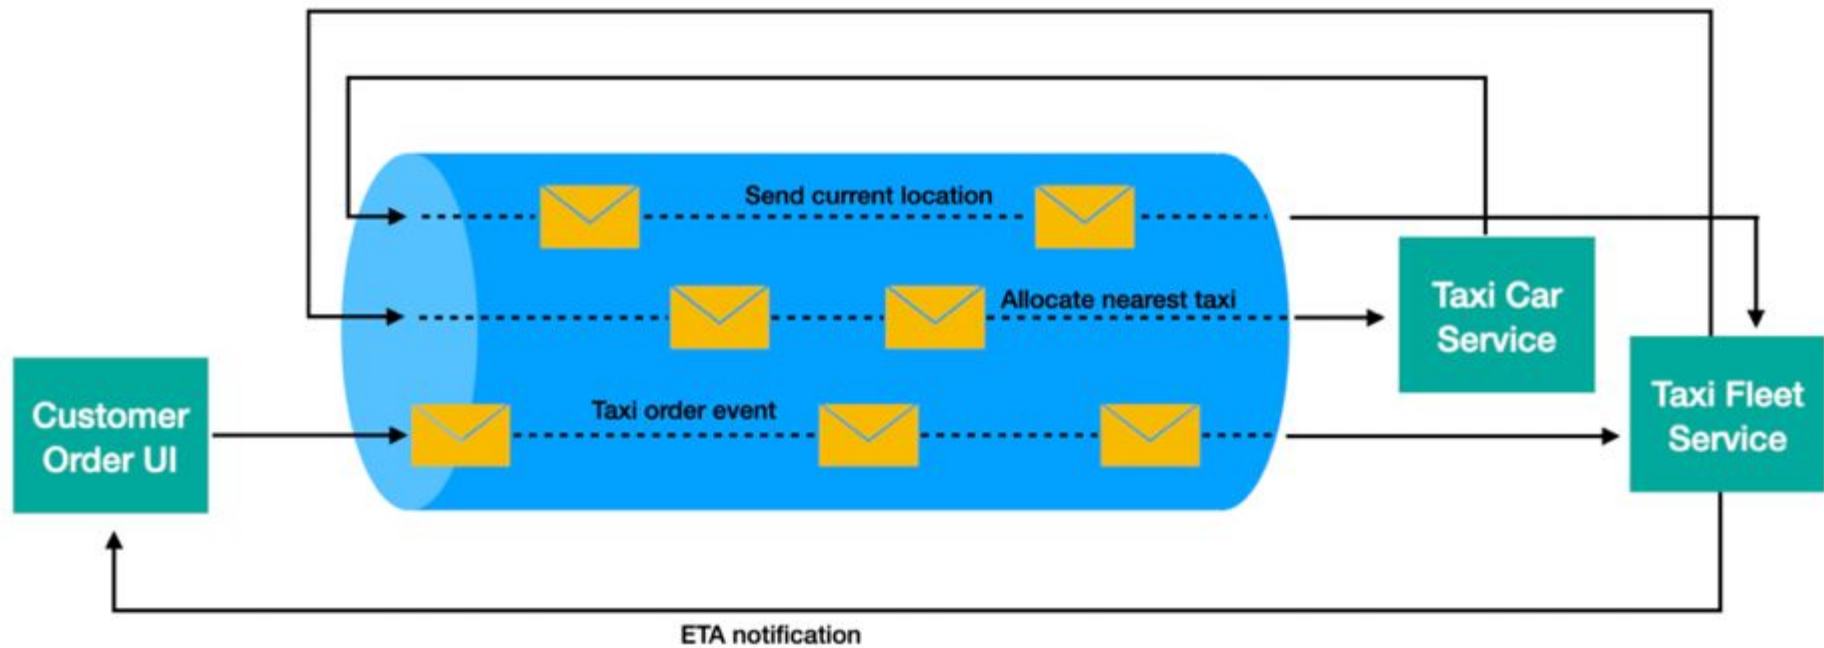


# Characteristics

- Event-Centric: Focuses on event production, detection, consumption
- Asynchronous Communication
- Decoupling: Event producers and consumers

# Use Cases

- Real-time Data Processing
  - Financial transactions, monitoring systems
- Distributed Systems
  - Microservices, cloud-based applications
- Complex Event Processing
  - Fraud detection, recommendation systems



# Benefits

- Scalability: Handles fluctuating loads effectively
- Flexibility: Easy to add/remove components
- Responsiveness: Real-time reaction to events



# Benefits

- Resilience: Reduces impact of a single component failure
- Maintainability: Simplifies maintenance and upgrades
- Improved User Experience
  - Faster, more interactive applications

# Challenges

- Event Consistency
  - Ensuring event reliability and ordering
- Complexity in Event Handling
  - Designing effective event consumers
- Integration with Legacy Systems

# Challenges

- Event Granularity
  - Balancing event size and informativeness
- Event Processing Logic
  - Distributing logic between producers and consumers

# Challenges

- Error Handling and Recovery
- Monitoring and Debugging
  - Tracing and logging event flows
- Security Concerns
  - Protecting event data in transit and at rest

# Streaming vs Batch Processing

# Big Data

- Exponential Growth of Data
  - Critical for organizations
  - Driven by technological advancements
- Big Data Processing Cycle
  - Includes: Collection, Preparation, Input, Processing, Output/Interpretation, Storage
- Importance of Data Analytics
  - Essential for extracting valuable insights
  - Supports decision-making and strategy

# Big Data

- Data Volume Increase
  - Estimated 50 trillion gigabytes (50 zettabytes) currently
  - Expected growth to 175 ZB by 2025
- Cloud Computing Influence
  - Facilitates data storage, processing, and analysis
  - Diverse Data Sources
  - Social, machine, transactional data
- <https://youtu.be/bqqCTC9nQDY?si=HICUQa-5H2jA6-mu>

# Challenges


- Security and privacy concerns
  - Managing Volume, Velocity, Variety, and Veracity
- Big Data Project Lifecycle
  - Stages: Problem Formulation, Data Collection, Data Preprocessing, Data Storage, Data Analysis, Data Visualization and Reporting, Data Quality
- Significance for Organizations
  - Enhancing growth and competitive edge



# Batch Processing

- Processing large data chunks
- Data already stored over time
- Jobs start and finish
- Processed in sequential order
- Advantages
  - Efficient for large jobs
  - Works offline, conserving resources

# Batch Processing



# Batch Processing Use Cases

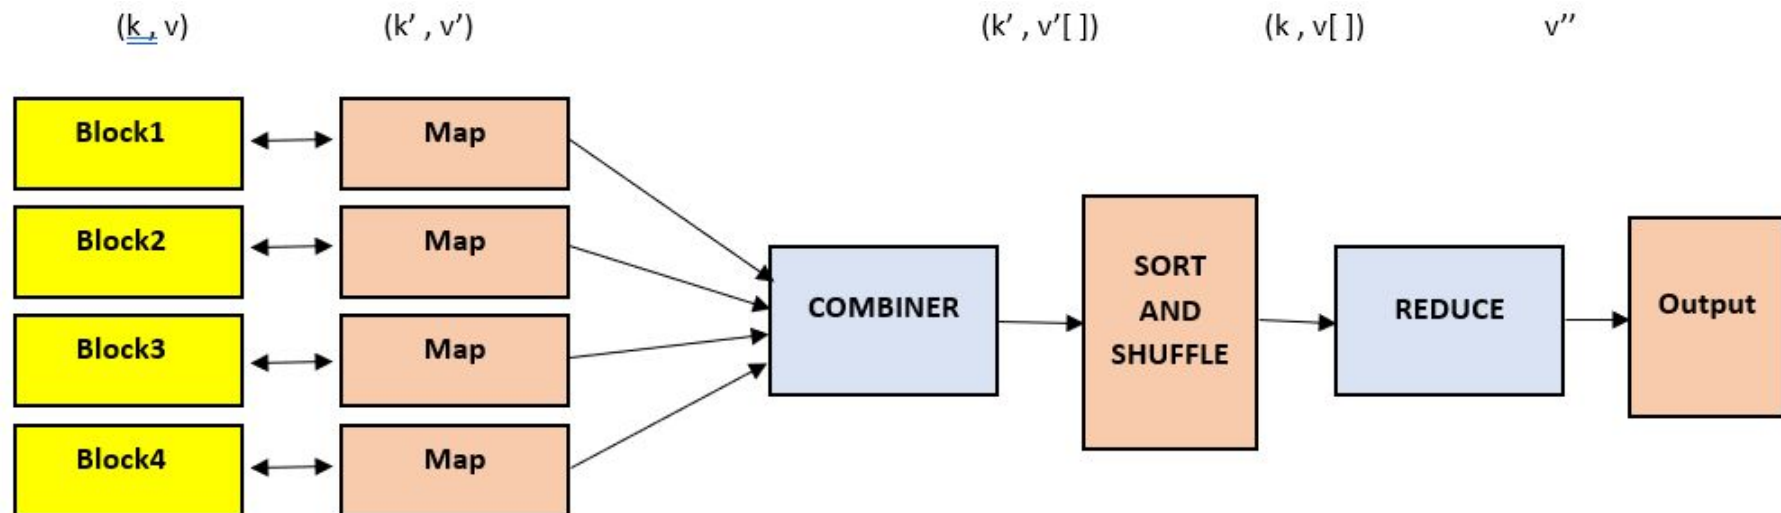
- Typical Applications
  - Financial transactions analysis
  - Payroll, billing, customer orders
- Processing Characteristics
  - High latency (minutes to hours)
  - Suitable for non-time-sensitive jobs
- Technology Examples
  - Hadoop
  - Cloudera (Hadoop distribution)

# Batch Processing in Big Data

- Crucial for structured, large-volume data
- Lifecycle Stages
  - Data collection and preparation
  - Data storage and analysis
  - Result generation and interpretation
- Challenges
  - Managing large data volumes
  - Ensuring data quality and security

# MapReduce

- Programming model for large data volumes
- Parallel processing of independent tasks
- Process
  - Input splitting
  - Map phase
  - Shuffle and sort phase
  - Reduce phase
- Key Features
  - Divide and conquer method
  - Key-value pair transformation



# Key Steps in MapReduce

- Input Splitting: Divides data into manageable pieces
- Map Phase: Applies map function to each piece
- Shuffle and Sort: Organizes output of map phase
- Reduce Phase: Aggregates and processes shuffled data
- Output
  - Final result compiled from reduce phase
- Advantages
  - Scalability
  - Fault tolerance

# Applications of MapReduce

- Large-scale data analysis
- Distributed computing tasks
- Technological Implementation
  - Hadoop: Popular framework implementing MapReduce
- Challenges and Considerations
  - Handling large data volumes
  - Optimizing performance



# Stream Processing

- Real-time data processing
- Characteristics
  - Low latency
  - Continuous data flow
- Importance
  - Immediate insights
  - Rapid decision-making

# Stream Processing

- Data Handling
  - High-velocity data streams
  - Continuous data acquisition
- Applications
  - Fraud detection
  - Online recommendations
  - Monitoring systems

# Stream Processing Challenges

- Data Volume and Velocity
  - Handling large, fast-moving data streams
- Real-time Analysis Requirements
  - Immediate data processing and analysis
- Resource Management
  - Efficient use of memory and processing power

# Apache Storm

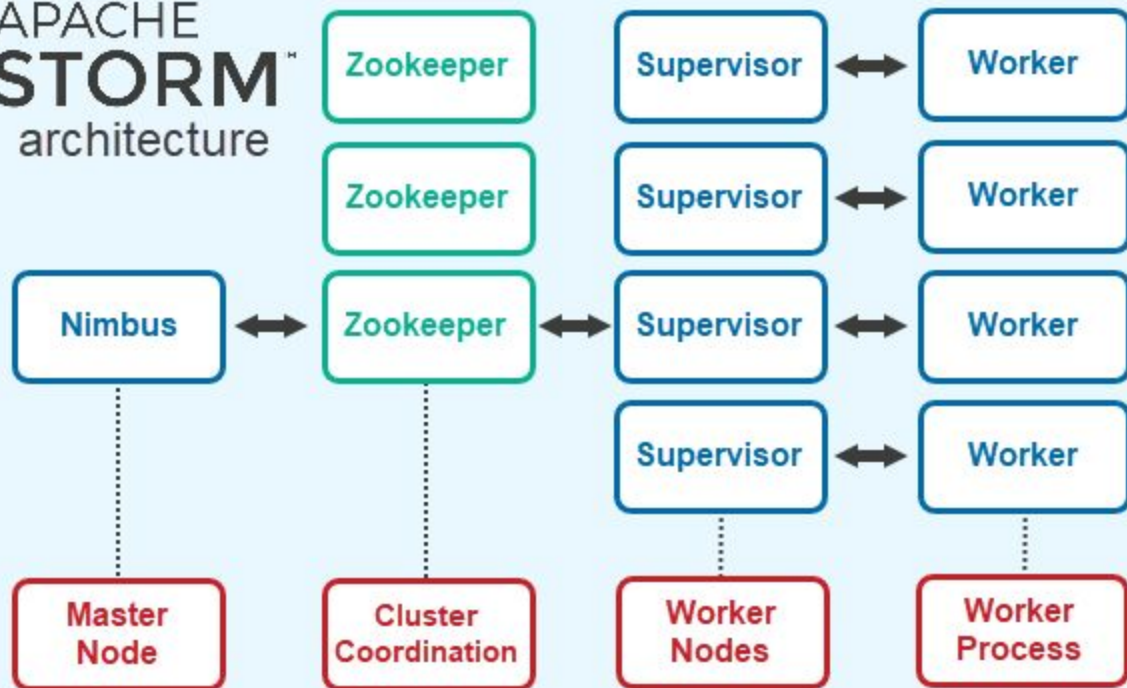
- Real-time computation system
- Processes data streams
- Key Components
  - Spouts: Data stream sources
  - Bolts: Data processing units
- Features
  - Scalability
  - Fault tolerance

# Apache Storm in Practice

- Use Cases
  - Real-time analytics
  - Distributed RPC (Remote Procedure Calls)
- Strengths
  - Easy to set up and operate
  - Integrates with various data sources
- Limitations
  - Complex state management
  - Limited data windowing capabilities



# APACHE STORM™ architecture



# Comparing Batch vs Stream

- Key Differentiators
  - Purpose
  - Data handling
  - Processing requirements
- Batch for large data volumes
- Stream for real-time analytics

# Comparing Batch vs Stream

- Batch Processing Paradigm
  - MapReduce
  - Offline analysis
- Stream Processing Paradigm
  - Stream processors
  - Online, continuous analysis



# Comparing Batch vs Stream

- Batch Processing
  - Scheduler-based job execution
  - Ideal for historical data analysis
- Stream Processing
  - Continuous data processing
  - Suited for real-time decision making

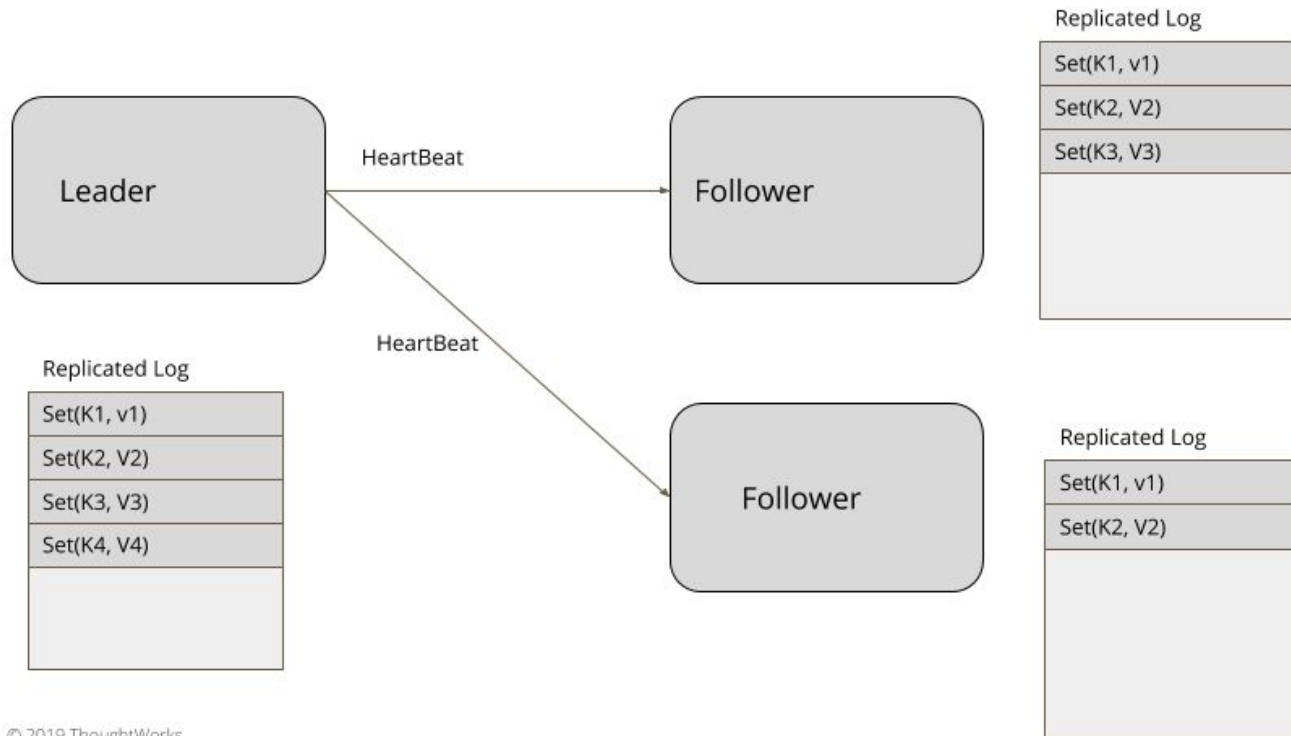
# Leader-Follower Architecture

# Leader Follower Architecture

- Definition: A concurrency pattern for distributed systems
- Core Concept
  - Dynamic role assignment among threads
  - Leader handles events, followers await their turn

# Components

- Thread Management
  - Single active leader thread
  - Pool of passive follower threads
- Event Handling
  - Leader demultiplexes and dispatches events



# Use Cases

- High-Volume Transaction Systems
  - OLTP, server-based applications
- Resource-Constrained Environments
  - Where thread overhead is a concern
- Complex Event Processing
  - Systems requiring real-time responsiveness

# Benefits

- Resource Efficiency
  - Reduced thread count
  - Lowered system overhead
- Scalability
  - Handles fluctuating workloads effectively
- Flexibility
  - Dynamic role switching enhances adaptability

# Benefits

- Simplified Synchronization
  - Reduced complexity in thread coordination
- Improved Throughput
  - Efficient event handling and processing
- Fault Tolerance
  - Leader failure leads to a new leader election



# Challenges

- Complexity in Design and Maintenance
  - Requires careful coordination of thread roles
- Bottlenecks
  - Leader thread can become a performance bottleneck
- Debugging and Monitoring
  - Tracing issues in dynamic role assignments

# Challenges

- Event Consistency and Ordering
  - Ensuring reliable event handling
- Failure Recovery Mechanisms
  - Handling leader thread failures
- Integration with Legacy Systems
  - Adapting the pattern to existing infrastructures

# Leader Election and Synchronization

- Leader Election Mechanisms
  - Techniques for assigning the leader role
- Thread Synchronization
  - Ensuring smooth role transitions
- Coordination with External Systems
  - Use of tools like Zookeeper for leader management

## Case Study: Apache Kafka

- Distributed streaming platform
- Use in handling real-time data feeds




# Leader-Follower in Kafka

- Basics of Kafka Architecture
  - Producers, Consumers, Brokers
- Leader-Follower in Brokers
  - Partition leadership among brokers

# Kafka Architecture

- Multiple brokers
- Leader and follower partitions
- Role of Zookeeper
  - Coordinates broker leadership



# Data Flow in Kafka

- Data Write Flow Diagram
  - Producers to leader partitions
  - Replication to followers
- Data Read Flow
  - Consumers reading from leaders or followers



# Leader Election in Kafka

- Mechanism of Leader Election
  - Broker failure scenarios
  - Zookeeper's role in electing new leaders
- Consistency Guarantees
  - How Kafka ensures data consistency during leadership changes

# Replication and Fault Tolerance

- Replication Strategy
  - In-sync replica sets
- Fault Tolerance
  - Handling broker downtimes
  - Ensuring data availability and integrity

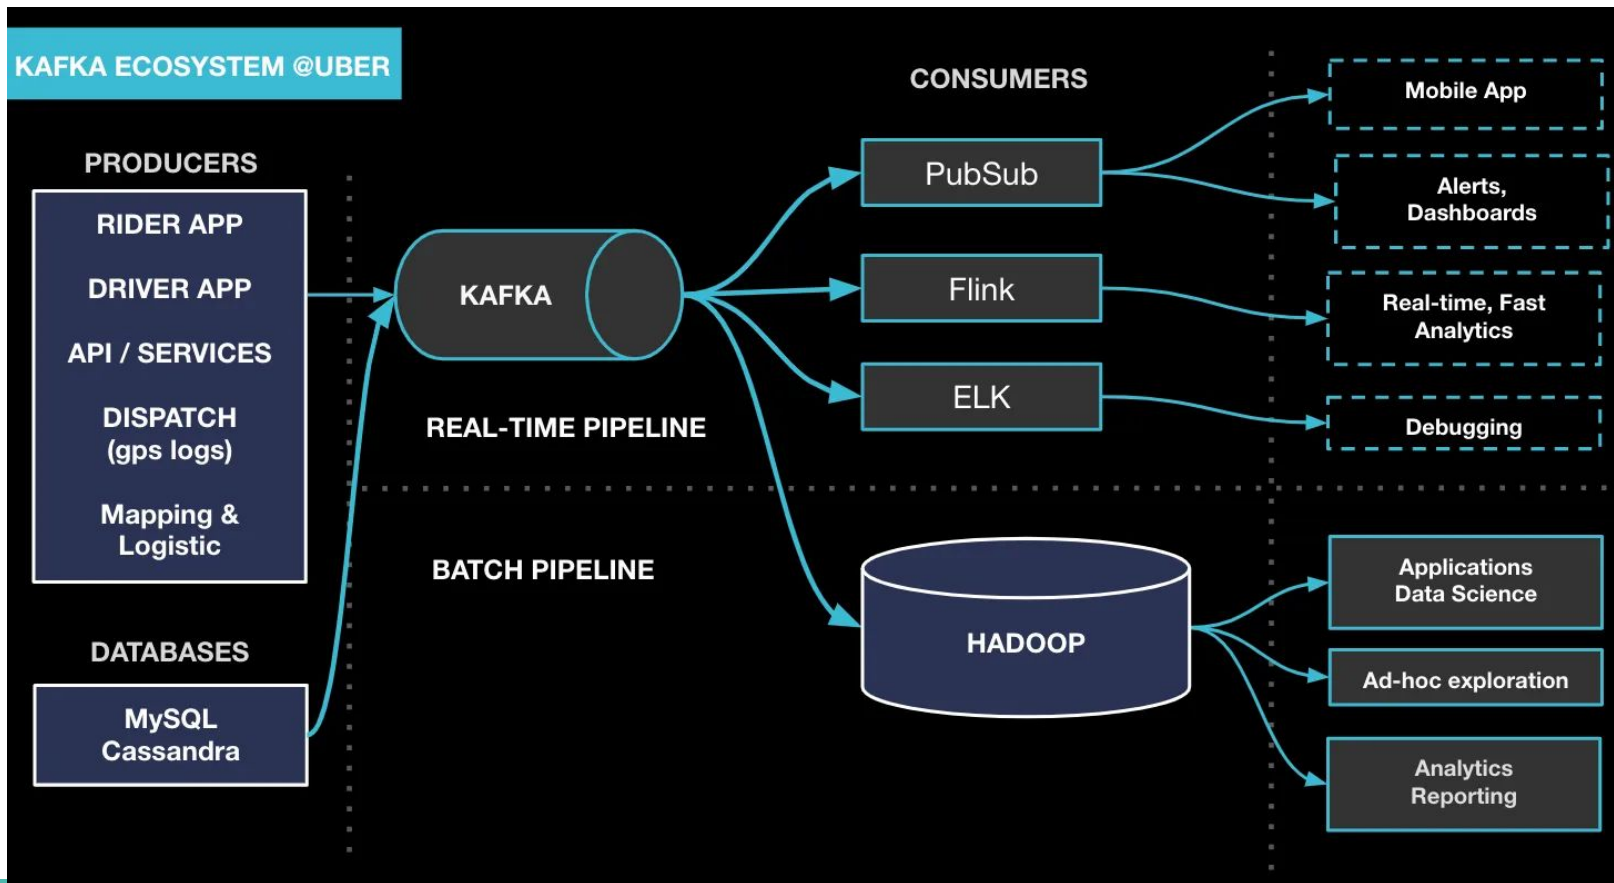
# Challenges

- Leader-Follower Impact on Throughput
  - Balancing load between leaders and followers
- Latency in Leader-Follower Replication
  - Factors affecting latency
- Bottlenecks in Leader Performance
  - High traffic and large data volumes
- Handling Follower Lag
  - Strategies to keep followers in sync

# Example Use Case for Kafka

- Scenario: E-commerce Platform Analytics
  - Real-time tracking of user behavior and transactions
  - Aggregation of data for instant analytics
- Kafka's Role
  - High-throughput data ingestion from various sources
  - Reliable data streaming to analytics engines
  - Ensures data consistency and availability
  - Facilitates rapid data processing and decision making

# Another Real-Life Example



# Sources

- <https://learn.microsoft.com/en-us/azure/architecture/guide/architecture-styles/event-driven>
- <https://www.dre.vanderbilt.edu/~schmidt/PDF/lf.pdf>
- <https://kafka.apache.org/21/documentation/streams/architecture.html>